Eleven-family LLM evaluation

Complete results, comparability audit, and paper preparation reference

Prepared 20 October 2026 from research-3-main.zip. Scope: 11 family folders, 37 checkpoints, 14,800 responses, 200 unique questions. This is an archival pilot report using the original saved scoring; no answers have been rescored and no inference has been rerun.

Main finding: all families used the same questions and prompt texts, but only 6,673/14,800 responses (45.1%) are usable under the recorded parser and truncation policy. The archive supports an evaluation-quality audit and selected descriptive comparisons. It does not yet support a causal decomposition of shared failure into difficulty, family, and lineage.

Reading guide: Sections 1-5 provide the audit, comparison figures, and interpretation. Section 6 contains each family and every model-condition accuracy. Section 7 reproduces all rows and columns of the saved pair, ladder, team, repeat, and baseline tables. Section 8 documents formulas, provenance, and next steps. NA means undefined or unavailable, never zero. PDF wide-table panels share row numbers; model aliases are defined below. The Markdown preserves original column names and numeric values to six significant digits.

1. Inventory and integrity

Run totals (family means are checkpoint-set dependent)

Column panel 1/2 - row numbers identify the same source rows.

Row	family	models	responses	usable	usable pct	truncated
1	Falcon	3	1200	535	44.5833	271
2	GLM	2	800	289	36.125	26
3	Gemma	3	1200	1101	91.75	15
4	Granite	2	800	438	54.75	96
5	InternLM	2	800	199	24.875	1
6	Llama	4	1600	628	39.25	107
7	Mistral	3	1200	269	22.4167	350
8	OLMo	3	1200	222	18.5	28
9	Phi	4	1600	800	50	51
10	Qwen	9	3600	1894	52.6111	234
11	Yi	2	800	298	37.25	189

Column panel 2/2 - row numbers identify the same source rows.

Row	truncated pct	correct	pipeline pct	tokens
1	22.5833	398	33.1667	831709
2	3.25	202	25.25	226116
3	1.25	817	68.0833	361092
4	12	210	26.25	357524
5	0.125	143	17.875	49047
6	6.6875	409	25.5625	923440
7	29.1667	117	9.75	786224
8	2.33333	117	9.75	240668
9	3.1875	593	37.0625	1165781
10	6.5	1542	42.8333	3406079
11	23.625	195	24.375	514969

Archive SHA-256: 7ffe3e16292ebdd5ba6e8c0c3c5f37b0554953a068d0ca209eb619772b7afd88. Across all runs: no duplicate model/question/prompt/seed/temperature records; all raw job-ID multisets match scored records; all saved accuracy values reconcile with recomputation; every file listed by each manifest is present and its hash matches. These checks establish internal consistency, not scientific validity of the parser or metadata.

Question-file SHA-256 is identical across all 11 runs: 93324a13e197e5ad65f52ec652a057556baec69281b2bfbd1c48605d7cbf1133. Each file contains 100 GSM8K and 100 MMLU-Pro questions, used in both prompt templates. The unrendered prompt text is identical across families for each item/template. Native chat wrappers vary by model as intended. There are two worker-code hashes: Mistral uses the token-ID formatting patch; the other runs share one worker version.

Integrity checklist

Column panel 1/2 - row numbers identify the same source rows.

Row	family	responses	raw rows	duplicates	accuracy mismatches	manifest mismatches
1	Falcon	1200	1200	0	0	0
2	GLM	800	800	0	0	0
3	Gemma	1200	1200	0	0	0
4	Granite	800	800	0	0	0
5	InternLM	800	800	0	0	0
6	Llama	1600	1600	0	0	0
7	Mistral	1200	1200	0	0	0
8	OLMo	1200	1200	0	0	0
9	Phi	1600	1600	0	0	0

Row	family	responses	raw rows	duplicates	accuracy mismatches	manifest mismatches
10	Qwen	3600	3600	0	0	0
11	Yi	800	800	0	0	0

Column panel 2/2 - row numbers identify the same source rows.

Row	raw scored match
1	True
2	True
3	True
4	True
5	True
6	True
7	True
8	True
9	True
10	True
11	True

2. Protocol, interpretation, and exclusions

Every checkpoint produced 400 responses: 2 benchmarks x 100 questions x 2 templates x one seed (17) x one temperature (0.0). This is a limited question sample, not a full-benchmark evaluation. Direct and reasoned outputs share questions and are paired, not independent replicates. A single seed does not measure sampling variability. Families have unequal checkpoint counts and mix base, instruction, reasoning, and different model generations; family averages are not controlled capability comparisons.

Dataset revisions

Row	benchmark	repo	revision	config	split	fingerprint
1	mmlu_pro	TIGER-Lab/MMLU-Pro	b189ec765aa7ed75c8acfea42df31fdae71f97be	NA	test	c54621f4b1475bdf
2	gsm8k	openai/gsm8k	740312add88f781978c0658806c59bc2815b9866	main	test	0b3e31f6665214ba

All configs set fit_bayesian=False, heldout=False, verification=False, and allow_truncated_analysis=False. The saved baseline files nevertheless contain a separate 50/50 item-holdout baseline routine where eligible teams exist. This is distinct from hierarchical held-out family/lineage validation, which was not performed. Self-consistency is absent because there is only one seed and temperature zero. Team sizes requested were 3 and 5.

Standalone DeepSeek, SmolLM, Pythia, and gpt-oss family folders are absent as requested. Three DeepSeek-R1-distilled checkpoints remain: two under Qwen and one under Llama. Thus standalone-family exclusion does not imply exclusion of DeepSeek-derived models. All three are marked DeepSeek-R1 lineage_verified=True; the other 34 checkpoints have unknown/unaudited lineage. Their labels are recorded metadata, not an independent provenance verification in this report. A teacher checkpoint is not present, so direct student-teacher error inheritance cannot be measured.

Family labels are developer groupings from the archive: Gemma includes Gemma 2 and 3; Phi includes Phi 3.5 and 4; Llama includes 3.1, 3.2 and an R1 distill; Qwen includes 2.5, Math, 3 and R1 distills. These are broader than a single training recipe. Every design audit flags lineage coverage/rank problems. Do not convert unknown lineage into one shared ancestor or infer a family/lineage variance component from pairwise correlations.

3. Evaluation quality and model-level comparisons

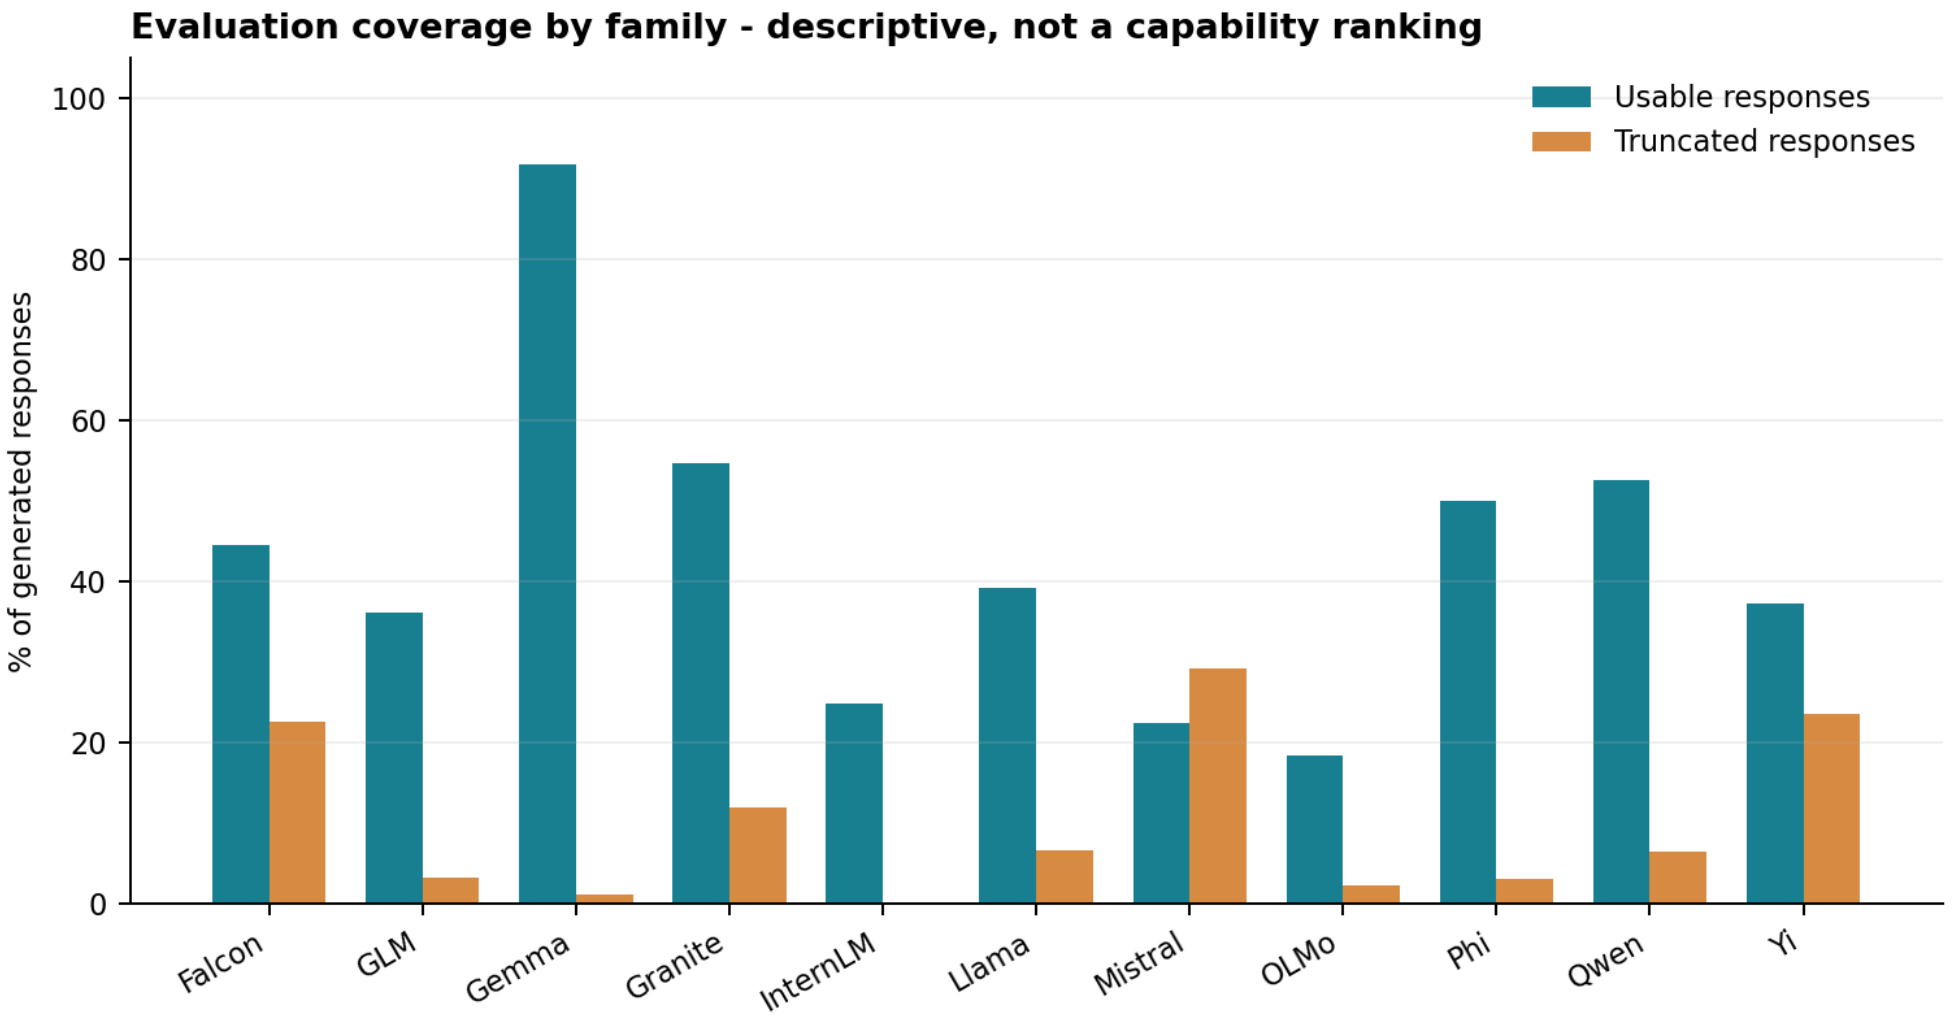
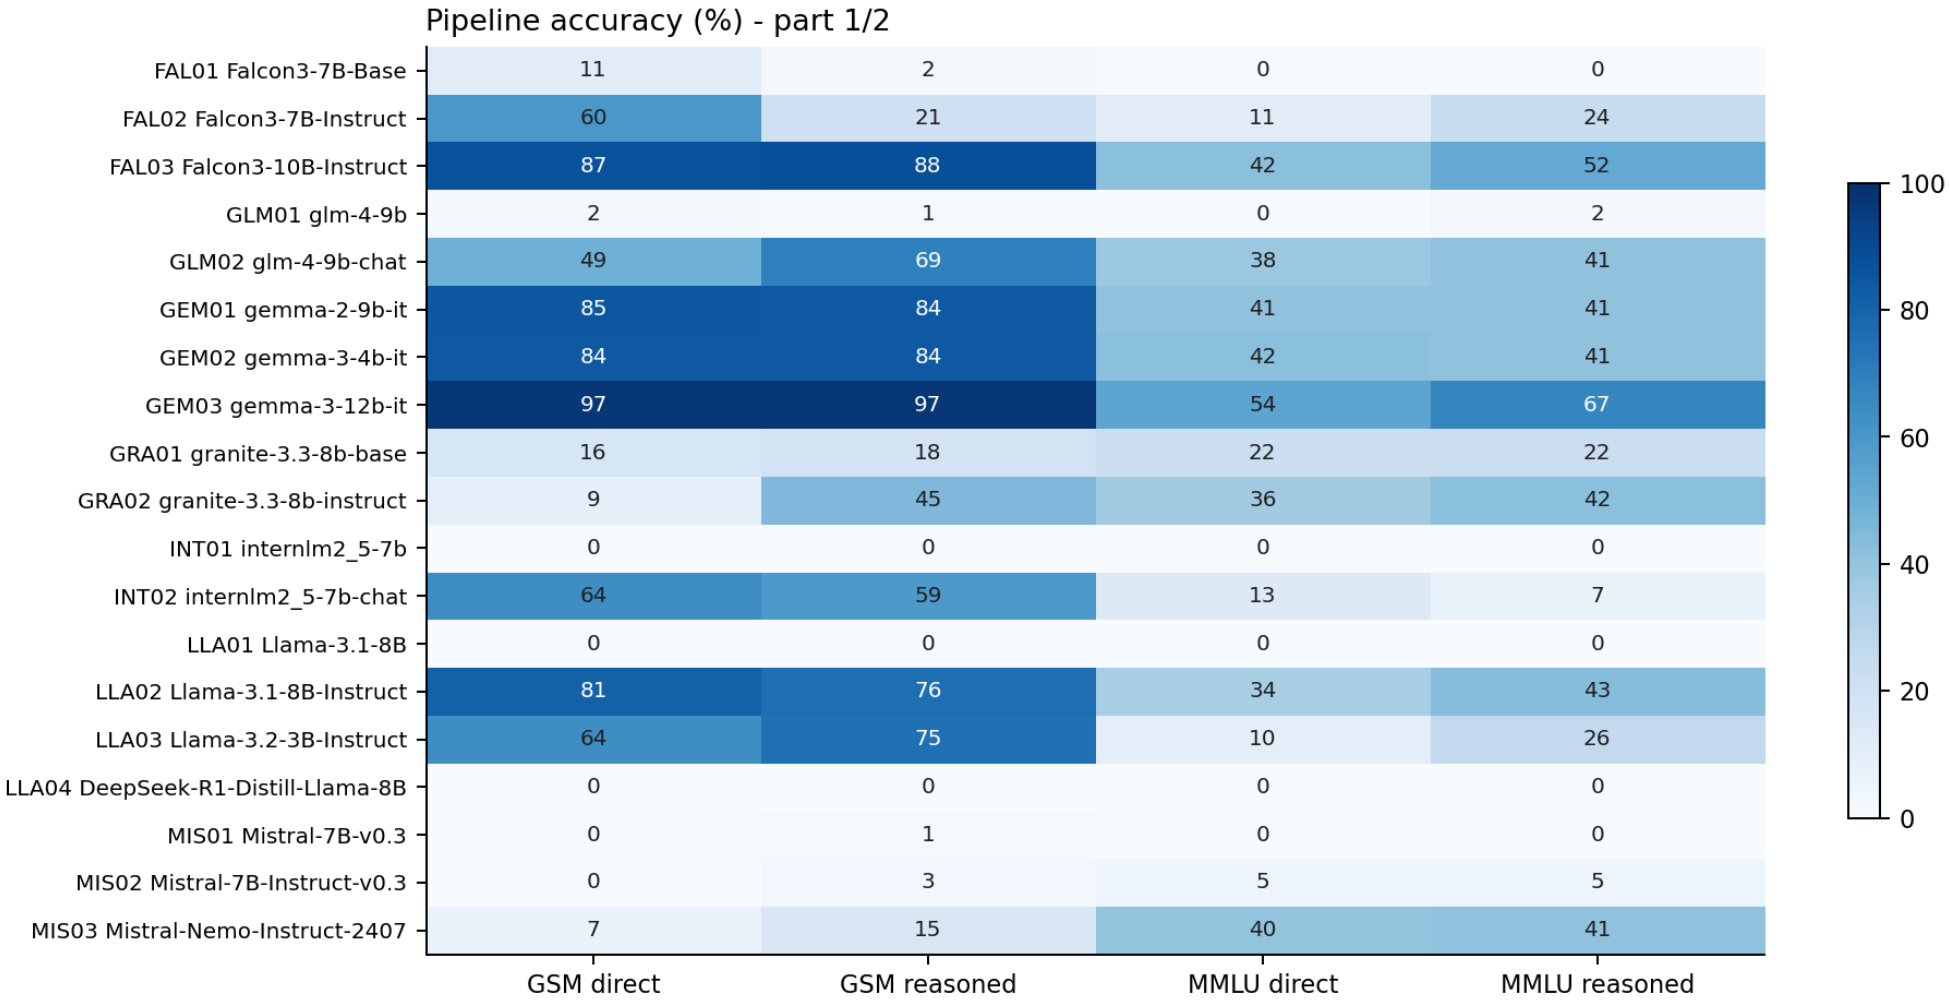


Figure 1. Usability and truncation by family. Usable means parsed and not truncated; failures overlap, so bars need not sum to 100%.

Overall, 1,368 responses (9.2%) are truncated. Parsing rejects additional responses, including some with clearly identifiable answers followed by units or explanations. Usable-only accuracy conditions on model-dependent formatting success and can be selection-biased. Pipeline accuracy is an operational outcome for this prompt/parser combination, not a clean estimate of intrinsic model capability.


Pipeline accuracy (%) - part 1/2
FAL01 Falcon3-7B-Base 11 2 0 0
FAL02 Falcon3-7B-Instruct 60 21 11 24
FAL03 Falcon3-10B-Instruct 87 88 42 52
GLM01 glm-4-9b 2 1 0 2
GLM02 glm-4-9b-chat 49 69 38 41
GEM01 gemma-2-9b-it 85 84 41 41
GEM02 gemma-3-4b-it 84 84 42 41
GEM03 gemma-3-12b-it 97 97 54 67
GRA01 granite-3.3-8b-base 16 18 22 22
GRA02 granite-3.3-8b-instruct 9 45 36 42
INT01 internlm2_5-7b 0 0 0 0
INT02 internlm2_5-7b-chat 64 59 13 7
LLA01 Llama-3.1-8B 0 0 0 0
LLA02 Llama-3.1-8B-Instruct 81 76 34 43
LLA03 Llama-3.2-3B-Instruct 64 75 10 26
LLA04 DeepSeek-R1-Distill-Llama-8B 0 0 0 0
MIS01 Mistral-7B-v0.3 0 1 0 0
MIS02 Mistral-7B-Instruct-v0.3 0 3 5 5
MIS03 Mistral-Nemo-Instruct-2407 7 15 40 41
GSM direct
GSM reasoned
MMLU direct
MMLU reasoned
100
80
60
40
20
0

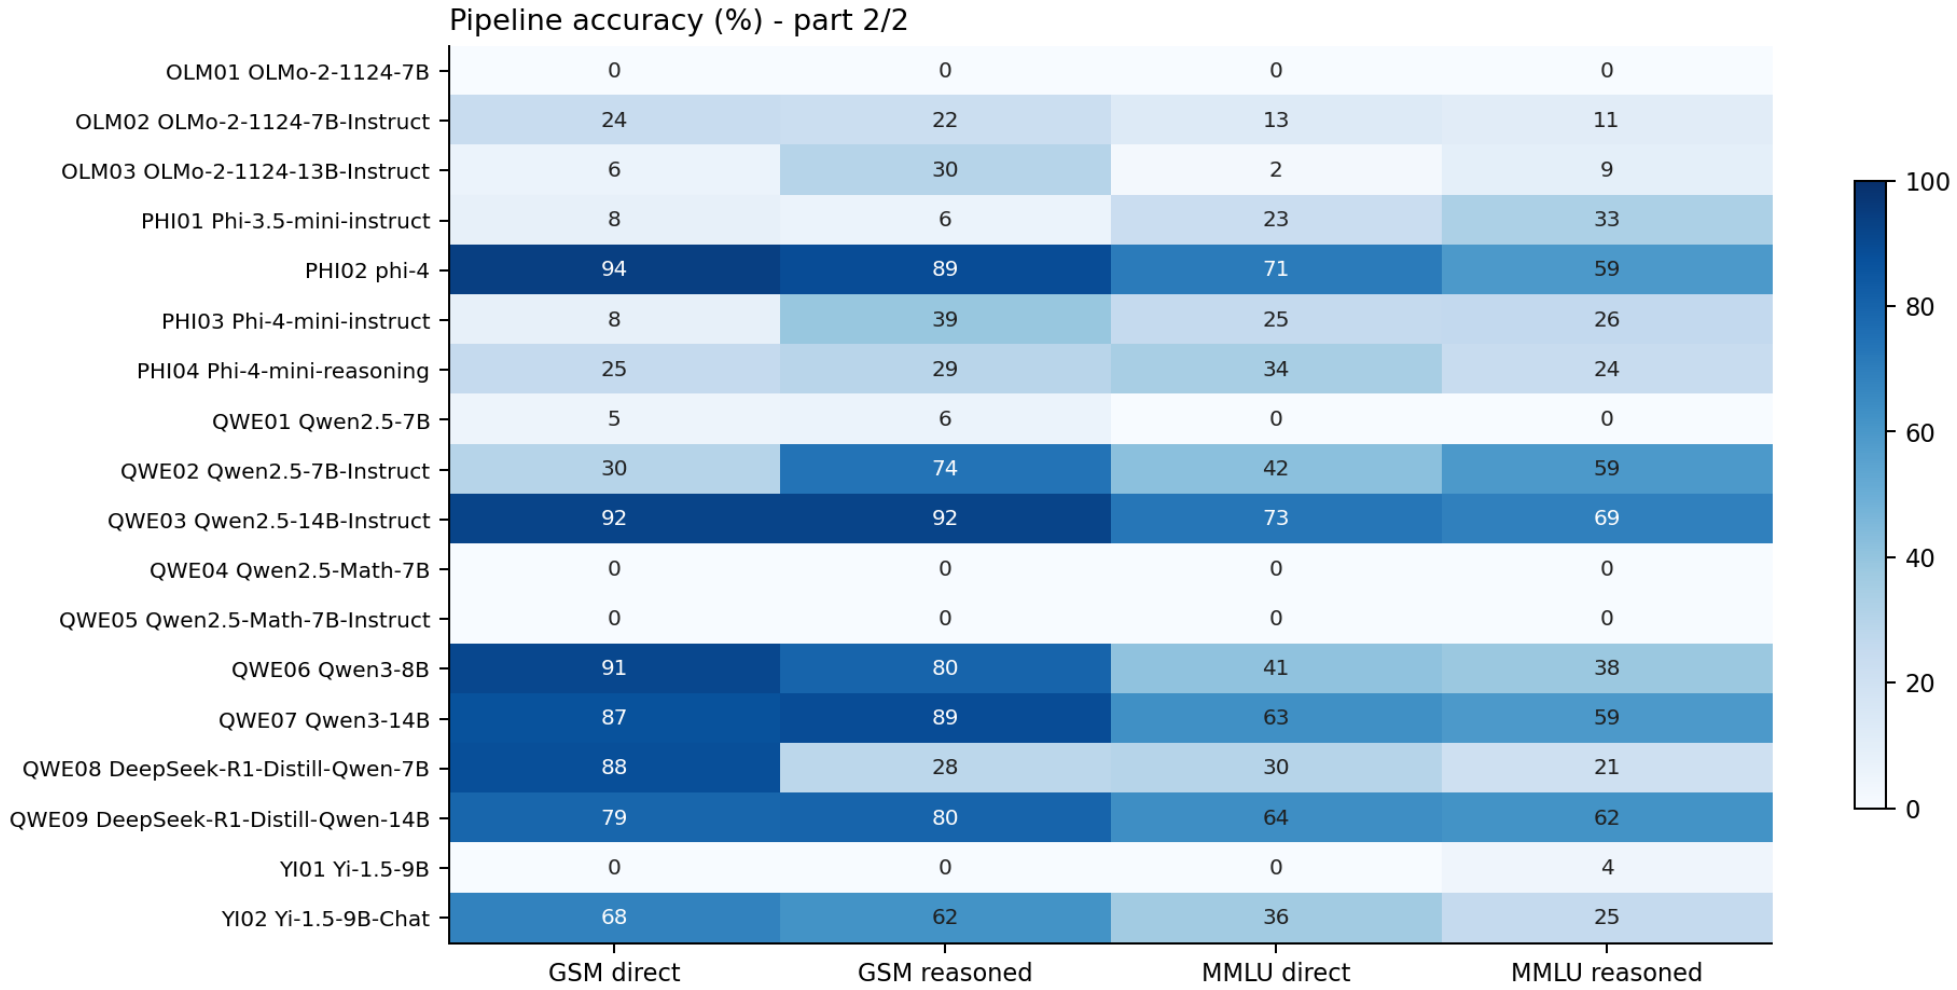
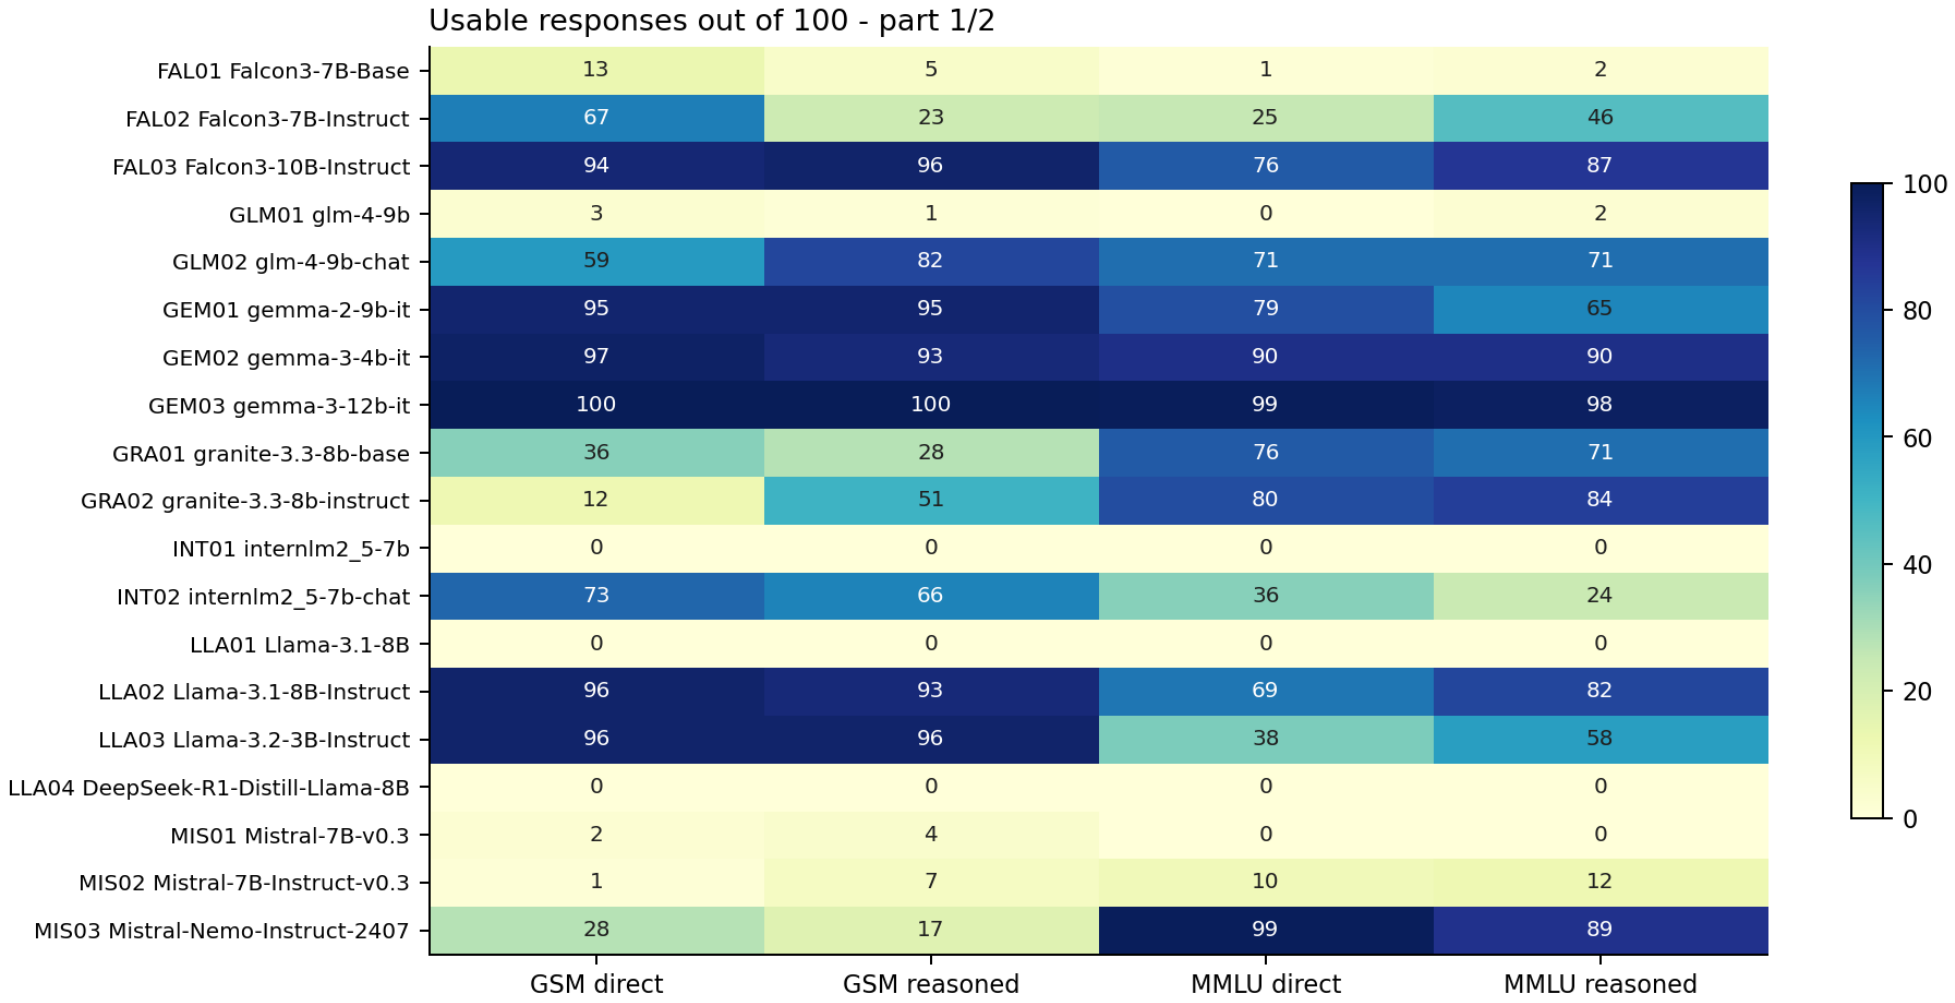


Figure 2. All 37 checkpoints: pipeline accuracy across the four conditions. Read with Figure 3; low coverage can make an otherwise capable model appear weak.


Usable responses out of 100 - part 1/2
FAL01 Falcon3-7B-Base 13 5 1 2
FAL02 Falcon3-7B-Instruct 67 23 25 46
FAL03 Falcon3-10B-Instruct 94 96 76 87
GLM01 glm-4-9b 3 1 0 2
GLM02 glm-4-9b-chat 59 82 71 71
GEM01 gemma-2-9b-it 95 95 79 65
GEM02 gemma-3-4b-it 97 93 90 90
GEM03 gemma-3-12b-it 100 100 99 98
GRA01 granite-3.3-8b-base 36 28 76 71
GRA02 granite-3.3-8b-instruct 12 51 80 84
INT01 internlm2_5-7b 0 0 0 0
INT02 internlm2_5-7b-chat 73 66 36 24
LLA01 Llama-3.1-8B 0 0 0 0
LLA02 Llama-3.1-8B-Instruct 96 93 69 82
LLA03 Llama-3.2-3B-Instruct 96 96 38 58
LLA04 DeepSeek-R1-Distill-Llama-8B 0 0 0 0
MIS01 Mistral-7B-v0.3 2 4 0 0
MIS02 Mistral-7B-Instruct-v0.3 1 7 10 12
MIS03 Mistral-Nemo-Instruct-2407 28 17 99 89
GSM direct
GSM reasoned
MMLU direct
MMLU reasoned
100
80
60
40
20
0

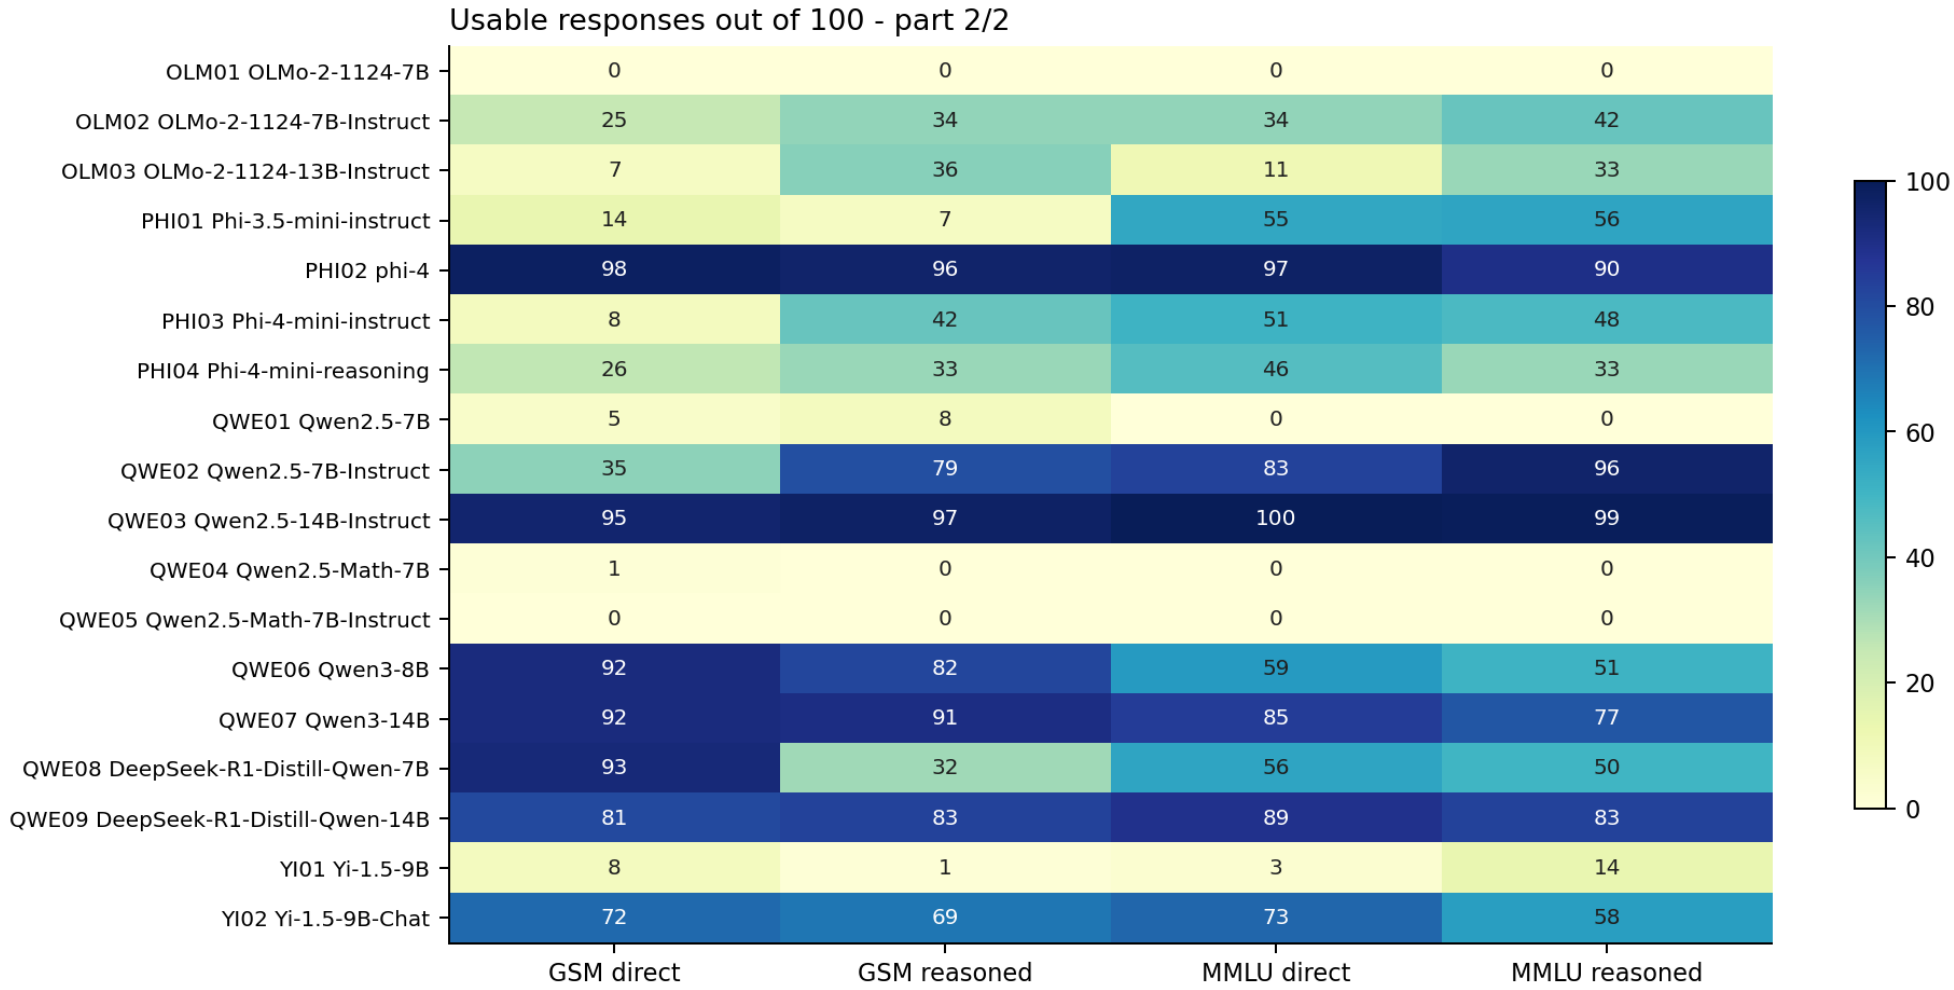


Figure 3. Usable counts out of 100 for each condition. This report uses >=90 as a descriptive screening flag, not a validated adequacy threshold.

Higher-coverage observations

All model-conditions with at least 90 usable responses

Column panel 1/2 - row numbers identify the same source rows.

Row	alias	model	benchmark	template	n	usable
1	FAL03	FAL03	gsm8k	direct	100	94
2	FAL03	FAL03	gsm8k	reasoned	100	96
3	GEM01	GEM01	gsm8k	direct	100	95
4	GEM01	GEM01	gsm8k	reasoned	100	95
5	GEM03	GEM03	gsm8k	direct	100	100
6	GEM03	GEM03	gsm8k	reasoned	100	100
7	GEM03	GEM03	mmlu_pro	direct	100	99
8	GEM03	GEM03	mmlu_pro	reasoned	100	98
9	GEM02	GEM02	gsm8k	direct	100	97
10	GEM02	GEM02	gsm8k	reasoned	100	93
11	GEM02	GEM02	mmlu_pro	direct	100	90
12	GEM02	GEM02	mmlu_pro	reasoned	100	90
13	LLA02	LLA02	gsm8k	direct	100	96
14	LLA02	LLA02	gsm8k	reasoned	100	93
15	LLA03	LLA03	gsm8k	direct	100	96
16	LLA03	LLA03	gsm8k	reasoned	100	96
17	MIS03	MIS03	mmlu_pro	direct	100	99
18	PHI02	PHI02	gsm8k	direct	100	98
19	PHI02	PHI02	gsm8k	reasoned	100	96
20	PHI02	PHI02	mmlu_pro	direct	100	97
21	PHI02	PHI02	mmlu_pro	reasoned	100	90
22	QWE03	QWE03	gsm8k	direct	100	95
23	QWE03	QWE03	gsm8k	reasoned	100	97
24	QWE03	QWE03	mmlu_pro	direct	100	100
25	QWE03	QWE03	mmlu_pro	reasoned	100	99
26	QWE02	QWE02	mmlu_pro	reasoned	100	96
27	QWE07	QWE07	gsm8k	direct	100	92
28	QWE07	QWE07	gsm8k	reasoned	100	91
29	QWE06	QWE06	gsm8k	direct	100	92
30	QWE08	QWE08	gsm8k	direct	100	93

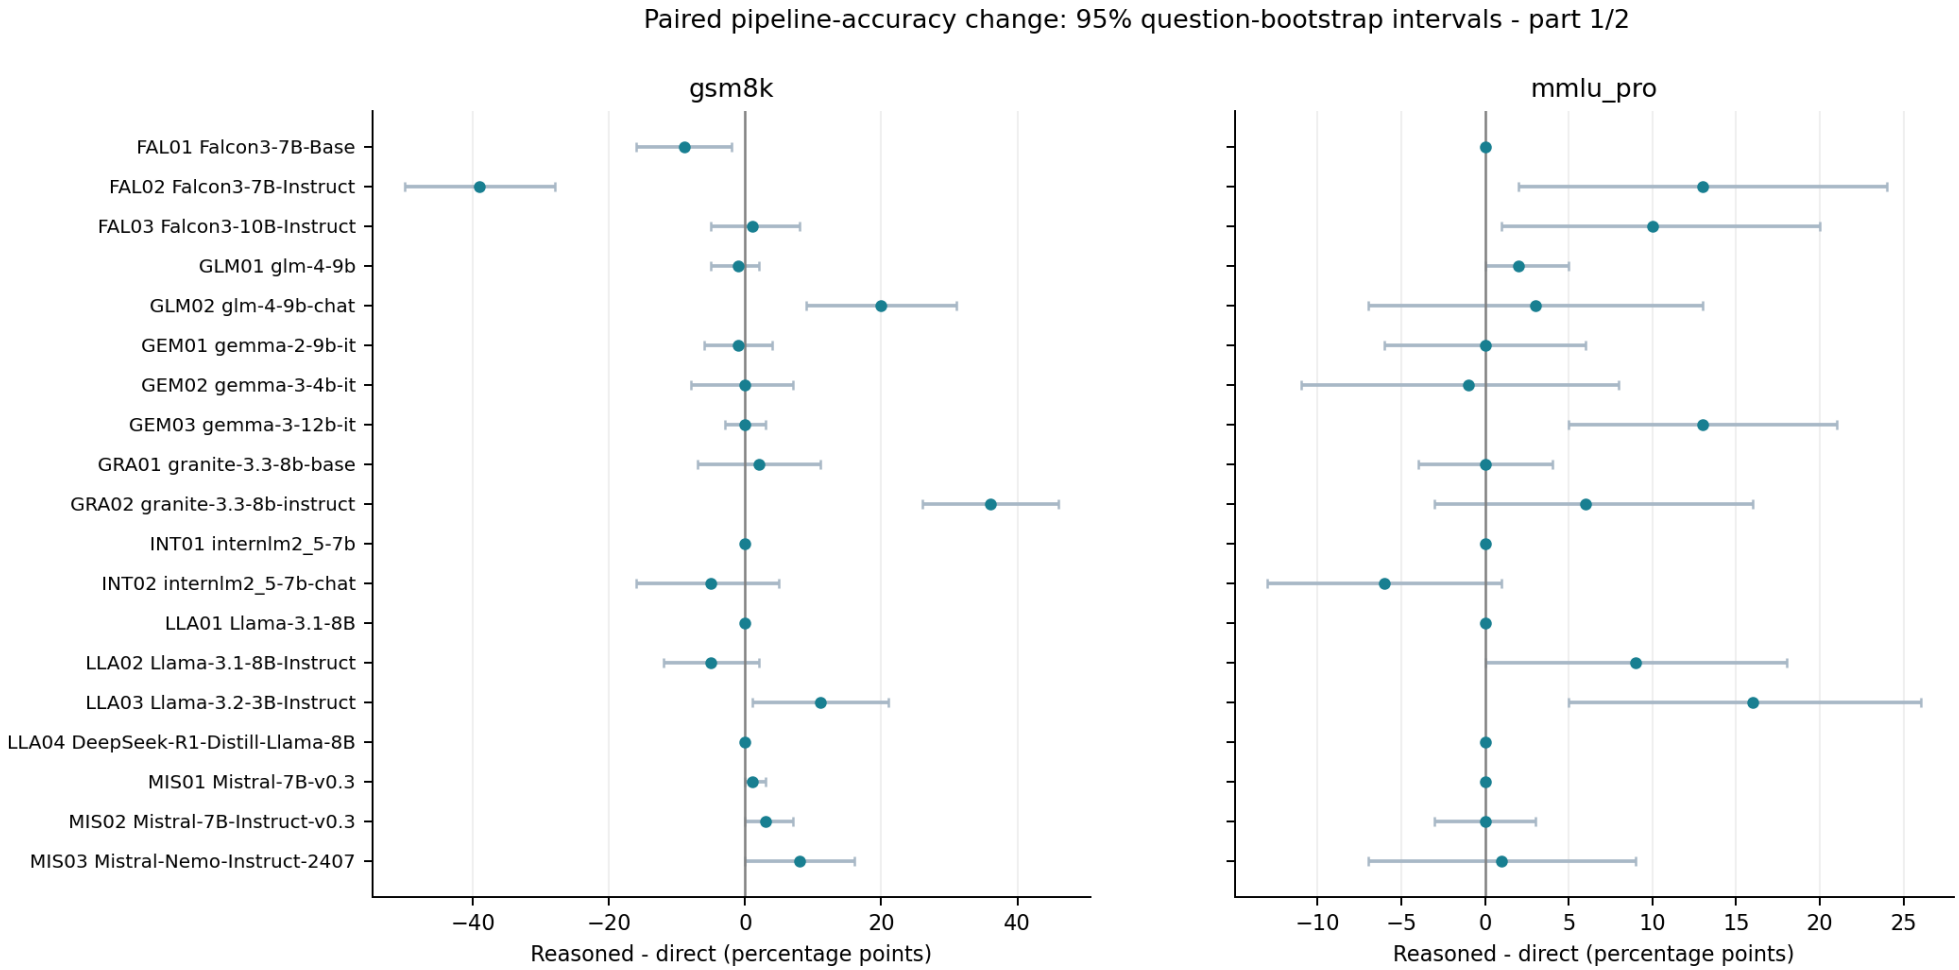
Column panel 2/2 - row numbers identify the same source rows.

Row	pipeline accuracy	usable accuracy	pipeline wilson low	pipeline wilson high
1	0.87	0.925532	0.790196	0.922428
2	0.88	0.916667	0.801879	0.930006
3	0.85	0.894737	0.767164	0.90694
4	0.84	0.884211	0.755797	0.899047
5	0.97	0.97	0.915481	0.989745
6	0.97	0.97	0.915481	0.989745
7	0.54	0.545455	0.442649	0.634392
8	0.67	0.683673	0.573053	0.754369
9	0.84	0.865979	0.755797	0.899047
10	0.84	0.903226	0.755797	0.899047
11	0.42	0.466667	0.327984	0.517935
12	0.41	0.455556	0.318673	0.507986
13	0.81	0.84375	0.722212	0.874852
14	0.76	0.817204	0.667677	0.833087
15	0.64	0.666667	0.542354	0.727288
16	0.75	0.78125	0.656955	0.824548
17	0.4	0.40404	0.309401	0.497997
18	0.94	0.959184	0.875232	0.972214
19	0.89	0.927083	0.813687	0.937458
20	0.71	0.731959	0.614611	0.789852
21	0.59	0.655556	0.492014	0.681327
22	0.92	0.968421	0.850019	0.958907
23	0.92	0.948454	0.850019	0.958907
24	0.73	0.73	0.635679	0.807304
25	0.69	0.69697	0.593739	0.772203
26	0.59	0.614583	0.492014	0.681327
27	0.87	0.945652	0.790196	0.922428
28	0.89	0.978022	0.813687	0.937458

Row	pipeline accuracy	usable accuracy	pipeline wilson low	pipeline wilson high
29	0.91	0.98913	0.837738	0.951927
30	0.88	0.946237	0.801879	0.930006

On this sample, Gemma-3-12B-it records 97/100 pipeline-correct GSM8K answers under both templates with 100 usable answers each. On MMLU-Pro direct, Qwen2.5-14B-Instruct records 73/100 with 100 usable, and phi-4 records 71/100 with 97 usable. On MMLU-Pro reasoned, Qwen2.5-14B-Instruct records 69/100 with 99 usable, and Gemma-3-12B-it records 67/100 with 98 usable. These are descriptive sample results, not established rankings or full-benchmark scores; the same questions induce dependence and uncertainty remains substantial at n=100.

4. Prompt effects and pooled shared-failure exploration



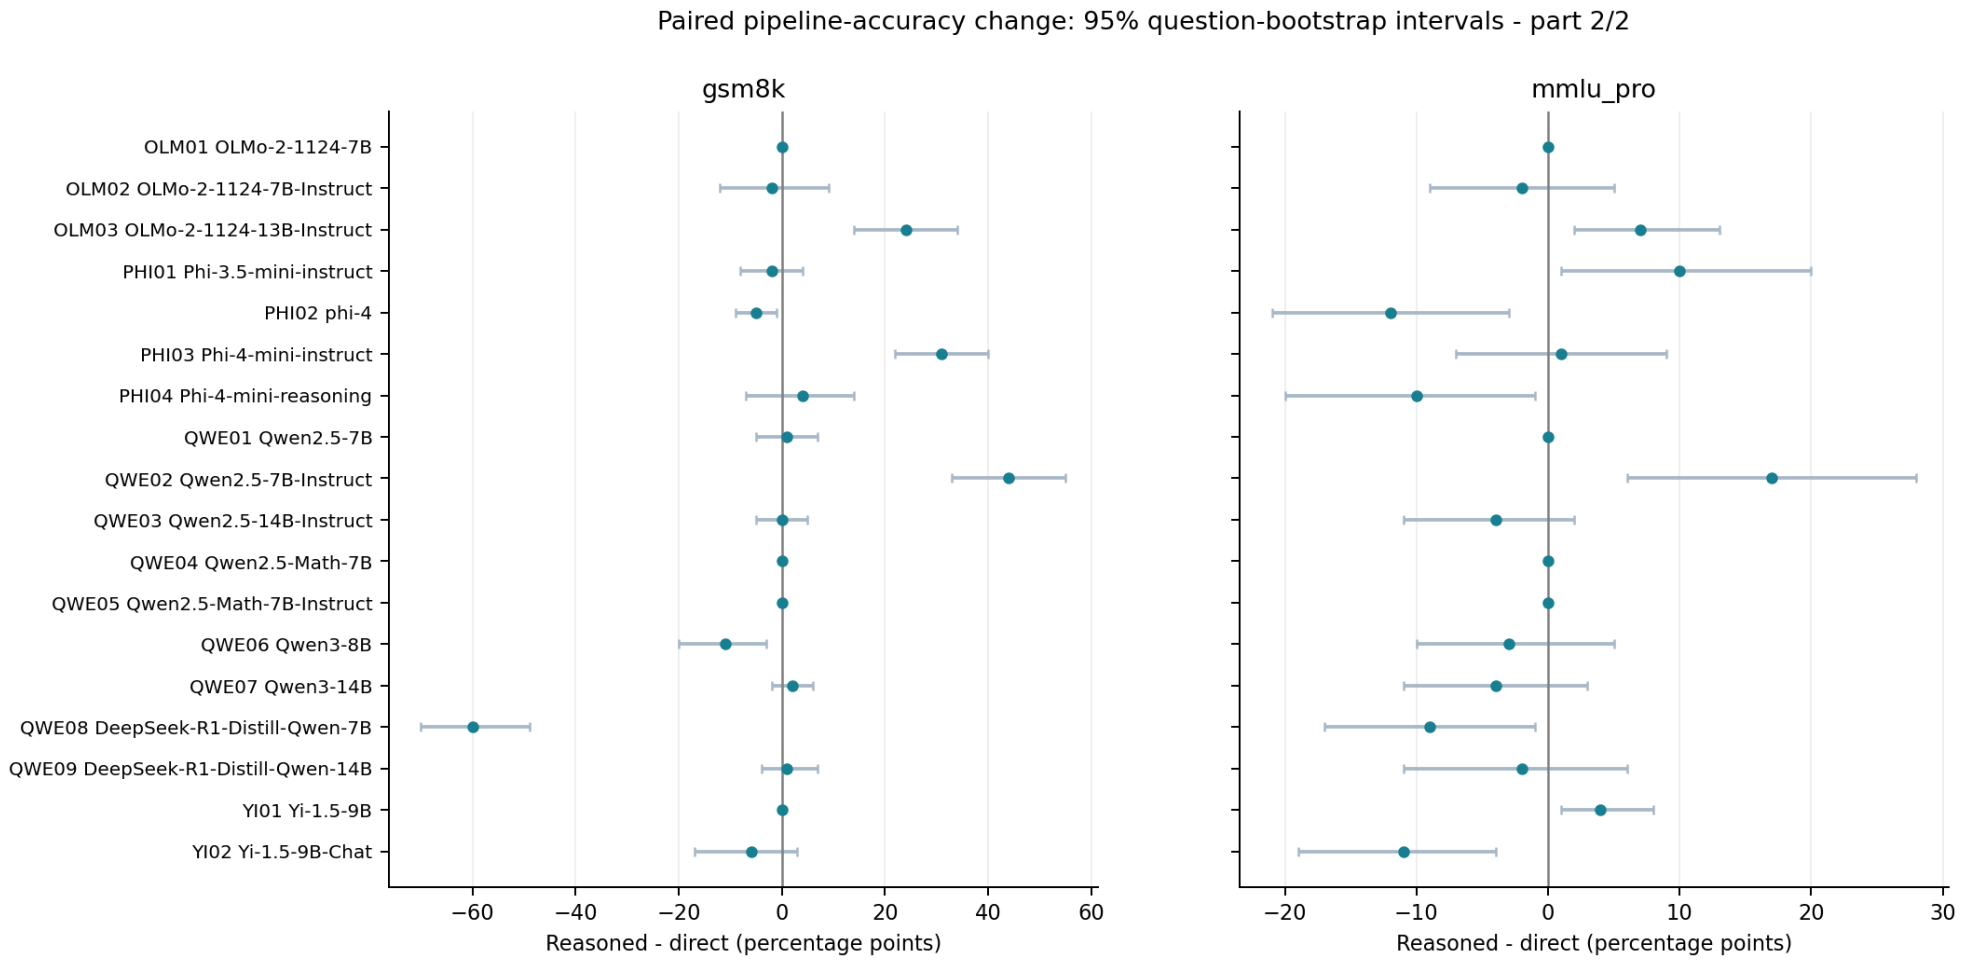


Figure 4. Paired prompt effects with 2,000 question-bootstrap draws per model/benchmark. Intervals are exploratory, unadjusted for multiple testing, and reflect the sampled questions only.

Reasoned prompting has heterogeneous effects and can change both correctness and formatting compliance. Report a paired difference with coverage for both conditions; do not infer that reasoning universally helps or hurts. The intervals do not capture training randomness or uncertainty across new benchmark distributions.

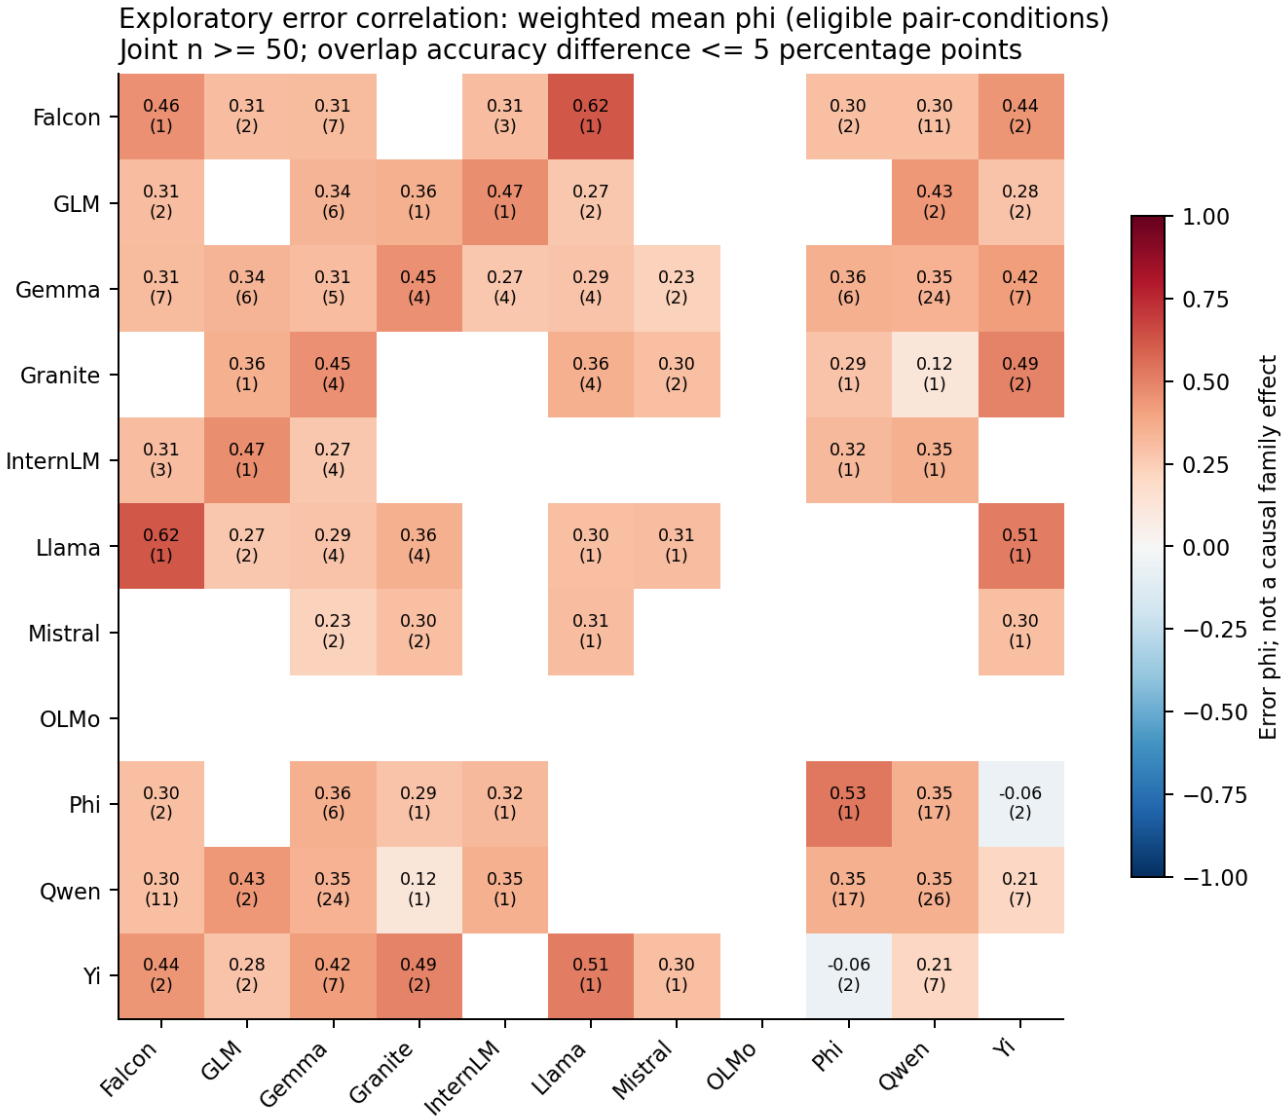


Figure 5. Newly recomputed pooled comparison from scored.json, not a pre-existing saved result. Pair-conditions are screened on joint coverage and overlap accuracy. Blank cells have no eligible pairs. Conditions are pooled descriptively; counts are not independent replications.

Pairwise calculations align identical item IDs and keep benchmark, format, template, seed, and temperature separate before aggregation. Phi is Pearson correlation of binary errors on the jointly usable subset. The figure averages phi using joint n as weights and requires n>=50 plus <=5 percentage-point accuracy difference on that subset. This is only a screening analysis: model selection, difficulty, unequal family composition, missingness, and overlapping questions/pairs remain confounded. No p-value or causal interpretation is assigned to

within-vs-cross-family differences. Numeric-answer wrong-agreement is reported separately from MC agreement; chance-uniform MC metrics do not transfer to numeric answers.

Pooled pair coverage across all 37 checkpoints

Row	benchmark	template	nonempty pairs	joint n ge50	joint n ge90
1	gsm8k	direct	452	117	42
2	gsm8k	reasoned	433	110	29
3	mmlu_pro	direct	369	133	7
4	mmlu_pro	reasoned	396	123	3

Exploratory screened within/cross-family summary

Column panel 1/2 - row numbers identify the same source rows.

Row	benchmark	template	relationship	pair conditions	joint n sum	weighted phi
1	gsm8k	direct	cross family	37	2930	0.279917
2	gsm8k	direct	within family	12	999	0.296846
3	gsm8k	reasoned	cross family	43	3269	0.330005
4	gsm8k	reasoned	within family	12	963	0.257228
5	mmlu_pro	direct	cross family	29	1914	0.349842
6	mmlu_pro	direct	within family	5	409	0.541106
7	mmlu_pro	reasoned	cross family	25	1707	0.398353
8	mmlu_pro	reasoned	within family	5	301	0.563324

Column panel 2/2 - row numbers identify the same source rows.

Row	median joint n
1	81
2	87
3	79
4	79
5	63
6	81
7	69

Row	median joint n
8	52

No cross-family voting teams were generated by the original notebooks. The pooled analysis here recomputes pairs only; it does not fabricate cross-family team voting, hierarchical posterior estimates, or verification outcomes. Complete pooled pair rows and paired prompt deltas are provided as companion CSVs in family_report_data/.

Recorded within-family voting: complete-case comparison

The following table summarizes all saved teams with joint n>=50, grouped by family and condition. Gains compare recorded answer-plurality accuracy with the strongest member on exactly the same team-complete-case questions. Team sizes can differ. Min/max gains describe the observed selected teams and are not confidence intervals. A family/condition without eligible teams is omitted; this is not evidence of zero gain. These are descriptive in-sample voting outcomes, not validated selection policies.

Within-family plurality gain over the same-item best member

Column panel 1/2 - row numbers identify the same source rows.

Row	family	benchmark	template	teams	minimum n	maximum n
1	Gemma	gsm8k	direct	1	92	92
2	Gemma	gsm8k	reasoned	1	90	90
3	Gemma	mmlu_pro	direct	1	75	75
4	Gemma	mmlu_pro	reasoned	1	61	61
5	Qwen	gsm8k	direct	11	72	86
6	Qwen	gsm8k	reasoned	11	63	79
7	Qwen	mmlu_pro	direct	10	54	81
8	Qwen	mmlu_pro	reasoned	7	50	80

Column panel 2/2 - row numbers identify the same source rows.

Row	mean gain pp	minimum gain pp	maximum gain pp
1	-1.44928	-1.44928	-1.44928
2	0.37037	0.37037	0.37037
3	4.88889	4.88889	4.88889
4	0.546448	0.546448	0.546448
5	-0.122194	-1.31579	1.2987
6	-0.41924	-1.5873	1.28205
7	-0.674981	-5.55556	2.92398

Row	mean gain pp	minimum gain pp	maximum gain pp
8	1.28045	-2.38095	4.66667

5. Conclusions usable for paper planning

Supported: 37 checkpoints from 11 family groups completed an aligned 200-question pilot, yielding 14,800 responses. Archive integrity and alignment are confirmed. Evaluation coverage varies sharply across models and families. Some higher-coverage checkpoints permit descriptive accuracy comparisons, and paired question-level records enable later rescoring and controlled shared-error analyses.

Not established: where shared failure comes from; family/lineage variance fractions or buy-back-able error; superiority of cross-family voting; student-teacher inheritance; calibrated predictions for held-out families/lineages; verifier false approval rates; sampled self-consistency gains. These require additional analysis or experiments beyond the archive.

Suggested paper wording: "We collected an aligned pilot of 37 checkpoints across 11 developer groups on 100 GSM8K and 100 MMLU-Pro items under direct and reasoned prompting. Under the original strict extraction and truncation policy, 6,673 of 14,800 responses were usable. We therefore treat current accuracy and dependence summaries as descriptive and prioritize response-extraction validation before drawing lineage-aware conclusions."

6. Complete family and checkpoint results

Each condition table below reproduces every saved accuracy.csv field. Values are proportions unless the header says otherwise. n is 100 per row. Model aliases keep PDF tables compact; exact model IDs and immutable revisions are listed for each family. Output tokens are summed, not token rates. Timing is summed once per model/batch_start from repeated batch_elapsed_s fields: generation calls only, excluding download, engine loading, dataset preparation, and analysis. Do not sum per-response batch times.

6.1 Falcon

Source: research-3-main/Families/Falcon/experiment_1bbbdfd1ab5811bb/. 3 checkpoints; 1200 responses; 535 usable (44.6%); 271 truncated; 398 pipeline-correct. Three checkpoints, including one base model. Base-model formatting and truncation need auditing before comparing the base/tuned relationship. Three-member team rows exist but must be read with their complete-case n.

gsm8k, direct: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is FAL03: 87/100 correct; 94/100 usable. This is a descriptive selection, not a significance-tested winner.

gsm8k, reasoned: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is FAL03: 88/100 correct; 96/100 usable. This is a descriptive selection, not a significance-tested winner.

mmlu_pro, direct: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

mmlu_pro, reasoned: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

Exact checkpoint inventory and settings

Column panel 1/3 - row numbers identify the same source rows.

Row	alias	model	kind	size B	revision	context
1	FAL01	tiiuae/Falcon3-7B-Base	base	7	bf3d7ed586cb22a921520e2d681a9d3d7642cde8	16384
2	FAL02	tiiuae/Falcon3-7B-Instruct	instruct	7	1e57a0ecd176c7c139f289c60a74e57f887c3dfb	16384
3	FAL03	tiiuae/Falcon3-10B-Instruct	instruct	10	8799bc6aec0152757221dc6b272d824642db6202	16384

Column panel 2/3 - row numbers identify the same source rows.

Row	max output	dtype	quantization	base	teacher	lineage
1	2048	auto	NA	NA	NA	NA
2	2048	auto	NA	NA	NA	NA
3	2048	auto	NA	NA	NA	NA

Column panel 3/3 - row numbers identify the same source rows.

Row	lineage verified	trust remote code	chat kwargs	engine kwargs	hub license	license reviewed
1	False	False	{}	{}	other	False
2	False	False	{}	{}	other	False
3	False	False	{}	{}	other	False

Every model-condition accuracy and coverage result

Column panel 1/4 - row numbers identify the same source rows.

Row	alias	model	benchmark	format	template	seed
1	FAL03	FAL03	gsm8k	numeric	direct	17
2	FAL03	FAL03	gsm8k	numeric	reasoned	17
3	FAL03	FAL03	mmlu_pro	mc	direct	17
4	FAL03	FAL03	mmlu_pro	mc	reasoned	17
5	FAL01	FAL01	gsm8k	numeric	direct	17
6	FAL01	FAL01	gsm8k	numeric	reasoned	17
7	FAL01	FAL01	mmlu_pro	mc	direct	17
8	FAL01	FAL01	mmlu_pro	mc	reasoned	17
9	FAL02	FAL02	gsm8k	numeric	direct	17
10	FAL02	FAL02	gsm8k	numeric	reasoned	17
11	FAL02	FAL02	mmlu_pro	mc	direct	17
12	FAL02	FAL02	mmlu_pro	mc	reasoned	17

Column panel 2/4 - row numbers identify the same source rows.

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
1	0	100	94	0.87	0.925532	0.06
2	0	100	96	0.88	0.916667	0.04
3	0	100	76	0.42	0.552632	0.24
4	0	100	87	0.52	0.597701	0.13
5	0	100	13	0.11	0.65	0.8
6	0	100	5	0.02	0.555556	0.73
7	0	100	1	0	0.666667	0.97
8	0	100	2	0	0	0.93
9	0	100	67	0.6	0.895522	0.33
10	0	100	23	0.21	0.913043	0.77
11	0	100	25	0.11	0.44	0.75

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
12	0	100	46	0.24	0.521739	0.54

Column panel 3/4 - row numbers identify the same source rows.

Row	usable accuracy	truncation rate	output tokens	pipeline wilson low	pipeline wilson high	usable wilson low
1	0.925532	0.01	24210	0.790196	0.922428	0.854188
2	0.916667	0.01	24584	0.801879	0.930006	0.844101
3	0.552632	0.02	23056	0.327984	0.517935	0.441005
4	0.597701	0	30200	0.423166	0.615354	0.492647
5	0.846154	0.44	131120	0.062542	0.186313	0.577654
6	0.4	0.68	171421	0.00550197	0.0700118	0.117621
7	0	0.77	169719	3.46945e-18	0.0369935	0
8	0	0.77	178457	3.46945e-18	0.0369935	0
9	0.895522	0	18229	0.502003	0.690599	0.799687
10	0.913043	0	13057	0.141657	0.2998	0.73204
11	0.44	0	13445	0.062542	0.186313	0.266656
12	0.521739	0.01	34211	0.166913	0.332323	0.381374

Column panel 4/4 - row numbers identify the same source rows.

Row	usable wilson high
1	0.963461
2	0.957169
3	0.659193
4	0.694492
5	0.956742
6	0.769276
7	0.793451
8	0.65762
9	0.948462
10	0.97582
11	0.629327

Row	usable wilson high
12	0.658753

Answer extraction status counts

Row	alias	model	parse status	responses
1	FAL03	FAL03	invalid_numeric	1
2	FAL03	FAL03	invalid_option	19
3	FAL03	FAL03	missing_final	27
4	FAL03	FAL03	ok	353
5	FAL01	FAL01	invalid_numeric	81
6	FAL01	FAL01	invalid_option	9
7	FAL01	FAL01	missing_final	253
8	FAL01	FAL01	ok	57
9	FAL02	FAL02	invalid_numeric	5
10	FAL02	FAL02	invalid_option	74
11	FAL02	FAL02	missing_final	160
12	FAL02	FAL02	ok	161

Measured token usage and generation-call time

Column panel 1/2 - row numbers identify the same source rows.

Row	alias	model	input tokens	output tokens	mean output tokens	generation seconds
1	FAL03	FAL03	65992	102050	255.125	600.79
2	FAL01	FAL01	61592	650717	1626.79	1425.06
3	FAL02	FAL02	65992	78942	197.355	369.358

Column panel 2/2 - row numbers identify the same source rows.

Row	generation tokens per s	completed first	completed last
1	169.86	2026-10-19T11:47:01Z	2026-10-19T11:56:32Z
2	456.623	2026-10-19T11:15:33Z	2026-10-19T11:37:57Z

Row	generation tokens per s	completed first	completed last
3	213.728	2026-10-19T11:39:19Z	2026-10-19T11:45:16Z

Analysis outputs present

Row	output	rows
1	accuracy.csv	12
2	pairs.csv	12
3	teams.csv	4
4	ladder.csv	14
5	repeat_pairs.csv	5
6	self_consistency.csv	0
7	baseline_predictions.csv	6

Saved design audit: Fewer than 20 models.; Fewer than six families.; Unknown or unaudited lineage entries; do not replace these with one shared unknown group.; Family/lineage covariance design is rank deficient.

6.2 GLM

Source: research-3-main/Families/GLM/experiment_2f5b62b0c3d89ef7/. 2 checkpoints; 800 responses; 289 usable (36.1%); 26 truncated; 202 pipeline-correct. Two checkpoints: no requested 3- or 5-member within-family team can be formed. Pair results use only available overlap; low coverage limits base/chat interpretation.

gsm8k, direct: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

gsm8k, reasoned: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

mmlu_pro, direct: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

mmlu_pro, reasoned: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

Exact checkpoint inventory and settings

Column panel 1/3 - row numbers identify the same source rows.

Row	alias	model	kind	size B	revision	context
1	GLM01	THUDM/glm-4-9b	base	9	8cd2b585357ba9e702647ac4e6fa4fafe5cc7bee	8192
2	GLM02	THUDM/glm-4-9b-chat	instruct	9	bd8234fe5e0c09c48637a92abb0c797cb5fa0e73	16384

Column panel 2/3 - row numbers identify the same source rows.

Row	max output	dtype	quantization	base	teacher	lineage
1	2048	auto	NA	NA	NA	NA
2	2048	auto	NA	NA	NA	NA

Column panel 3/3 - row numbers identify the same source rows.

Row	lineage verified	trust remote code	chat kwargs	engine kwargs	hub license	license reviewed
1	False	True	{}	{}	other	False
2	False	True	{}	{}	other	False

Every model-condition accuracy and coverage result

Column panel 1/4 - row numbers identify the same source rows.

Row	alias	model	benchmark	format	template	seed
1	GLM01	GLM01	gsm8k	numeric	direct	17
2	GLM01	GLM01	gsm8k	numeric	reasoned	17
3	GLM01	GLM01	mmlu_pro	mc	direct	17
4	GLM01	GLM01	mmlu_pro	mc	reasoned	17
5	GLM02	GLM02	gsm8k	numeric	direct	17
6	GLM02	GLM02	gsm8k	numeric	reasoned	17
7	GLM02	GLM02	mmlu_pro	mc	direct	17
8	GLM02	GLM02	mmlu_pro	mc	reasoned	17

Column panel 2/4 - row numbers identify the same source rows.

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
1	0	100	3	0.02	0.666667	0.97
2	0	100	1	0.01	1	0.99
3	0	100	0	0	NA	1
4	0	100	2	0.02	1	0.98
5	0	100	59	0.49	0.830508	0.41
6	0	100	82	0.69	0.841463	0.18
7	0	100	71	0.38	0.535211	0.29
8	0	100	71	0.41	0.577465	0.29

Column panel 3/4 - row numbers identify the same source rows.

Row	usable accuracy	truncation rate	output tokens	pipeline wilson low	pipeline wilson high	usable wilson low
1	0.666667	0	5423	0.00550197	0.0700118	0.20766
2	1	0	11151	0.00176743	0.0544862	0.206549
3	NA	0.04	38026	3.46945e-18	0.0369935	NA
4	1	0.09	53631	0.00550197	0.0700118	0.34238
5	0.830508	0	17858	0.39422	0.58652	0.715368
6	0.841463	0	19874	0.593739	0.772203	0.747421
7	0.535211	0.07	31763	0.290976	0.477902	0.420392
8	0.577465	0.06	48390	0.318673	0.507986	0.461507

Column panel 4/4 - row numbers identify the same source rows.

Row	usable wilson high
1	0.938508
2	1
3	NA
4	1
5	0.905242
6	0.904944
7	0.646416
8	0.68547

Answer extraction status counts

Row	alias	model	parse status	responses
1	GLM01	GLM01	invalid_numeric	2
2	GLM01	GLM01	invalid_option	1
3	GLM01	GLM01	missing_final	391
4	GLM01	GLM01	ok	6
5	GLM02	GLM02	invalid_numeric	59
6	GLM02	GLM02	invalid_option	5
7	GLM02	GLM02	missing_final	53
8	GLM02	GLM02	ok	283

Measured token usage and generation-call time

Column panel 1/2 - row numbers identify the same source rows.

Row	alias	model	input tokens	output tokens	mean output tokens	generation seconds
1	GLM01	GLM01	57416	108231	270.577	626.644
2	GLM02	GLM02	59416	117885	294.712	582.189

Column panel 2/2 - row numbers identify the same source rows.

Row	generation tokens per s	completed first	completed last
1	172.715	2026-10-20T15:57:06Z	2026-10-20T16:06:03Z
2	202.486	2026-10-20T16:07:54Z	2026-10-20T16:17:03Z

Analysis outputs present

Row	output	rows
1	accuracy.csv	8
2	pairs.csv	3
3	teams.csv	0
4	ladder.csv	5
5	repeat_pairs.csv	2
6	self_consistency.csv	0
7	baseline_predictions.csv	0

Saved design audit: Fewer than 20 models.; Fewer than six families.; Unknown or unaudited lineage entries; do not replace these with one shared unknown group.; Family/lineage covariance design is rank deficient.

6.3 Gemma

Source: research-3-main/Families/Gemma/experiment_c08c743d00ac8744/. 3 checkpoints; 1200 responses; 1101 usable (91.8%); 15 truncated; 817 pipeline-correct. Highest aggregate usable coverage in this archive (91.8%). All three are instruction-tuned; the group spans Gemma 2 and 3. This is a relatively useful pilot subset, but it does not isolate a single-generation family effect.

gsm8k, direct: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is GEM03: 97/100 correct; 100/100 usable. This is a descriptive selection, not a significance-tested winner.

gsm8k, reasoned: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is GEM03: 97/100 correct; 100/100 usable. This is a descriptive selection, not a significance-tested winner.

mmlu_pro, direct: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is GEM03: 54/100 correct; 99/100 usable. This is a descriptive selection, not a significance-tested winner.

mmlu_pro, reasoned: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is GEM03: 67/100 correct; 98/100 usable. This is a descriptive selection, not a significance-tested winner.

Exact checkpoint inventory and settings

Column panel 1/3 - row numbers identify the same source rows.

Row	alias	model	kind	size B	revision	context
1	GEM01	google/gemma-2-9b-it	instruct	9	11c9b309abf73637e4b6f9a3fa1e92e615547819	8192
2	GEM02	google/gemma-3-4b-it	instruct	4	093f9f388b31de276ce2de164bdc2081324b9767	16384
3	GEM03	google/gemma-3-12b-it	instruct	12	96b6f1eccf38110c56df3a15bffe176da04bfd80	16384

Column panel 2/3 - row numbers identify the same source rows.

Row	max output	dtype	quantization	base	teacher	lineage
1	2048	auto	NA	NA	NA	NA
2	2048	auto	NA	NA	NA	NA
3	2048	auto	NA	NA	NA	NA

Column panel 3/3 - row numbers identify the same source rows.

Row	lineage verified	trust remote code	chat kwargs	engine kwargs	hub license	license reviewed
1	False	False	{}	{}	gemma	False

Row	lineage verified	trust remote code	chat kwargs	engine kwargs	hub license	license reviewed
2	False	False	{}	{}	gemma	False
3	False	False	{}	{}	gemma	False

Every model-condition accuracy and coverage result

Column panel 1/4 - row numbers identify the same source rows.

Row	alias	model	benchmark	format	template	seed
1	GEM01	GEM01	gsm8k	numeric	direct	17
2	GEM01	GEM01	gsm8k	numeric	reasoned	17
3	GEM01	GEM01	mmlu_pro	mc	direct	17
4	GEM01	GEM01	mmlu_pro	mc	reasoned	17
5	GEM03	GEM03	gsm8k	numeric	direct	17
6	GEM03	GEM03	gsm8k	numeric	reasoned	17
7	GEM03	GEM03	mmlu_pro	mc	direct	17
8	GEM03	GEM03	mmlu_pro	mc	reasoned	17
9	GEM02	GEM02	gsm8k	numeric	direct	17
10	GEM02	GEM02	gsm8k	numeric	reasoned	17
11	GEM02	GEM02	mmlu_pro	mc	direct	17
12	GEM02	GEM02	mmlu_pro	mc	reasoned	17

Column panel 2/4 - row numbers identify the same source rows.

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
1	0	100	95	0.85	0.894737	0.05
2	0	100	95	0.84	0.884211	0.05
3	0	100	79	0.41	0.518987	0.21
4	0	100	65	0.41	0.630769	0.35
5	0	100	100	0.97	0.97	0
6	0	100	100	0.97	0.97	0
7	0	100	99	0.54	0.545455	0.01
8	0	100	98	0.67	0.683673	0.02

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
9	0	100	97	0.84	0.865979	0.03
10	0	100	93	0.84	0.903226	0.07
11	0	100	90	0.42	0.466667	0.1
12	0	100	90	0.41	0.455556	0.1

Column panel 3/4 - row numbers identify the same source rows.

Row	usable accuracy	truncation rate	output tokens	pipeline wilson low	pipeline wilson high	usable wilson low
1	0.894737	0	11968	0.767164	0.90694	0.816979
2	0.884211	0	12716	0.755797	0.899047	0.804455
3	0.518987	0	13254	0.318673	0.507986	0.410511
4	0.630769	0	15257	0.318673	0.507986	0.509238
5	0.97	0	20085	0.915481	0.989745	0.915481
6	0.97	0	20603	0.915481	0.989745	0.915481
7	0.545455	0.01	44211	0.442649	0.634392	0.447507
8	0.683673	0.02	51555	0.573053	0.754369	0.586161
9	0.865979	0.02	32119	0.755797	0.899047	0.7841
10	0.903226	0.03	38154	0.755797	0.899047	0.826213
11	0.466667	0.03	46589	0.327984	0.517935	0.367084
12	0.455556	0.04	54581	0.318673	0.507986	0.356596

Column panel 4/4 - row numbers identify the same source rows.

Row	usable wilson high
1	0.941811
2	0.934101
3	0.625703
4	0.737706
5	0.989745
6	0.989745
7	0.640006
8	0.767329

Row	usable wilson high
9	0.919975
10	0.948248
11	0.568978
12	0.558154

Answer extraction status counts

Row	alias	model	parse status	responses
1	GEM01	GEM01	invalid_numeric	8
2	GEM01	GEM01	invalid_option	12
3	GEM01	GEM01	missing_final	46
4	GEM01	GEM01	ok	334
5	GEM03	GEM03	missing_final	3
6	GEM03	GEM03	ok	397
7	GEM02	GEM02	missing_final	30
8	GEM02	GEM02	ok	370

Measured token usage and generation-call time

Column panel 1/2 - row numbers identify the same source rows.

Row	alias	model	input tokens	output tokens	mean output tokens	generation seconds
1	GEM01	GEM01	63778	53195	132.988	404.984
2	GEM03	GEM03	63814	136454	341.135	1690.63
3	GEM02	GEM02	63814	171443	428.608	1786.7

Column panel 2/2 - row numbers identify the same source rows.

Row	generation tokens per s	completed first	completed last
1	131.351	2026-10-18T11:19:38Z	2026-10-18T11:25:19Z
2	80.7121	2026-10-18T11:59:36Z	2026-10-18T12:27:15Z
3	95.9551	2026-10-18T11:28:20Z	2026-10-18T11:56:59Z

Analysis outputs present

Row	output	rows
1	accuracy.csv	12
2	pairs.csv	12
3	teams.csv	4
4	ladder.csv	16
5	repeat_pairs.csv	6
6	self_consistency.csv	0
7	baseline_predictions.csv	8

Saved design audit: Fewer than 20 models.; Fewer than six families.; Unknown or unaudited lineage entries; do not replace these with one shared unknown group.; Family/lineage covariance design is rank deficient.

6.4 Granite

Source: research-3-main/Families/Granite/experiment_54a2a5a5b75dd874/. 2 checkpoints; 800 responses; 438 usable (54.8%); 96 truncated; 210 pipeline-correct. Base and instruct pair only. Requested team sizes cannot be formed. Strict answer extraction and truncation need inspection before interpreting tuning effects.

gsm8k, direct: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

gsm8k, reasoned: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

mmlu_pro, direct: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

mmlu_pro, reasoned: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

Exact checkpoint inventory and settings

Column panel 1/3 - row numbers identify the same source rows.

Row	alias	model	kind	size B	revision	context
1	GRA01	ibm-granite/granite-3.3-8b-base	base	8	cfb7adb44a974653cbb2ff883653971c54dba578	16384
2	GRA02	ibm-granite/granite-3.3-8b-instruct	instruct	8	51dd4bc2ade4059a6bd87649d68aa11e4fb2529b	16384

Column panel 2/3 - row numbers identify the same source rows.

Row	max output	dtype	quantization	base	teacher	lineage
1	2048	auto	NA	NA	NA	NA
2	2048	auto	NA	NA	NA	NA

Column panel 3/3 - row numbers identify the same source rows.

Row	lineage verified	trust remote code	chat kwargs	engine kwargs	hub license	license reviewed
1	False	False	{}	{}	apache-2.0	False
2	False	False	{}	{}	apache-2.0	False

Every model-condition accuracy and coverage result

Column panel 1/4 - row numbers identify the same source rows.

Row	alias	model	benchmark	format	template	seed
1	GRA01	GRA01	gsm8k	numeric	direct	17
2	GRA01	GRA01	gsm8k	numeric	reasoned	17
3	GRA01	GRA01	mmlu_pro	mc	direct	17
4	GRA01	GRA01	mmlu_pro	mc	reasoned	17
5	GRA02	GRA02	gsm8k	numeric	direct	17
6	GRA02	GRA02	gsm8k	numeric	reasoned	17
7	GRA02	GRA02	mmlu_pro	mc	direct	17
8	GRA02	GRA02	mmlu_pro	mc	reasoned	17

Column panel 2/4 - row numbers identify the same source rows.

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
1	0	100	36	0.16	0.375	0.52
2	0	100	28	0.18	0.545455	0.67
3	0	100	76	0.22	0.260417	0.04
4	0	100	71	0.22	0.315789	0.05
5	0	100	12	0.09	0.75	0.88
6	0	100	51	0.45	0.882353	0.49
7	0	100	80	0.36	0.45	0.2
8	0	100	84	0.42	0.5	0.16

Column panel 3/4 - row numbers identify the same source rows.

Row	usable accuracy	truncation rate	output tokens	pipeline wilson low	pipeline wilson high	usable wilson low
1	0.444444	0.31	72187	0.100953	0.244203	0.295413
2	0.642857	0.18	48666	0.117002	0.266674	0.458303
3	0.289474	0.21	47840	0.150013	0.310704	0.19961
4	0.309859	0.24	54584	0.150013	0.310704	0.214397
5	0.75	0	24550	0.0480725	0.162262	0.467695
6	0.882353	0	23875	0.356145	0.547554	0.766192
7	0.45	0	34892	0.272712	0.457646	0.345777
8	0.5	0.02	50930	0.327984	0.517935	0.395439

Column panel 4/4 - row numbers identify the same source rows.

Row	usable wilson high
1	0.604189
2	0.792942
3	0.399596
4	0.42484
5	0.911058
6	0.944949
7	0.558805
8	0.604561

Answer extraction status counts

Row	alias	model	parse status	responses
1	GRA01	GRA01	invalid_numeric	113
2	GRA01	GRA01	invalid_option	7
3	GRA01	GRA01	missing_final	8
4	GRA01	GRA01	ok	272
5	GRA02	GRA02	invalid_numeric	115
6	GRA02	GRA02	invalid_option	19
7	GRA02	GRA02	missing_final	39
8	GRA02	GRA02	ok	227

Measured token usage and generation-call time

Column panel 1/2 - row numbers identify the same source rows.

Row	alias	model	input tokens	output tokens	mean output tokens	generation seconds
1	GRA01	GRA01	66426	223277	558.192	1373.9
2	GRA02	GRA02	90028	134247	335.618	553.089

Column panel 2/2 - row numbers identify the same source rows.

Row	generation tokens per s	completed first	completed last
1	162.513	2026-10-19T11:10:50Z	2026-10-19T11:32:16Z
2	242.722	2026-10-19T11:33:48Z	2026-10-19T11:42:41Z

Analysis outputs present

Row	output	rows
1	accuracy.csv	8
2	pairs.csv	4
3	teams.csv	0
4	ladder.csv	4
5	repeat_pairs.csv	4
6	self_consistency.csv	0
7	baseline_predictions.csv	0

Saved design audit: Fewer than 20 models.; Fewer than six families.; Unknown or unaudited lineage entries; do not replace these with one shared unknown group.; Family/lineage covariance design is rank deficient.

6.5 InternLM

Source: research-3-main/Families/InternLM/experiment_b563ce3b896d28b8/. 2 checkpoints; 800 responses; 199 usable (24.9%); 1 truncated; 143 pipeline-correct. Base/chat pair has no saved pair rows despite completed inference: no joint usable questions under the original scorer. Do not interpret an empty pairs.csv as zero error correlation.

gsm8k, direct: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

gsm8k, reasoned: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

mmlu_pro, direct: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

mmlu_pro, reasoned: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

Exact checkpoint inventory and settings

Column panel 1/3 - row numbers identify the same source rows.

Row	alias	model	kind	size B	revision	context
1	INT01	internlm/internlm2_5-7b	base	7	0481d88c24f938d97226eb8556cf1de89ba60772	16384
2	INT02	internlm/internlm2_5-7b-chat	instruct	7	eb72b541689b0432dade0435feb339b89fdd39ff	16384

Column panel 2/3 - row numbers identify the same source rows.

Row	max output	dtype	quantization	base	teacher	lineage
1	2048	auto	NA	NA	NA	NA
2	2048	auto	NA	NA	NA	NA

Column panel 3/3 - row numbers identify the same source rows.

Row	lineage verified	trust remote code	chat kwargs	engine kwargs	hub license	license reviewed
1	False	True	{}	{"trust_remote_code":true}	NA	False
2	False	True	{}	{"trust_remote_code":true}	other	False

Every model-condition accuracy and coverage result

Column panel 1/4 - row numbers identify the same source rows.

Row	alias	model	benchmark	format	template	seed
1	INT01	INT01	gsm8k	numeric	direct	17
2	INT01	INT01	gsm8k	numeric	reasoned	17
3	INT01	INT01	mmlu_pro	mc	direct	17
4	INT01	INT01	mmlu_pro	mc	reasoned	17
5	INT02	INT02	gsm8k	numeric	direct	17
6	INT02	INT02	gsm8k	numeric	reasoned	17
7	INT02	INT02	mmlu_pro	mc	direct	17
8	INT02	INT02	mmlu_pro	mc	reasoned	17

Column panel 2/4 - row numbers identify the same source rows.

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
1	0	100	0	0	NA	1
2	0	100	0	0	NA	1
3	0	100	0	0	NA	1
4	0	100	0	0	NA	1
5	0	100	73	0.64	0.876712	0.27
6	0	100	66	0.59	0.893939	0.34
7	0	100	36	0.13	0.361111	0.64
8	0	100	24	0.07	0.291667	0.76

Column panel 3/4 - row numbers identify the same source rows.

Row	usable accuracy	truncation rate	output tokens	pipeline wilson low	pipeline wilson high	usable wilson low
1	NA	0	303	3.46945e-18	0.0369935	NA
2	NA	0	100	3.46945e-18	0.0369935	NA
3	NA	0	100	3.46945e-18	0.0369935	NA
4	NA	0	100	3.46945e-18	0.0369935	NA
5	0.876712	0	16146	0.542354	0.727288	0.781997
6	0.893939	0	14353	0.492014	0.681327	0.796877
7	0.361111	0.01	9059	0.0775717	0.209804	0.224756
8	0.291667	0	8886	0.0343193	0.137495	0.149146

Column panel 4/4 - row numbers identify the same source rows.

Row	usable wilson high
1	NA
2	NA
3	NA
4	NA
5	0.933762
6	0.947667
7	0.524249
8	0.491677

Answer extraction status counts

Row	alias	model	parse status	responses
1	INT01	INT01	missing_final	400
2	INT02	INT02	invalid_numeric	34
3	INT02	INT02	invalid_option	7
4	INT02	INT02	missing_final	160
5	INT02	INT02	ok	199

Measured token usage and generation-call time

Column panel 1/2 - row numbers identify the same source rows.

Row	alias	model	input tokens	output tokens	mean output tokens	generation seconds
1	INT01	INT01	58470	603	1.5075	86.4345
2	INT02	INT02	62470	48444	121.11	222.693

Column panel 2/2 - row numbers identify the same source rows.

Row	generation tokens per s	completed first	completed last
1	6.97638	2026-10-20T16:20:08Z	2026-10-20T16:20:38Z
2	217.537	2026-10-20T16:22:22Z	2026-10-20T16:25:52Z

Analysis outputs present

Row	output	rows
1	accuracy.csv	8
2	pairs.csv	0
3	teams.csv	0
4	ladder.csv	0
5	repeat_pairs.csv	2
6	self_consistency.csv	0
7	baseline_predictions.csv	0

Saved design audit: Fewer than 20 models.; Fewer than six families.; Unknown or unaudited lineage entries; do not replace these with one shared unknown group.; Family/lineage covariance design is rank deficient.

6.6 Llama

Source: research-3-main/Families/Llama/experiment_ba84666668bc63e1/. 4 checkpoints; 1600 responses; 628 usable (39.2%); 107 truncated; 409 pipeline-correct. Includes a DeepSeek-R1-distilled Llama checkpoint. Recorded lineage spans Qwen and Llama through distills, but the teacher is missing. No saved team rows passed complete-case requirements.

gsm8k, direct: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is LLA02: 81/100 correct; 96/100 usable. This is a descriptive selection, not a significance-tested winner.

gsm8k, reasoned: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is LLA02: 76/100 correct; 93/100 usable. This is a descriptive selection, not a significance-tested winner.

mmlu_pro, direct: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

mmlu_pro, reasoned: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

Exact checkpoint inventory and settings

Column panel 1/3 - row numbers identify the same source rows.

Row	alias	model	kind	size B	revision	context
1	LLA01	meta-llama/Llama-3.1-8B	base	8	d04e592bb4f6aa9cfee91e2e20afa771667e1d4b	16384
2	LLA02	meta-llama/Llama-3.1-8B-Instruct	instruct	8	0e9e39f249a16976918f6564b8830bc894c89659	16384
3	LLA03	meta-llama/Llama-3.2-3B-Instruct	instruct	3	0cb88a4f764b7a12671c53f0838cd831a0843b95	16384
4	LLA04	deepseek-ai/DeepSeek-R1-Distill-Llama-8B	reasoning	8	6a6f4aa4197940add57724a7707d069478df56b1	16384

Column panel 2/3 - row numbers identify the same source rows.

Row	max output	dtype	quantization	base	teacher	lineage
1	2048	auto	NA	NA	NA	NA
2	2048	auto	NA	LLA01	NA	NA
3	2048	auto	NA	NA	NA	NA
4	8192	auto	NA	LLA01	deepseek-ai/DeepSeek-R1	DeepSeek-R1

Column panel 3/3 - row numbers identify the same source rows.

Row	lineage verified	trust remote code	chat kwargs	engine kwargs	hub license	license reviewed
1	False	False	{}	{}	llama3.1	False
2	False	False	{}	{}	llama3.1	False
3	False	False	{}	{}	llama3.2	False
4	True	False	{}	{}	mit	False

Every model-condition accuracy and coverage result

Column panel 1/4 - row numbers identify the same source rows.

Row	alias	model	benchmark	format	template	seed
1	LLA04	LLA04	gsm8k	numeric	direct	17
2	LLA04	LLA04	gsm8k	numeric	reasoned	17
3	LLA04	LLA04	mmlu_pro	mc	direct	17
4	LLA04	LLA04	mmlu_pro	mc	reasoned	17
5	LLA01	LLA01	gsm8k	numeric	direct	17
6	LLA01	LLA01	gsm8k	numeric	reasoned	17
7	LLA01	LLA01	mmlu_pro	mc	direct	17
8	LLA01	LLA01	mmlu_pro	mc	reasoned	17
9	LLA02	LLA02	gsm8k	numeric	direct	17
10	LLA02	LLA02	gsm8k	numeric	reasoned	17
11	LLA02	LLA02	mmlu_pro	mc	direct	17
12	LLA02	LLA02	mmlu_pro	mc	reasoned	17
13	LLA03	LLA03	gsm8k	numeric	direct	17
14	LLA03	LLA03	gsm8k	numeric	reasoned	17
15	LLA03	LLA03	mmlu_pro	mc	direct	17
16	LLA03	LLA03	mmlu_pro	mc	reasoned	17

Column panel 2/4 - row numbers identify the same source rows.

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
1	0	100	0	0	NA	1
2	0	100	0	0	NA	1

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
3	0	100	0	0	NA	1
4	0	100	0	0	NA	1
5	0	100	0	0	NA	1
6	0	100	0	0	NA	1
7	0	100	0	0	NA	1
8	0	100	0	0	NA	1
9	0	100	96	0.81	0.84375	0.04
10	0	100	93	0.76	0.817204	0.07
11	0	100	69	0.34	0.492754	0.31
12	0	100	82	0.43	0.52439	0.18
13	0	100	96	0.64	0.666667	0.04
14	0	100	96	0.75	0.78125	0.04
15	0	100	38	0.1	0.263158	0.62
16	0	100	58	0.26	0.448276	0.42

Column panel 3/4 - row numbers identify the same source rows.

Row	usable accuracy	truncation rate	output tokens	pipeline wilson low	pipeline wilson high	usable wilson low
1	NA	0.01	50148	3.46945e-18	0.0369935	NA
2	NA	0.01	51095	3.46945e-18	0.0369935	NA
3	NA	0.23	287830	3.46945e-18	0.0369935	NA
4	NA	0.23	283605	3.46945e-18	0.0369935	NA
5	NA	0.03	6241	3.46945e-18	0.0369935	NA
6	NA	0.04	8349	3.46945e-18	0.0369935	NA
7	NA	0	101	3.46945e-18	0.0369935	NA
8	NA	0.03	6354	3.46945e-18	0.0369935	NA
9	0.84375	0.01	17285	0.722212	0.874852	0.758085
10	0.817204	0.01	19507	0.667677	0.833087	0.726622
11	0.492754	0.18	48512	0.254615	0.437223	0.378324
12	0.52439	0.12	55183	0.337333	0.527846	0.417647
13	0.666667	0	13434	0.542354	0.727288	0.567565

Row	usable accuracy	truncation rate	output tokens	pipeline wilson low	pipeline wilson high	usable wilson low
14	0.78125	0	18356	0.656955	0.824548	0.688621
15	0.263158	0.03	7563	0.0552291	0.174366	0.149716
16	0.448276	0.14	49877	0.184047	0.35371	0.327499

Column panel 4/4 - row numbers identify the same source rows.

Row	usable wilson high
1	NA
2	NA
3	NA
4	NA
5	NA
6	NA
7	NA
8	NA
9	0.902963
10	0.882621
11	0.607947
12	0.62895
13	0.752943
14	0.852237
15	0.420088
16	0.575479

Answer extraction status counts

Row	alias	model	parse status	responses
1	LLA04	LLA04	missing_final	400
2	LLA01	LLA01	missing_final	400
3	LLA02	LLA02	invalid_numeric	5
4	LLA02	LLA02	missing_final	55

Row	alias	model	parse status	responses
5	LLA02	LLA02	ok	340
6	LLA03	LLA03	invalid_numeric	1
7	LLA03	LLA03	invalid_option	4
8	LLA03	LLA03	missing_final	107
9	LLA03	LLA03	ok	288

Measured token usage and generation-call time

Column panel 1/2 - row numbers identify the same source rows.

Row	alias	model	input tokens	output tokens	mean output tokens	generation seconds
1	LLA04	LLA04	59358	672678	1681.69	2733.06
2	LLA01	LLA01	56798	21045	52.6125	306.964
3	LLA02	LLA02	70798	140487	351.217	671.173
4	LLA03	LLA03	70798	89230	223.075	363.056

Column panel 2/2 - row numbers identify the same source rows.

Row	generation tokens per s	completed first	completed last
1	246.126	2026-10-18T11:20:18Z	2026-10-18T12:04:54Z
2	68.5586	2026-10-18T10:54:48Z	2026-10-18T10:58:58Z
3	209.315	2026-10-18T11:00:34Z	2026-10-18T11:11:17Z
4	245.775	2026-10-18T11:12:30Z	2026-10-18T11:18:14Z

Analysis outputs present

Row	output	rows
1	accuracy.csv	16
2	pairs.csv	4
3	teams.csv	0
4	ladder.csv	10
5	repeat_pairs.csv	4

Row	output	rows
6	self_consistency.csv	0
7	baseline_predictions.csv	0

Saved design audit: Fewer than 20 models.; Fewer than six families.; Unknown or unaudited lineage entries; do not replace these with one shared unknown group.; Some lineages do not span two families.; Family/lineage covariance design is rank deficient.

6.7 Mistral

Source: research-3-main/Families/Mistral/experiment_91f6579fc6f5ebd7/. 3 checkpoints; 1200 responses; 269 usable (22.4%); 350 truncated; 117 pipeline-correct. Original parser rejects answers such as FINAL: F (15.00 cm) and FINAL: 44 quilt blocks. Base-model truncation is severe. All saved MMLU pair rows have joint n=10; large correlations are not reliable family evidence.

gsm8k, direct: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

gsm8k, reasoned: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

mmlu_pro, direct: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is MIS03: 40/100 correct; 99/100 usable. This is a descriptive selection, not a significance-tested winner.

mmlu_pro, reasoned: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

Exact checkpoint inventory and settings

Column panel 1/3 - row numbers identify the same source rows.

Row	alias	model	kind	size B	revision	context
1	MIS01	mistralai/Mistral-7B-v0.3	base	7	caa1feb0e54d415e2df31207e5f4e273e33509b1	16384
2	MIS02	mistralai/Mistral-7B-Instruct-v0.3	instruct	7	c170c708c41dac9275d15a8fff4eca08d52bab71	16384
3	MIS03	mistralai/Mistral-Nemo-Instruct-2407	instruct	12	04d8a90549d23fc6bd7f642064003592df51e9b3	16384

Column panel 2/3 - row numbers identify the same source rows.

Row	max output	dtype	quantization	base	teacher	lineage
1	2048	auto	NA	NA	NA	NA
2	2048	auto	NA	MIS01	NA	NA
3	2048	auto	NA	NA	NA	NA

Column panel 3/3 - row numbers identify the same source rows.

Row	lineage verified	trust remote code	chat kwargs	engine kwargs	hub license	license reviewed
1	False	False	{}	{}	apache-2.0	False
2	False	False	{}	{}	apache-2.0	False
3	False	False	{}	{}	apache-2.0	False

Every model-condition accuracy and coverage result

Column panel 1/4 - row numbers identify the same source rows.

Row	alias	model	benchmark	format	template	seed
1	MIS02	MIS02	gsm8k	numeric	direct	17
2	MIS02	MIS02	gsm8k	numeric	reasoned	17
3	MIS02	MIS02	mmlu_pro	mc	direct	17
4	MIS02	MIS02	mmlu_pro	mc	reasoned	17
5	MIS01	MIS01	gsm8k	numeric	direct	17
6	MIS01	MIS01	gsm8k	numeric	reasoned	17
7	MIS01	MIS01	mmlu_pro	mc	direct	17
8	MIS01	MIS01	mmlu_pro	mc	reasoned	17
9	MIS03	MIS03	gsm8k	numeric	direct	17
10	MIS03	MIS03	gsm8k	numeric	reasoned	17
11	MIS03	MIS03	mmlu_pro	mc	direct	17
12	MIS03	MIS03	mmlu_pro	mc	reasoned	17

Column panel 2/4 - row numbers identify the same source rows.

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
1	0	100	1	0	0	0.99
2	0	100	7	0.03	0.428571	0.93
3	0	100	10	0.05	0.5	0.9
4	0	100	12	0.05	0.416667	0.88
5	0	100	2	0	0	0.95
6	0	100	4	0.01	0.111111	0.91
7	0	100	0	0	0	0.95
8	0	100	0	0	0.285714	0.93
9	0	100	28	0.07	0.25	0.72
10	0	100	17	0.15	0.882353	0.83
11	0	100	99	0.4	0.40404	0.01
12	0	100	89	0.41	0.460674	0.11

Column panel 3/4 - row numbers identify the same source rows.

Row	usable accuracy	truncation rate	output tokens	pipeline wilson low	pipeline wilson high	usable wilson low
1	0	0	4793	3.46945e-18	0.0369935	0
2	0.428571	0.01	13129	0.0102545	0.0845194	0.15822
3	0.5	0	6359	0.0215437	0.11175	0.236593
4	0.416667	0.02	19664	0.0215437	0.11175	0.19326
5	0	0.94	194652	3.46945e-18	0.0369935	0
6	0.25	0.9	187697	0.00176743	0.0544862	0.0455873
7	NA	0.83	172511	3.46945e-18	0.0369935	NA
8	NA	0.79	165387	3.46945e-18	0.0369935	NA
9	0.25	0	2593	0.0343193	0.137495	0.126765
10	0.882353	0	10973	0.0930599	0.232836	0.656636
11	0.40404	0	1155	0.309401	0.497997	0.312719
12	0.460674	0.01	7311	0.318673	0.507986	0.360897

Column panel 4/4 - row numbers identify the same source rows.

Row	usable wilson high
1	0.793451
2	0.749542
3	0.763407
4	0.680489
5	0.65762
6	0.699358
7	NA
8	NA
9	0.433557
10	0.96712
11	0.502531
12	0.563705

Answer extraction status counts

Row	alias	model	parse status	responses
1	MIS02	MIS02	invalid_numeric	190
2	MIS02	MIS02	invalid_option	175
3	MIS02	MIS02	missing_final	5
4	MIS02	MIS02	ok	30
5	MIS01	MIS01	invalid_numeric	1
6	MIS01	MIS01	invalid_option	1
7	MIS01	MIS01	missing_final	372
8	MIS01	MIS01	ok	26
9	MIS03	MIS03	invalid_numeric	119
10	MIS03	MIS03	invalid_option	2
11	MIS03	MIS03	missing_final	46
12	MIS03	MIS03	ok	233

Measured token usage and generation-call time

Column panel 1/2 - row numbers identify the same source rows.

Row	alias	model	input tokens	output tokens	mean output tokens	generation seconds
1	MIS02	MIS02	66196	43945	109.862	298.507
2	MIS01	MIS01	64996	720247	1800.62	1393.14
3	MIS03	MIS03	60314	22032	55.08	203.172

Column panel 2/2 - row numbers identify the same source rows.

Row	generation tokens per s	completed first	completed last
1	147.216	2026-10-18T08:10:02Z	2026-10-18T08:14:34Z
2	516.995	2026-10-18T07:46:14Z	2026-10-18T08:08:59Z
3	108.44	2026-10-18T08:16:01Z	2026-10-18T08:19:19Z

Analysis outputs present

Row	output	rows
1	accuracy.csv	12
2	pairs.csv	6
3	teams.csv	0
4	ladder.csv	14
5	repeat_pairs.csv	4
6	self_consistency.csv	0
7	baseline_predictions.csv	0

Saved design audit: Fewer than 20 models.; Fewer than six families.; Unknown or unaudited lineage entries; do not replace these with one shared unknown group.; Family/lineage covariance design is rank deficient.

6.8 OLMo

Source: research-3-main/Families/OLMo/experiment_6d9605f600a5f1d9/. 3 checkpoints; 1200 responses; 222 usable (18.5%); 28 truncated; 117 pipeline-correct. Actual checkpoints are OLMo-2, not OLMo-3. Aggregate usability is only 18.5%; this family is especially sensitive to extraction rules. Do not label these results as OLMo-3 in a paper.

gsm8k, direct: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

gsm8k, reasoned: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

mmlu_pro, direct: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

mmlu_pro, reasoned: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

Exact checkpoint inventory and settings

Column panel 1/3 - row numbers identify the same source rows.

Row	alias	model	kind	size B	revision	context
1	OLM01	allenai/OLMo-2-1124-7B	base	7	7df9a82518afdecae4e8c026b27adccc8c1f0032	4096
2	OLM02	allenai/OLMo-2-1124-7B-Instruct	instruct	7	470b1fba1ae01581f270116362ee4aa1b97f4c84	4096
3	OLM03	allenai/OLMo-2-1124-13B-Instruct	instruct	13	3a5c85baefbb1896a54d56fe2e76c0395627ddf4	4096

Column panel 2/3 - row numbers identify the same source rows.

Row	max output	dtype	quantization	base	teacher	lineage
1	2048	auto	NA	NA	NA	NA
2	2048	auto	NA	OLM01	NA	NA
3	2048	auto	NA	NA	NA	NA

Column panel 3/3 - row numbers identify the same source rows.

Row	lineage verified	trust remote code	chat kwargs	engine kwargs	hub license	license reviewed
1	False	False	{}	{}	apache-2.0	False
2	False	False	{}	{}	apache-2.0	False
3	False	False	{}	{}	apache-2.0	False

Every model-condition accuracy and coverage result

Column panel 1/4 - row numbers identify the same source rows.

Row	alias	model	benchmark	format	template	seed
1	OLM03	OLM03	gsm8k	numeric	direct	17
2	OLM03	OLM03	gsm8k	numeric	reasoned	17
3	OLM03	OLM03	mmlu_pro	mc	direct	17
4	OLM03	OLM03	mmlu_pro	mc	reasoned	17
5	OLM01	OLM01	gsm8k	numeric	direct	17
6	OLM01	OLM01	gsm8k	numeric	reasoned	17
7	OLM01	OLM01	mmlu_pro	mc	direct	17
8	OLM01	OLM01	mmlu_pro	mc	reasoned	17
9	OLM02	OLM02	gsm8k	numeric	direct	17
10	OLM02	OLM02	gsm8k	numeric	reasoned	17
11	OLM02	OLM02	mmlu_pro	mc	direct	17
12	OLM02	OLM02	mmlu_pro	mc	reasoned	17

Column panel 2/4 - row numbers identify the same source rows.

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
1	0	100	7	0.06	0.857143	0.93
2	0	100	36	0.3	0.833333	0.64
3	0	100	11	0.02	0.181818	0.89
4	0	100	33	0.09	0.272727	0.67
5	0	100	0	0	NA	1
6	0	100	0	0	NA	1
7	0	100	0	0	0	0.99
8	0	100	0	0	NA	1
9	0	100	25	0.24	0.96	0.75
10	0	100	34	0.22	0.647059	0.66
11	0	100	34	0.13	0.382353	0.66
12	0	100	42	0.11	0.261905	0.58

Column panel 3/4 - row numbers identify the same source rows.

Row	usable accuracy	truncation rate	output tokens	pipeline wilson low	pipeline wilson high	usable wilson low
1	0.857143	0	15202	0.0277861	0.124768	0.486872
2	0.833333	0	14868	0.218949	0.395849	0.681092
3	0.181818	0.02	25423	0.00550197	0.0700118	0.0513677
4	0.272727	0.01	30905	0.0480725	0.162262	0.150674
5	NA	0	100	3.46945e-18	0.0369935	NA
6	NA	0	17998	3.46945e-18	0.0369935	NA
7	NA	0.09	18750	3.46945e-18	0.0369935	NA
8	NA	0.09	23614	3.46945e-18	0.0369935	NA
9	0.96	0	18609	0.166913	0.332323	0.804559
10	0.647059	0	14639	0.150013	0.310704	0.47914
11	0.382353	0.03	27365	0.0775717	0.209804	0.239002
12	0.261905	0.04	33195	0.062542	0.186313	0.153026

Column panel 4/4 - row numbers identify the same source rows.

Row	usable wilson high
1	0.97432
2	0.921296
3	0.476981
4	0.442176
5	NA
6	NA
7	NA
8	NA
9	0.992904
10	0.785121
11	0.54959
12	0.410687

Answer extraction status counts

Row	alias	model	parse status	responses
1	OLM03	OLM03	invalid_numeric	121
2	OLM03	OLM03	invalid_option	88
3	OLM03	OLM03	missing_final	104
4	OLM03	OLM03	ok	87
5	OLM01	OLM01	missing_final	399
6	OLM01	OLM01	ok	1
7	OLM02	OLM02	invalid_numeric	6
8	OLM02	OLM02	invalid_option	10
9	OLM02	OLM02	missing_final	249
10	OLM02	OLM02	ok	135

Measured token usage and generation-call time

Column panel 1/2 - row numbers identify the same source rows.

Row	alias	model	input tokens	output tokens	mean output tokens	generation seconds
1	OLM03	OLM03	60822	86398	215.995	714.39
2	OLM01	OLM01	56422	60462	151.155	649.404
3	OLM02	OLM02	60822	93808	234.52	492.86

Column panel 2/2 - row numbers identify the same source rows.

Row	generation tokens per s	completed first	completed last
1	120.939	2026-10-19T01:41:49Z	2026-10-19T01:53:14Z
2	93.1039	2026-10-19T01:21:03Z	2026-10-19T01:30:26Z
3	190.334	2026-10-19T01:31:51Z	2026-10-19T01:39:44Z

Analysis outputs present

Row	output	rows
1	accuracy.csv	12
2	pairs.csv	4

Row	output	rows
3	teams.csv	0
4	ladder.csv	14
5	repeat_pairs.csv	4
6	self_consistency.csv	0
7	baseline_predictions.csv	0

Saved design audit: Fewer than 20 models.; Fewer than six families.; Unknown or unaudited lineage entries; do not replace these with one shared unknown group.; Family/lineage covariance design is rank deficient.

6.9 Phi

Source: research-3-main/Families/Phi/experiment_ed6e929ca419ed5e/. 4 checkpoints; 1600 responses; 800 usable (50.0%); 51 truncated; 593 pipeline-correct. Actual group includes Phi-3.5-mini, phi-4, Phi-4-mini-instruct, and Phi-4-mini-reasoning. Different training types and budgets confound an aggregate family average. Team/baseline rows exist.

gsm8k, direct: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is PHI02: 94/100 correct; 98/100 usable. This is a descriptive selection, not a significance-tested winner.

gsm8k, reasoned: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is PHI02: 89/100 correct; 96/100 usable. This is a descriptive selection, not a significance-tested winner.

mmlu_pro, direct: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is PHI02: 71/100 correct; 97/100 usable. This is a descriptive selection, not a significance-tested winner.

mmlu_pro, reasoned: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is PHI02: 59/100 correct; 90/100 usable. This is a descriptive selection, not a significance-tested winner.

Exact checkpoint inventory and settings

Column panel 1/3 - row numbers identify the same source rows.

Row	alias	model	kind	size B	revision	context
1	PHI01	microsoft/Phi-3.5-mini-instruct	instruct	3.8	2fe192450127e6a83f7441aef6e3ca586c338b77	16384
2	PHI02	microsoft/phi-4	instruct	14	2db69c1c3e91a05d2c64a3185acfbaf36f744e25	16384
3	PHI03	microsoft/Phi-4-mini-instruct	instruct	3.8	cfbefacb99257ffa30c83adab238a50856ac3083	16384
4	PHI04	microsoft/Phi-4-mini-reasoning	reasoning	3.8	0e3b1e2d02ee478a3743abe3f629e9c0cb722e0a	16384

Column panel 2/3 - row numbers identify the same source rows.

Row	max output	dtype	quantization	base	teacher	lineage
1	2048	auto	NA	NA	NA	NA
2	2048	auto	NA	NA	NA	NA
3	2048	auto	NA	NA	NA	NA
4	8192	auto	NA	NA	NA	NA

Column panel 3/3 - row numbers identify the same source rows.

Row	lineage verified	trust remote code	chat kwargs	engine kwargs	hub license	license reviewed
1	False	False	{}	{}	mit	False
2	False	False	{}	{}	mit	False
3	False	False	{}	{}	mit	False
4	False	False	{}	{}	mit	False

Every model-condition accuracy and coverage result

Column panel 1/4 - row numbers identify the same source rows.

Row	alias	model	benchmark	format	template	seed
1	PHI01	PHI01	gsm8k	numeric	direct	17
2	PHI01	PHI01	gsm8k	numeric	reasoned	17
3	PHI01	PHI01	mmlu_pro	mc	direct	17
4	PHI01	PHI01	mmlu_pro	mc	reasoned	17
5	PHI03	PHI03	gsm8k	numeric	direct	17
6	PHI03	PHI03	gsm8k	numeric	reasoned	17
7	PHI03	PHI03	mmlu_pro	mc	direct	17
8	PHI03	PHI03	mmlu_pro	mc	reasoned	17
9	PHI04	PHI04	gsm8k	numeric	direct	17
10	PHI04	PHI04	gsm8k	numeric	reasoned	17
11	PHI04	PHI04	mmlu_pro	mc	direct	17
12	PHI04	PHI04	mmlu_pro	mc	reasoned	17
13	PHI02	PHI02	gsm8k	numeric	direct	17
14	PHI02	PHI02	gsm8k	numeric	reasoned	17
15	PHI02	PHI02	mmlu_pro	mc	direct	17
16	PHI02	PHI02	mmlu_pro	mc	reasoned	17

Column panel 2/4 - row numbers identify the same source rows.

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
1	0	100	14	0.08	0.571429	0.86
2	0	100	7	0.06	0.857143	0.93

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
3	0	100	55	0.23	0.410714	0.44
4	0	100	56	0.33	0.589286	0.44
5	0	100	8	0.08	1	0.92
6	0	100	42	0.39	0.928571	0.58
7	0	100	51	0.25	0.490196	0.49
8	0	100	48	0.26	0.541667	0.52
9	0	100	26	0.25	0.961538	0.74
10	0	100	33	0.29	0.878788	0.67
11	0	100	46	0.34	0.73913	0.54
12	0	100	33	0.24	0.727273	0.67
13	0	100	98	0.94	0.959184	0.02
14	0	100	96	0.89	0.927083	0.04
15	0	100	97	0.71	0.731959	0.03
16	0	100	90	0.59	0.655556	0.1

Column panel 3/4 - row numbers identify the same source rows.

Row	usable accuracy	truncation rate	output tokens	pipeline wilson low	pipeline wilson high	usable wilson low
1	0.571429	0	10864	0.0410935	0.149981	0.325906
2	0.857143	0	19932	0.0277861	0.124768	0.486872
3	0.418182	0.02	23046	0.158433	0.321544	0.297378
4	0.589286	0.03	38623	0.245631	0.426947	0.458776
5	1	0.01	14511	0.0410935	0.149981	0.675592
6	0.928571	0.01	20042	0.300169	0.48797	0.809906
7	0.490196	0.01	12416	0.175452	0.343045	0.358575
8	0.541667	0.04	22990	0.184047	0.35371	0.402911
9	0.961538	0.01	97523	0.175452	0.343045	0.811072
10	0.878788	0.01	101833	0.210148	0.385389	0.726745
11	0.73913	0.18	344722	0.254615	0.437223	0.597408
12	0.727273	0.18	344351	0.166913	0.332323	0.557824
13	0.959184	0	20353	0.875232	0.972214	0.899712

Row	usable accuracy	truncation rate	output tokens	pipeline wilson low	pipeline wilson high	usable wilson low
14	0.927083	0	21905	0.813687	0.937458	0.85707
15	0.731959	0.01	32701	0.614611	0.789852	0.636221
16	0.655556	0	39969	0.492014	0.681327	0.552835

Column panel 4/4 - row numbers identify the same source rows.

Row	usable wilson high
1	0.786192
2	0.97432
3	0.549669
4	0.708332
5	1
6	0.97541
7	0.623191
8	0.674248
9	0.993178
10	0.951838
11	0.843991
12	0.849326
13	0.984015
14	0.964232
15	0.810024
16	0.745541

Answer extraction status counts

Row	alias	model	parse status	responses
1	PHI01	PHI01	invalid_numeric	166
2	PHI01	PHI01	invalid_option	56
3	PHI01	PHI01	missing_final	45
4	PHI01	PHI01	ok	133

Row	alias	model	parse status	responses
5	PHI03	PHI03	invalid_numeric	115
6	PHI03	PHI03	invalid_option	90
7	PHI03	PHI03	missing_final	46
8	PHI03	PHI03	ok	149
9	PHI04	PHI04	invalid_numeric	49
10	PHI04	PHI04	invalid_option	2
11	PHI04	PHI04	missing_final	211
12	PHI04	PHI04	ok	138
13	PHI02	PHI02	invalid_numeric	6
14	PHI02	PHI02	invalid_option	9
15	PHI02	PHI02	missing_final	4
16	PHI02	PHI02	ok	381

Measured token usage and generation-call time

Column panel 1/2 - row numbers identify the same source rows.

Row	alias	model	input tokens	output tokens	mean output tokens	generation seconds
1	PHI01	PHI01	67556	92465	231.162	471.846
2	PHI03	PHI03	57188	69959	174.898	403.714
3	PHI04	PHI04	63188	888429	2221.07	3632.29
4	PHI02	PHI02	59222	114928	287.32	721.852

Column panel 2/2 - row numbers identify the same source rows.

Row	generation tokens per s	completed first	completed last
1	195.965	2026-10-19T01:10:18Z	2026-10-19T01:17:39Z
2	173.288	2026-10-19T01:33:12Z	2026-10-19T01:39:25Z
3	244.592	2026-10-19T01:42:20Z	2026-10-19T02:40:53Z
4	159.213	2026-10-19T01:21:07Z	2026-10-19T01:31:49Z

Analysis outputs present

Row	output	rows
1	accuracy.csv	16
2	pairs.csv	24
3	teams.csv	15
4	ladder.csv	16
5	repeat_pairs.csv	8
6	self_consistency.csv	0
7	baseline_predictions.csv	24

Saved design audit: Fewer than 20 models.; Fewer than six families.; Unknown or unaudited lineage entries; do not replace these with one shared unknown group.; Family/lineage covariance design is rank deficient.

6.10 Qwen

Source: research-3-main/Families/Qwen/experiment_a812191ca0cf2dbb/. 9 checkpoints; 3600 responses; 1894 usable (52.6%); 234 truncated; 1542 pipeline-correct. Nine checkpoints include base, Math, instruct, Qwen3, and two DeepSeek-R1 distills. This group contributes far more pair/team rows than smaller groups. Do not let its number of combinations dominate a family-level statistic.

gsm8k, direct: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is QWE03: 92/100 correct; 95/100 usable. This is a descriptive selection, not a significance-tested winner.

gsm8k, reasoned: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is QWE03: 92/100 correct; 97/100 usable. This is a descriptive selection, not a significance-tested winner.

mmlu_pro, direct: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is QWE03: 73/100 correct; 100/100 usable. This is a descriptive selection, not a significance-tested winner.

mmlu_pro, reasoned: highest pipeline accuracy among this family's checkpoints meeting the >=90 usable screening rule is QWE03: 69/100 correct; 99/100 usable. This is a descriptive selection, not a significance-tested winner.

Exact checkpoint inventory and settings

Column panel 1/3 - row numbers identify the same source rows.

Row	alias	model	kind	size B	revision	context
1	QWE01	Qwen/Qwen2.5-7B	base	7	d149729398750b98c0af14eb82c78cfe92750796	16384
2	QWE02	Qwen/Qwen2.5-7B-Instruct	instruct	7	a09a35458c702b33eeacc393d103063234e8bc28	16384
3	QWE03	Qwen/Qwen2.5-14B-Instruct	instruct	14	cf98f3b3bbb457ad9e2bb7baf9a0125b6b88caa8	16384
4	QWE04	Qwen/Qwen2.5-Math-7B	base	7	b101308fe89651ea5ce025f25317fea6fc07e96e	4096
5	QWE05	Qwen/Qwen2.5-Math-7B-Instruct	instruct	7	ef9926d75ab1d54532f6a30dd5e760355eb9aa4d	4096
6	QWE06	Qwen/Qwen3-8B	reasoning	8	b968826d9c46dd6066d109eabc6255188de91218	16384
7	QWE07	Qwen/Qwen3-14B	reasoning	14	40c069824f4251a91eefaf281ebe4c544efd3e18	16384
8	QWE08	deepseek-ai/DeepSeek-R1-Distill-Qwen-7B	reasoning	7	916b56a44061fd5cd7d6a8fb632557ed4f724f60	16384
9	QWE09	deepseek-ai/DeepSeek-R1-Distill-Qwen-14B	reasoning	14	1df8507178afcc1bef68cd8c393f61a886323761	16384

Column panel 2/3 - row numbers identify the same source rows.

Row	max output	dtype	quantization	base	teacher	lineage
1	2048	auto	NA	NA	NA	NA
2	2048	auto	NA	QWE01	NA	NA
3	2048	auto	NA	Qwen/Qwen2.5-14B	NA	NA
4	2048	auto	NA	NA	NA	NA
5	2048	auto	NA	QWE04	NA	NA
6	8192	auto	NA	NA	NA	NA
7	8192	auto	NA	NA	NA	NA
8	8192	auto	NA	QWE04	deepseek-ai/DeepSeek-R1	DeepSeek-R1
9	8192	auto	NA	Qwen/Qwen2.5-14B	deepseek-ai/DeepSeek-R1	DeepSeek-R1

Column panel 3/3 - row numbers identify the same source rows.

Row	lineage verified	trust remote code	chat kwargs	engine kwargs	hub license	license reviewed
1	False	False	{}	{}	apache-2.0	False
2	False	False	{}	{}	apache-2.0	False
3	False	False	{}	{}	apache-2.0	False
4	False	False	{}	{}	apache-2.0	False
5	False	False	{}	{}	apache-2.0	False
6	False	False	{}	{}	apache-2.0	False
7	False	False	{}	{}	apache-2.0	False
8	True	False	{}	{}	mit	False
9	True	False	{}	{}	mit	False

Every model-condition accuracy and coverage result

Column panel 1/4 - row numbers identify the same source rows.

Row	alias	model	benchmark	format	template	seed
1	QWE03	QWE03	gsm8k	numeric	direct	17
2	QWE03	QWE03	gsm8k	numeric	reasoned	17
3	QWE03	QWE03	mmlu_pro	mc	direct	17

Row	alias	model	benchmark	format	template	seed
4	QWE03	QWE03	mmlu_pro	mc	reasoned	17
5	QWE01	QWE01	gsm8k	numeric	direct	17
6	QWE01	QWE01	gsm8k	numeric	reasoned	17
7	QWE01	QWE01	mmlu_pro	mc	direct	17
8	QWE01	QWE01	mmlu_pro	mc	reasoned	17
9	QWE02	QWE02	gsm8k	numeric	direct	17
10	QWE02	QWE02	gsm8k	numeric	reasoned	17
11	QWE02	QWE02	mmlu_pro	mc	direct	17
12	QWE02	QWE02	mmlu_pro	mc	reasoned	17
13	QWE04	QWE04	gsm8k	numeric	direct	17
14	QWE04	QWE04	gsm8k	numeric	reasoned	17
15	QWE04	QWE04	mmlu_pro	mc	direct	17
16	QWE04	QWE04	mmlu_pro	mc	reasoned	17
17	QWE05	QWE05	gsm8k	numeric	direct	17
18	QWE05	QWE05	gsm8k	numeric	reasoned	17
19	QWE05	QWE05	mmlu_pro	mc	direct	17
20	QWE05	QWE05	mmlu_pro	mc	reasoned	17
21	QWE07	QWE07	gsm8k	numeric	direct	17
22	QWE07	QWE07	gsm8k	numeric	reasoned	17
23	QWE07	QWE07	mmlu_pro	mc	direct	17
24	QWE07	QWE07	mmlu_pro	mc	reasoned	17
25	QWE06	QWE06	gsm8k	numeric	direct	17
26	QWE06	QWE06	gsm8k	numeric	reasoned	17
27	QWE06	QWE06	mmlu_pro	mc	direct	17
28	QWE06	QWE06	mmlu_pro	mc	reasoned	17
29	QWE09	QWE09	gsm8k	numeric	direct	17
30	QWE09	QWE09	gsm8k	numeric	reasoned	17
31	QWE09	QWE09	mmlu_pro	mc	direct	17
32	QWE09	QWE09	mmlu_pro	mc	reasoned	17
33	QWE08	QWE08	gsm8k	numeric	direct	17

Row	alias	model	benchmark	format	template	seed
34	QWE08	QWE08	gsm8k	numeric	reasoned	17
35	QWE08	QWE08	mmlu_pro	mc	direct	17
36	QWE08	QWE08	mmlu_pro	mc	reasoned	17

Column panel 2/4 - row numbers identify the same source rows.

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
1	0	100	95	0.92	0.968421	0.05
2	0	100	97	0.92	0.948454	0.03
3	0	100	100	0.73	0.73	0
4	0	100	99	0.69	0.69697	0.01
5	0	100	5	0.05	1	0.95
6	0	100	8	0.06	0.75	0.92
7	0	100	0	0	NA	1
8	0	100	0	0	NA	1
9	0	100	35	0.3	0.857143	0.65
10	0	100	79	0.74	0.936709	0.21
11	0	100	83	0.42	0.506024	0.17
12	0	100	96	0.59	0.614583	0.04
13	0	100	1	0	0	0.99
14	0	100	0	0	NA	1
15	0	100	0	0	NA	1
16	0	100	0	0	NA	1
17	0	100	0	0	NA	1
18	0	100	0	0	NA	1
19	0	100	0	0	NA	1
20	0	100	0	0	NA	1
21	0	100	92	0.87	0.945652	0.08
22	0	100	91	0.89	0.978022	0.09
23	0	100	85	0.63	0.741176	0.15
24	0	100	77	0.59	0.766234	0.23

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
25	0	100	92	0.91	0.98913	0.08
26	0	100	82	0.8	0.97561	0.18
27	0	100	59	0.41	0.694915	0.41
28	0	100	51	0.38	0.745098	0.49
29	0	100	81	0.79	0.975309	0.19
30	0	100	83	0.8	0.963855	0.17
31	0	100	89	0.64	0.719101	0.11
32	0	100	83	0.62	0.746988	0.17
33	0	100	93	0.88	0.946237	0.07
34	0	100	32	0.28	0.875	0.68
35	0	100	56	0.3	0.535714	0.44
36	0	100	50	0.21	0.42	0.5

Column panel 3/4 - row numbers identify the same source rows.

Row	usable accuracy	truncation rate	output tokens	pipeline wilson low	pipeline wilson high	usable wilson low
1	0.968421	0	20584	0.850019	0.958907	0.911229
2	0.948454	0	26540	0.850019	0.958907	0.884956
3	0.73	0	27695	0.635679	0.807304	0.635679
4	0.69697	0	47084	0.593739	0.772203	0.600487
5	1	0.01	11261	0.0215437	0.11175	0.565518
6	0.75	0	6673	0.0277861	0.124768	0.409275
7	NA	0.37	81015	3.46945e-18	0.0369935	NA
8	NA	0.09	20005	3.46945e-18	0.0369935	NA
9	0.857143	0	16098	0.218949	0.395849	0.706244
10	0.936709	0	22662	0.64629	0.815953	0.860251
11	0.506024	0	21580	0.327984	0.517935	0.400604
12	0.614583	0	33378	0.492014	0.681327	0.514607
13	0	0	9421	3.46945e-18	0.0369935	0
14	NA	0.04	31736	3.46945e-18	0.0369935	NA
15	NA	0.36	85799	3.46945e-18	0.0369935	NA

Row	usable accuracy	truncation rate	output tokens	pipeline wilson low	pipeline wilson high	usable wilson low
16	NA	0.09	34875	3.46945e-18	0.0369935	NA
17	NA	0.01	32691	3.46945e-18	0.0369935	NA
18	NA	0	31506	3.46945e-18	0.0369935	NA
19	NA	0.04	61140	3.46945e-18	0.0369935	NA
20	NA	0.04	60307	3.46945e-18	0.0369935	NA
21	0.945652	0	66499	0.790196	0.922428	0.879015
22	0.978022	0	74215	0.813687	0.937458	0.923368
23	0.741176	0.09	237996	0.532205	0.718176	0.639077
24	0.766234	0.11	266384	0.492014	0.681327	0.660462
25	0.98913	0	87585	0.837738	0.951927	0.940972
26	0.97561	0.01	108452	0.711171	0.866633	0.915366
27	0.694915	0.18	321436	0.318673	0.507986	0.568537
28	0.745098	0.17	331648	0.290976	0.477902	0.611317
29	0.975309	0	54226	0.7002	0.858343	0.914373
30	0.963855	0.01	74350	0.711171	0.866633	0.899041
31	0.719101	0.1	202456	0.542354	0.727288	0.618166
32	0.746988	0.16	235976	0.522098	0.709024	0.643977
33	0.946237	0	48341	0.801879	0.930006	0.880252
34	0.875	0.01	60666	0.201397	0.37488	0.719317
35	0.535714	0.22	263081	0.218949	0.395849	0.407042
36	0.42	0.23	290718	0.141657	0.2998	0.29375

Column panel 4/4 - row numbers identify the same source rows.

Row	usable wilson high
1	0.989203
2	0.977784
3	0.807304
4	0.778737
5	1
6	0.928521

Row	usable wilson high
7	NA
8	NA
9	0.937398
10	0.972665
11	0.610911
12	0.705742
13	0.793451
14	NA
15	NA
16	NA
17	NA
18	NA
19	NA
20	NA
21	0.976565
22	0.993952
23	0.822419
24	0.846704
25	0.998079
26	0.993286
27	0.797464
28	0.844542
29	0.993202
30	0.987632
31	0.801905
32	0.828148
33	0.976819
34	0.950299
35	0.659801
36	0.557666

Answer extraction status counts

Row	alias	model	parse status	responses
1	QWE03	QWE03	invalid_numeric	8
2	QWE03	QWE03	missing_final	1
3	QWE03	QWE03	ok	391
4	QWE01	QWE01	invalid_numeric	68
5	QWE01	QWE01	invalid_option	28
6	QWE01	QWE01	missing_final	291
7	QWE01	QWE01	ok	13
8	QWE02	QWE02	invalid_numeric	27
9	QWE02	QWE02	missing_final	80
10	QWE02	QWE02	ok	293
11	QWE04	QWE04	invalid_numeric	3
12	QWE04	QWE04	invalid_option	1
13	QWE04	QWE04	missing_final	395
14	QWE04	QWE04	ok	1
15	QWE05	QWE05	missing_final	400
16	QWE07	QWE07	invalid_numeric	15
17	QWE07	QWE07	missing_final	40
18	QWE07	QWE07	ok	345
19	QWE06	QWE06	invalid_numeric	5
20	QWE06	QWE06	missing_final	111
21	QWE06	QWE06	ok	284
22	QWE09	QWE09	invalid_numeric	31
23	QWE09	QWE09	missing_final	33
24	QWE09	QWE09	ok	336
25	QWE08	QWE08	invalid_numeric	6
26	QWE08	QWE08	invalid_option	1
27	QWE08	QWE08	missing_final	162

Row	alias	model	parse status	responses
28	QWE08	QWE08	ok	231

Measured token usage and generation-call time

Column panel 1/2 - row numbers identify the same source rows.

Row	alias	model	input tokens	output tokens	mean output tokens	generation seconds
1	QWE03	QWE03	70586	121903	304.757	712.153
2	QWE01	QWE01	58986	118954	297.385	723.165
3	QWE02	QWE02	70586	93718	234.295	313.402
4	QWE04	QWE04	58986	161831	404.577	813.76
5	QWE05	QWE05	70186	185644	464.11	543.181
6	QWE07	QWE07	62186	645094	1612.73	4832.02
7	QWE06	QWE06	62186	849121	2122.8	3486.51
8	QWE09	QWE09	60986	567008	1417.52	4575.37
9	QWE08	QWE08	60986	662806	1657.02	2708.95

Column panel 2/2 - row numbers identify the same source rows.

Row	generation tokens per s	completed first	completed last
1	171.175	2026-10-18T07:06:14Z	2026-10-18T07:17:50Z
2	164.491	2026-10-18T06:47:05Z	2026-10-18T06:57:47Z
3	299.035	2026-10-18T06:59:04Z	2026-10-18T07:04:11Z
4	198.868	2026-10-18T07:19:08Z	2026-10-18T07:32:37Z
5	341.772	2026-10-18T07:34:01Z	2026-10-18T07:42:55Z
6	133.504	2026-10-18T08:47:38Z	2026-10-18T10:04:56Z
7	243.545	2026-10-18T07:46:32Z	2026-10-18T08:42:34Z
8	123.926	2026-10-18T10:55:41Z	2026-10-18T12:09:25Z
9	244.672	2026-10-18T10:07:57Z	2026-10-18T10:51:17Z

Analysis outputs present

Row	output	rows
1	accuracy.csv	36
2	pairs.csv	78
3	teams.csv	179
4	ladder.csv	31
5	repeat_pairs.csv	13
6	self_consistency.csv	0
7	baseline_predictions.csv	220

Saved design audit: Fewer than 20 models.; Fewer than six families.; Unknown or unaudited lineage entries; do not replace these with one shared unknown group.; Some lineages do not span two families.; Family/lineage covariance design is rank deficient.

6.11 Yi

Source: research-3-main/Families/Yi/experiment_e65f0d9272abef25/. 2 checkpoints; 800 responses; 298 usable (37.2%); 189 truncated; 195 pipeline-correct. Actual checkpoints are Yi-1.5-9B and Yi-1.5-9B-Chat. Two-member group cannot form requested 3/5 teams; base truncation and limited joint coverage constrain pair comparisons.

gsm8k, direct: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

gsm8k, reasoned: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

mmlu_pro, direct: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

mmlu_pro, reasoned: no checkpoint in this family reaches 90 usable responses; do not use a raw accuracy ordering as a clean capability comparison.

Exact checkpoint inventory and settings

Column panel 1/3 - row numbers identify the same source rows.

Row	alias	model	kind	size B	revision	context
1	YI01	01-ai/Yi-1.5-9B	base	9	80d5471b1eae28beae33e06eadbd4b48e74d4ce1	4096
2	YI02	01-ai/Yi-1.5-9B-Chat	instruct	9	1a0fc698cf883c4f5c325f026ca79f0ebd9955a5	4096

Column panel 2/3 - row numbers identify the same source rows.

Row	max output	dtype	quantization	base	teacher	lineage
1	2048	auto	NA	NA	NA	NA
2	2048	auto	NA	NA	NA	NA

Column panel 3/3 - row numbers identify the same source rows.

Row	lineage verified	trust remote code	chat kwargs	engine kwargs	hub license	license reviewed
1	False	False	{}	{}	apache-2.0	False
2	False	False	{}	{}	apache-2.0	False

Every model-condition accuracy and coverage result

Column panel 1/4 - row numbers identify the same source rows.

Row	alias	model	benchmark	format	template	seed
1	YI01	YI01	gsm8k	numeric	direct	17
2	YI01	YI01	gsm8k	numeric	reasoned	17
3	YI01	YI01	mmlu_pro	mc	direct	17
4	YI01	YI01	mmlu_pro	mc	reasoned	17
5	YI02	YI02	gsm8k	numeric	direct	17
6	YI02	YI02	gsm8k	numeric	reasoned	17
7	YI02	YI02	mmlu_pro	mc	direct	17
8	YI02	YI02	mmlu_pro	mc	reasoned	17

Column panel 2/4 - row numbers identify the same source rows.

Row	temperature	n	usable	pipeline accuracy	parsed accuracy	parse failure
1	0	100	8	0	0.1	0.8
2	0	100	1	0	0	0.97
3	0	100	3	0	0.2	0.95
4	0	100	14	0.04	0.4	0.8
5	0	100	72	0.68	0.944444	0.28
6	0	100	69	0.62	0.898551	0.31
7	0	100	73	0.36	0.493151	0.27
8	0	100	58	0.25	0.431034	0.42

Column panel 3/4 - row numbers identify the same source rows.

Row	usable accuracy	truncation rate	output tokens	pipeline wilson low	pipeline wilson high	usable wilson low
1	0	0.45	98340	3.46945e-18	0.0369935	0
2	0	0.33	74660	3.46945e-18	0.0369935	0
3	0	0.48	104592	3.46945e-18	0.0369935	5.55112e-17
4	0.285714	0.63	136595	0.0156633	0.0983707	0.117214
5	0.944444	0	20214	0.583374	0.763309	0.86568
6	0.898551	0	23284	0.522098	0.709024	0.805081
7	0.493151	0	25415	0.272712	0.457646	0.381709
8	0.431034	0	31869	0.175452	0.343045	0.311818

Column panel 4/4 - row numbers identify the same source rows.

Row	usable wilson high
1	0.324408
2	0.793451
3	0.561497
4	0.546491
5	0.978186
6	0.949983
7	0.605278
8	0.558819

Answer extraction status counts

Row	alias	model	parse status	responses
1	YI01	YI01	invalid_numeric	3
2	YI01	YI01	invalid_option	3
3	YI01	YI01	missing_final	346
4	YI01	YI01	ok	48
5	YI02	YI02	invalid_numeric	59
6	YI02	YI02	invalid_option	59
7	YI02	YI02	missing_final	10
8	YI02	YI02	ok	272

Measured token usage and generation-call time

Column panel 1/2 - row numbers identify the same source rows.

Row	alias	model	input tokens	output tokens	mean output tokens	generation seconds
1	YI01	YI01	62190	414187	1035.47	1625.39
2	YI02	YI02	70266	100782	251.955	463.277

Column panel 2/2 - row numbers identify the same source rows.

Row	generation tokens per s	completed first	completed last
1	254.824	2026-10-19T12:13:02Z	2026-10-19T12:39:33Z
2	217.541	2026-10-19T12:41:00Z	2026-10-19T12:48:26Z

Analysis outputs present

Row	output	rows
1	accuracy.csv	8
2	pairs.csv	3
3	teams.csv	0
4	ladder.csv	3
5	repeat_pairs.csv	2
6	self_consistency.csv	0
7	baseline_predictions.csv	0

Saved design audit: Fewer than 20 models.; Fewer than six families.; Unknown or unaudited lineage entries; do not replace these with one shared unknown group.; Family/lineage covariance design is rank deficient.

7. Complete saved dependence, team, and baseline tables

This appendix preserves all rows and all columns of the six non-accuracy result tables for every family. Empty files are explicitly identified. No rows have been discarded for low n here; inspect n/test_items before interpretation. All family-local results were computed on available complete cases, which can vary from row to row. PDF model IDs are replaced by inventory aliases. Markdown uses the full original IDs. The original ladder intervals bootstrap shared questions but do not resolve missingness or small-sample problems.

Baseline notes: independence and single_factor_tetrachoric are fitted on an item training half and evaluated on jointly usable held-out items. beta_mae and majority_mae are absolute errors for that row, not averages over independent teams. Teams share members and questions. copula_rho is a calibrated latent-factor parameter, not observed Pearson phi; effective_voters is a variance heuristic and cannot identify an all-wrong tail. The baseline routine uses a pooled single factor rather than an estimated family/lineage decomposition.

7.1 Falcon - saved analysis tables

pairs.csv

Column panel 1/4 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	a
1	gsm8k	numeric	direct	17	0	FAL03
2	gsm8k	numeric	direct	17	0	FAL03
3	gsm8k	numeric	direct	17	0	FAL01
4	gsm8k	numeric	reasoned	17	0	FAL03
5	gsm8k	numeric	reasoned	17	0	FAL03
6	gsm8k	numeric	reasoned	17	0	FAL01
7	mmlu_pro	mc	direct	17	0	FAL03
8	mmlu_pro	mc	direct	17	0	FAL03
9	mmlu_pro	mc	direct	17	0	FAL01
10	mmlu_pro	mc	reasoned	17	0	FAL03
11	mmlu_pro	mc	reasoned	17	0	FAL03
12	mmlu_pro	mc	reasoned	17	0	FAL01

Column panel 2/4 - row numbers identify the same source rows.

Row	b	relationship	accuracy matched	n	error a	error b
1	FAL01	same_family|different_size	False	12	0.0833333	0.166667
2	FAL02	same_family|different_size	True	63	0.0793651	0.111111
3	FAL02	same_family	True	10	0.2	0.2
4	FAL01	same_family|different_size	False	4	0	0.5
5	FAL02	same_family|different_size	True	23	0.130435	0.0869565
6	FAL02	same_family	False	2	0.5	0
7	FAL01	same_family|different_size	True	1	1	1
8	FAL02	same_family|different_size	False	21	0.428571	0.619048
9	FAL02	same_family	True	1	1	1
10	FAL01	same_family|different_size	False	2	0.5	1

Row	b	relationship	accuracy matched	n	error a	error b
11	FAL02	same_family|different_size	False	42	0.357143	0.428571
12	FAL02	same_family	False	1	1	0

Column panel 3/4 - row numbers identify the same source rows.

Row	double fault	excess cofailure	phi	adjusted phi	kappa	wrong agreement
1	0	-0.0138889	-0.13484	-0.2	-0.125	NA
2	0.047619	0.0388007	0.45675	0.55	0.44898	0
3	0.1	0.06	0.375	0.375	0.375	0
4	0	0	NA	NA	0	NA
5	0.0434783	0.0321361	0.338648	0.425	0.330097	0
6	0	0	NA	NA	0	NA
7	1	0	NA	NA	NA	0
8	0.380952	0.115646	0.481218	0.708333	0.447368	0.375
9	1	0	NA	NA	NA	0
10	0.5	0	NA	NA	0	0
11	0.190476	0.037415	0.157788	0.183333	0.156028	0.25
12	0	0	NA	NA	0	NA

Column panel 4/4 - row numbers identify the same source rows.

Row	wrong chance uniform	both wrong n	capa onehot
1	NA	0	NA
2	NA	3	NA
3	NA	1	NA
4	NA	0	NA
5	NA	1	NA
6	NA	0	NA
7	0.111111	1	-0.125
8	0.111111	8	0.303312
9	0.111111	1	-0.125
10	0.111111	1	-0.0627306

Row	wrong chance uniform	both wrong n	capa onehot
11	0.111111	8	0.108623
12	NA	0	0

ladder.csv

Column panel 1/2 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	relationship
1	gsm8k	numeric	direct	17	0	same_family
2	gsm8k	numeric	direct	17	0	same_family
3	gsm8k	numeric	direct	17	0	different_size
4	gsm8k	numeric	direct	17	0	different_size
5	gsm8k	numeric	reasoned	17	0	same_family
6	gsm8k	numeric	reasoned	17	0	same_family
7	gsm8k	numeric	reasoned	17	0	different_size
8	gsm8k	numeric	reasoned	17	0	different_size
9	mmlu_pro	mc	direct	17	0	same_family
10	mmlu_pro	mc	direct	17	0	same_family
11	mmlu_pro	mc	direct	17	0	different_size
12	mmlu_pro	mc	direct	17	0	different_size
13	mmlu_pro	mc	reasoned	17	0	same_family
14	mmlu_pro	mc	reasoned	17	0	different_size

Column panel 2/2 - row numbers identify the same source rows.

Row	matched only	pairs	mean adjusted phi	low	high
1	False	3	0.241667	-0.0827824	1
2	True	2	0.4625	-0.00294409	1
3	False	2	0.175	-0.169229	1
4	True	1	0.55	0.107946	1
5	False	3	0.425	-0.2	1
6	True	1	0.425	-0.253571	1

Row	matched only	pairs	mean adjusted phi	low	high
7	False	2	0.425	-0.210088	1
8	True	1	0.425	-0.222222	1
9	False	3	0.708333	0.0764286	1
10	True	2	NA	NA	NA
11	False	2	0.708333	0.141558	1
12	True	1	NA	NA	NA
13	False	3	0.183333	-0.196152	0.540977
14	False	2	0.183333	-0.166667	0.549009

teams.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	team
1	gsm8k	numeric	direct	17	0	["FAL03", "FAL01", "FAL02"]
2	gsm8k	numeric	reasoned	17	0	["FAL03", "FAL01", "FAL02"]
3	mmlu_pro	mc	direct	17	0	["FAL03", "FAL01", "FAL02"]
4	mmlu_pro	mc	reasoned	17	0	["FAL03", "FAL01", "FAL02"]

Column panel 2/3 - row numbers identify the same source rows.

Row	n	beta	beta low	beta high	selection ceiling	correctness majority
1	10	0	0	0	1	0.9
2	2	0	0	0	1	1
3	1	1	1	1	0	0
4	1	0	0	0	1	1

Column panel 3/3 - row numbers identify the same source rows.

Row	plurality accuracy	independence beta
1	0.933333	0.004
2	1	0
3	0	1

Row	plurality accuracy	independence beta
4	1	0

repeat_pairs.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	model	benchmark	format	condition a	condition b	n
1	FAL03	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	92
2	FAL03	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	70
3	FAL01	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	1
4	FAL02	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	20
5	FAL02	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	12

Column panel 2/3 - row numbers identify the same source rows.

Row	error a	error b	double fault	excess cofailure	phi	adjusted phi
1	0.076087	0.0869565	0.0543478	0.0477316	0.638908	0.687075
2	0.414286	0.442857	0.357143	0.173673	0.709783	0.752431
3	1	1	1	0	NA	NA
4	0.15	0.1	0.05	0.035	0.326732	0.411765
5	0.666667	0.666667	0.583333	0.138889	0.625	0.625

Column panel 3/3 - row numbers identify the same source rows.

Row	kappa	wrong agreement	wrong chance uniform	both wrong n	capa onehot
1	0.637224	0.6	NA	5	NA
2	0.708576	0.64	0.12254	25	0.583372
3	NA	1	0.111111	1	1
4	0.318182	0	NA	1	NA
5	0.625	0.428571	0.115646	7	0.403576

self_consistency.csv

No rows were recorded in this output.

baseline_predictions.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	condition	team	method	train items	test items	observed beta
1	["gsm8k", "numeric", "direct", 17, 0.0]	["FAL03", "FAL01", "FAL02"]	independence	50	5	0
2	["gsm8k", "numeric", "direct", 17, 0.0]	["FAL03", "FAL01", "FAL02"]	single_factor_tetrachoric	50	5	0
3	["gsm8k", "numeric", "reasoned", 17, 0.0]	["FAL03", "FAL01", "FAL02"]	independence	50	2	0
4	["gsm8k", "numeric", "reasoned", 17, 0.0]	["FAL03", "FAL01", "FAL02"]	single_factor_tetrachoric	50	2	0
5	["mmlu_pro", "mc", "direct", 17, 0.0]	["FAL03", "FAL01", "FAL02"]	independence	50	1	1
6	["mmlu_pro", "mc", "direct", 17, 0.0]	["FAL03", "FAL01", "FAL02"]	single_factor_tetrachoric	50	1	1

Column panel 2/3 - row numbers identify the same source rows.

Row	beta low	beta high	predicted beta	beta mae	observed majority	predicted majority
1	0	0	0.00107143	0.00107143	0.8	0.967143
2	0	0	0.00107143	0.00107143	0.8	0.9705
3	0	0	0.0116043	0.0116043	1	0.793056
4	0	0	0.0116043	0.0116043	1	0.809
5	1	1	NA	NA	0	NA
6	1	1	NA	NA	0	0.177

Column panel 3/3 - row numbers identify the same source rows.

Row	majority mae	copula rho	effective voters	policy
1	0.167143	0	NA	item holdout; capabilities calibrated on training items
2	0.1705	0	NA	item holdout; capabilities calibrated on training items
3	0.206944	0	NA	item holdout; capabilities calibrated on training items
4	0.191	0	NA	item holdout; capabilities calibrated on training items
5	NA	0	NA	item holdout; capabilities calibrated on training items
6	0.177	0	NA	item holdout; capabilities calibrated on training items

7.2 GLM - saved analysis tables

pairs.csv

Column panel 1/4 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	a
1	gsm8k	numeric	direct	17	0	GLM01
2	gsm8k	numeric	reasoned	17	0	GLM01
3	mmlu_pro	mc	reasoned	17	0	GLM01

Column panel 2/4 - row numbers identify the same source rows.

Row	b	relationship	accuracy matched	n	error a	error b
1	GLM02	same_family	False	1	1	0
2	GLM02	same_family	True	1	0	0
3	GLM02	same_family	True	2	0	0

Column panel 3/4 - row numbers identify the same source rows.

Row	double fault	excess cofailure	phi	adjusted phi	kappa	wrong agreement
1	0	0	NA	NA	0	NA
2	0	0	NA	NA	NA	NA
3	0	0	NA	NA	NA	NA

Column panel 4/4 - row numbers identify the same source rows.

Row	wrong chance uniform	both wrong n	capa onehot
1	NA	0	NA
2	NA	0	NA
3	NA	0	NA

ladder.csv

Column panel 1/2 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	relationship
1	gsm8k	numeric	direct	17	0	same_family
2	gsm8k	numeric	reasoned	17	0	same_family
3	gsm8k	numeric	reasoned	17	0	same_family
4	mmlu_pro	mc	reasoned	17	0	same_family
5	mmlu_pro	mc	reasoned	17	0	same_family

Column panel 2/2 - row numbers identify the same source rows.

Row	matched only	pairs	mean adjusted phi	low	high
1	False	1	NA	NA	NA
2	False	1	NA	NA	NA
3	True	1	NA	NA	NA
4	False	1	NA	NA	NA
5	True	1	NA	NA	NA

teams.csv

No rows were recorded in this output.

repeat_pairs.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	model	benchmark	format	condition a	condition b	n
1	GLM02	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	59
2	GLM02	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	59

Column panel 2/3 - row numbers identify the same source rows.

Row	error a	error b	double fault	excess cofailure	phi	adjusted phi
1	0.169492	0.20339	0.0847458	0.0502729	0.33289	0.37234
2	0.474576	0.389831	0.288136	0.103131	0.423467	0.503506

Column panel 3/3 - row numbers identify the same source rows.

Row	kappa	wrong agreement	wrong chance uniform	both wrong n	capa onehot
1	0.330813	0.6	NA	5	NA
2	0.417199	0.411765	0.111111	17	0.302361

self_consistency.csv

No rows were recorded in this output.

baseline_predictions.csv

No rows were recorded in this output.

7.3 Gemma - saved analysis tables

pairs.csv

Column panel 1/4 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	a
1	gsm8k	numeric	direct	17	0	GEM01
2	gsm8k	numeric	direct	17	0	GEM01
3	gsm8k	numeric	direct	17	0	GEM03
4	gsm8k	numeric	reasoned	17	0	GEM01
5	gsm8k	numeric	reasoned	17	0	GEM01
6	gsm8k	numeric	reasoned	17	0	GEM03
7	mmlu_pro	mc	direct	17	0	GEM01
8	mmlu_pro	mc	direct	17	0	GEM01
9	mmlu_pro	mc	direct	17	0	GEM03
10	mmlu_pro	mc	reasoned	17	0	GEM01
11	mmlu_pro	mc	reasoned	17	0	GEM01
12	mmlu_pro	mc	reasoned	17	0	GEM03

Column panel 2/4 - row numbers identify the same source rows.

Row	b	relationship	accuracy matched	n	error a	error b
1	GEM03	same_family|different_size	False	95	0.105263	0.0315789
2	GEM02	same_family|different_size	True	92	0.0869565	0.130435
3	GEM02	same_family|different_size	False	97	0.0206186	0.134021
4	GEM03	same_family|different_size	False	95	0.115789	0.0315789
5	GEM02	same_family|different_size	True	90	0.122222	0.0888889
6	GEM02	same_family|different_size	False	93	0.0322581	0.0967742
7	GEM03	same_family|different_size	True	79	0.481013	0.455696
8	GEM02	same_family|different_size	True	75	0.466667	0.506667
9	GEM02	same_family|different_size	False	89	0.460674	0.52809
10	GEM03	same_family|different_size	True	65	0.369231	0.353846

Row	b	relationship	accuracy matched	n	error a	error b
11	GEM02	same_family|different_size	False	61	0.344262	0.491803
12	GEM02	same_family|different_size	False	89	0.325843	0.539326

Column panel 3/4 - row numbers identify the same source rows.

Row	double fault	excess cofailure	phi	adjusted phi	kappa	wrong agreement
1	0.0210526	0.0177285	0.330336	0.627451	0.27234	0.5
2	0.0326087	0.0212665	0.224105	0.28125	0.218447	0.333333
3	0.0206186	0.0178552	0.368826	1	0.239487	1
4	0.0105263	0.00686981	0.122773	0.246032	0.0981013	1
5	0.0222222	0.011358	0.12185	0.14557	0.119948	0
6	0.0215054	0.0183836	0.351928	0.630952	0.299435	0.5
7	0.341772	0.122577	0.492598	0.518293	0.491961	0.481481
8	0.32	0.0835556	0.334997	0.362934	0.333925	0.25
9	0.325843	0.0825653	0.331812	0.379791	0.328808	0.413793
10	0.230769	0.100118	0.433865	0.448568	0.433624	0.533333
11	0.295082	0.125773	0.529499	0.718894	0.505673	0.388889
12	0.247191	0.0714556	0.305865	0.47603	0.278201	0.409091

Column panel 4/4 - row numbers identify the same source rows.

Row	wrong chance uniform	both wrong n	capa onehot
1	NA	2	NA
2	NA	3	NA
3	NA	2	NA
4	NA	1	NA
5	NA	2	NA
6	NA	2	NA
7	0.128748	27	0.376444
8	0.115079	24	0.189358
9	0.112206	29	0.261974
10	0.130159	15	0.358716

Row	wrong chance uniform	both wrong n	capa onehot
11	0.13933	18	0.339166
12	0.124098	22	0.226039

ladder.csv

Column panel 1/2 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	relationship
1	gsm8k	numeric	direct	17	0	same_family
2	gsm8k	numeric	direct	17	0	same_family
3	gsm8k	numeric	direct	17	0	different_size
4	gsm8k	numeric	direct	17	0	different_size
5	gsm8k	numeric	reasoned	17	0	same_family
6	gsm8k	numeric	reasoned	17	0	same_family
7	gsm8k	numeric	reasoned	17	0	different_size
8	gsm8k	numeric	reasoned	17	0	different_size
9	mmlu_pro	mc	direct	17	0	same_family
10	mmlu_pro	mc	direct	17	0	same_family
11	mmlu_pro	mc	direct	17	0	different_size
12	mmlu_pro	mc	direct	17	0	different_size
13	mmlu_pro	mc	reasoned	17	0	same_family
14	mmlu_pro	mc	reasoned	17	0	same_family
15	mmlu_pro	mc	reasoned	17	0	different_size
16	mmlu_pro	mc	reasoned	17	0	different_size

Column panel 2/2 - row numbers identify the same source rows.

Row	matched only	pairs	mean adjusted phi	low	high
1	False	3	0.636234	0.318309	0.854476
2	True	1	0.28125	-0.104779	0.622887
3	False	3	0.636234	0.159057	0.883227
4	True	1	0.28125	-0.100147	0.766139

Row	matched only	pairs	mean adjusted phi	low	high
5	False	3	0.340851	-0.108593	0.778656
6	True	1	0.14557	-0.154	0.569637
7	False	3	0.340851	-0.120717	0.753838
8	True	1	0.14557	-0.135375	0.530713
9	False	3	0.420339	0.282457	0.591613
10	True	2	0.440614	0.31427	0.66762
11	False	3	0.420339	0.268859	0.612035
12	True	2	0.440614	0.299746	0.655564
13	False	3	0.547831	0.336814	0.769941
14	True	1	0.448568	0.260762	0.718301
15	False	3	0.547831	0.342857	0.737884
16	True	1	0.448568	0.177879	0.775391

teams.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	team
1	gsm8k	numeric	direct	17	0	["GEM01", "GEM03", "GEM02"]
2	gsm8k	numeric	reasoned	17	0	["GEM01", "GEM03", "GEM02"]
3	mmlu_pro	mc	direct	17	0	["GEM01", "GEM03", "GEM02"]
4	mmlu_pro	mc	reasoned	17	0	["GEM01", "GEM03", "GEM02"]

Column panel 2/3 - row numbers identify the same source rows.

Row	n	beta	beta low	beta high	selection ceiling	correctness majority
1	92	0.0108696	0	0.0326087	0.98913	0.956522
2	90	0.0111111	0	0.0336111	0.988889	0.966667
3	75	0.266667	0.16	0.373333	0.733333	0.546667
4	61	0.213115	0.114754	0.312295	0.786885	0.639344

Column panel 3/3 - row numbers identify the same source rows.

Row	plurality accuracy	independence beta
1	0.963768	0.000246569
2	0.97037	0.00036214
3	0.582222	0.110341
4	0.661202	0.0610624

repeat_pairs.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	model	benchmark	format	condition a	condition b	n
1	GEM01	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	93
2	GEM01	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	63
3	GEM03	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	100
4	GEM03	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	97
5	GEM02	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	92
6	GEM02	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	86

Column panel 2/3 - row numbers identify the same source rows.

Row	error a	error b	double fault	excess cofailure	phi	adjusted phi
1	0.0860215	0.107527	0.0645161	0.0552665	0.63626	0.71988
2	0.428571	0.349206	0.349206	0.199546	0.845841	1
3	0.03	0.03	0.02	0.0191	0.656357	0.656357
4	0.453608	0.319588	0.298969	0.154001	0.663363	0.881923
5	0.130435	0.0869565	0.0543478	0.0430057	0.453191	0.56875
6	0.523256	0.534884	0.406977	0.127096	0.510177	0.522222

Column panel 3/3 - row numbers identify the same source rows.

Row	kappa	wrong agreement	wrong chance uniform	both wrong n	capa onehot
1	0.63144	1	NA	6	NA
2	0.834123	0.818182	0.135642	22	0.765407
3	0.656357	1	NA	2	NA

Row	kappa	wrong agreement	wrong chance uniform	both wrong n	capa onehot
4	0.637343	0.551724	0.120963	29	0.493425
5	0.441748	0.8	NA	5	NA
6	0.510038	0.714286	0.114512	35	0.515304

self_consistency.csv

No rows were recorded in this output.

baseline_predictions.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	condition	team	method	train items	test items	observed beta
1	["gsm8k", "numeric", "direct", 17, 0.0]	["GEM01", "GEM03", "GEM02"]	independence	50	46	0
2	["gsm8k", "numeric", "direct", 17, 0.0]	["GEM01", "GEM03", "GEM02"]	single_factor_tetrachoric	50	46	0
3	["gsm8k", "numeric", "reasoned", 17, 0.0]	["GEM01", "GEM03", "GEM02"]	independence	50	45	0.0222222
4	["gsm8k", "numeric", "reasoned", 17, 0.0]	["GEM01", "GEM03", "GEM02"]	single_factor_tetrachoric	50	45	0.0222222
5	["mmlu_pro", "mc", "direct", 17, 0.0]	["GEM01", "GEM03", "GEM02"]	independence	50	42	0.261905
6	["mmlu_pro", "mc", "direct", 17, 0.0]	["GEM01", "GEM03", "GEM02"]	single_factor_tetrachoric	50	42	0.261905
7	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["GEM01", "GEM03", "GEM02"]	independence	50	30	0.166667
8	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["GEM01", "GEM03", "GEM02"]	single_factor_tetrachoric	50	30	0.166667

Column panel 2/3 - row numbers identify the same source rows.

Row	beta low	beta high	predicted beta	beta mae	observed majority	predicted majority
1	0	0	0.000625271	0.000625271	0.956522	0.975614
2	0	0	0.0174912	0.0174912	0.956522	0.9385
3	0	0.0888889	1.10816e-06	0.0222211	0.955556	0.9889
4	0	0.0888889	2.3845e-05	0.0221984	0.955556	0.979
5	0.142262	0.404762	0.126	0.135905	0.5	0.49
6	0.142262	0.404762	0.236835	0.0250697	0.5	0.493

Row	beta low	beta high	predicted beta	beta mae	observed majority	predicted majority
7	0.0666667	0.333333	0.0675177	0.0991489	0.666667	0.620567
8	0.0666667	0.333333	0.198196	0.0315289	0.666667	0.594

Column panel 3/3 - row numbers identify the same source rows.

Row	majority mae	copula rho	effective voters	policy
1	0.0190927	0.598119	NA	item holdout; capabilities calibrated on training items
2	0.0180217	0.598119	NA	item holdout; capabilities calibrated on training items
3	0.033344	0.318064	NA	item holdout; capabilities calibrated on training items
4	0.0234444	0.318064	NA	item holdout; capabilities calibrated on training items
5	0.01	0.478946	NA	item holdout; capabilities calibrated on training items
6	0.007	0.478946	NA	item holdout; capabilities calibrated on training items
7	0.0460993	0.62839	NA	item holdout; capabilities calibrated on training items
8	0.0726667	0.62839	NA	item holdout; capabilities calibrated on training items

7.4 Granite - saved analysis tables

pairs.csv

Column panel 1/4 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	a
1	gsm8k	numeric	direct	17	0	GRA01
2	gsm8k	numeric	reasoned	17	0	GRA01
3	mmlu_pro	mc	direct	17	0	GRA01
4	mmlu_pro	mc	reasoned	17	0	GRA01

Column panel 2/4 - row numbers identify the same source rows.

Row	b	relationship	accuracy matched	n	error a	error b
1	GRA02	same_family	False	7	0.714286	0.142857
2	GRA02	same_family	False	18	0.388889	0.111111
3	GRA02	same_family	False	64	0.6875	0.5625
4	GRA02	same_family	False	63	0.650794	0.507937

Column panel 3/4 - row numbers identify the same source rows.

Row	double fault	excess cofailure	phi	adjusted phi	kappa	wrong agreement
1	0.142857	0.0408163	0.258199	1	0.125	0
2	0.111111	0.0679012	0.443203	1	0.328358	0
3	0.484375	0.0976562	0.424706	0.555556	0.409836	0.387097
4	0.380952	0.0503905	0.211432	0.284091	0.202532	0.291667

Column panel 4/4 - row numbers identify the same source rows.

Row	wrong chance uniform	both wrong n	capa onehot
1	NA	1	NA
2	NA	2	NA
3	0.114183	31	0.291292
4	0.112434	24	0.151919

ladder.csv

Column panel 1/2 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	relationship
1	gsm8k	numeric	direct	17	0	same_family
2	gsm8k	numeric	reasoned	17	0	same_family
3	mmlu_pro	mc	direct	17	0	same_family
4	mmlu_pro	mc	reasoned	17	0	same_family

Column panel 2/2 - row numbers identify the same source rows.

Row	matched only	pairs	mean adjusted phi	low	high
1	False	1	1	1	1
2	False	1	1	1	1
3	False	1	0.555556	0.270195	0.875103
4	False	1	0.284091	-0.0061036	0.588827

teams.csv

No rows were recorded in this output.

repeat_pairs.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	model	benchmark	format	condition a	condition b	n
1	GRA01	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	16
2	GRA01	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	65
3	GRA02	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	9
4	GRA02	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	71

Column panel 2/3 - row numbers identify the same source rows.

Row	error a	error b	double fault	excess cofailure	phi	adjusted phi
1	0.4375	0.4375	0.3125	0.121094	0.492063	0.492063

Row	error a	error b	double fault	excess cofailure	phi	adjusted phi
2	0.692308	0.661538	0.661538	0.20355	0.932034	1
3	0.222222	0.111111	0.111111	0.0864198	0.661438	1
4	0.549296	0.507042	0.43662	0.158104	0.635573	0.69184

Column panel 3/3 - row numbers identify the same source rows.

Row	kappa	wrong agreement	wrong chance uniform	both wrong n	capa onehot
1	0.492063	0.6	NA	5	NA
2	0.92973	0.883721	0.12034	43	0.871456
3	0.608696	1	NA	1	NA
4	0.633294	0.451613	0.120328	31	0.432008

self_consistency.csv

No rows were recorded in this output.

baseline_predictions.csv

No rows were recorded in this output.

7.5 InternLM - saved analysis tables

pairs.csv

No rows were recorded in this output.

ladder.csv

No rows were recorded in this output.

teams.csv

No rows were recorded in this output.

repeat_pairs.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	model	benchmark	format	condition a	condition b	n
1	INT02	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	55
2	INT02	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	19

Column panel 2/3 - row numbers identify the same source rows.

Row	error a	error b	double fault	excess cofailure	phi	adjusted phi
1	0.0909091	0.109091	0.0545455	0.0446281	0.497955	0.55102
2	0.736842	0.684211	0.631579	0.127424	0.622532	0.707692

Column panel 3/3 - row numbers identify the same source rows.

Row	kappa	wrong agreement	wrong chance uniform	both wrong n	capa onehot
1	0.495413	0.666667	NA	3	NA
2	0.61745	0.75	0.113757	12	0.632189

self_consistency.csv

No rows were recorded in this output.

baseline_predictions.csv

No rows were recorded in this output.

7.6 Llama - saved analysis tables

pairs.csv

Column panel 1/4 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	a
1	gsm8k	numeric	direct	17	0	LLA02
2	gsm8k	numeric	reasoned	17	0	LLA02
3	mmlu_pro	mc	direct	17	0	LLA02
4	mmlu_pro	mc	reasoned	17	0	LLA02

Column panel 2/4 - row numbers identify the same source rows.

Row	b	relationship	accuracy matched	n	error a	error b
1	LLA03	same_family|different_size	False	93	0.16129	0.322581
2	LLA03	same_family|different_size	True	91	0.175824	0.186813
3	LLA03	same_family|different_size	False	25	0.6	0.72
4	LLA03	same_family|different_size	False	50	0.42	0.54

Column panel 3/4 - row numbers identify the same source rows.

Row	double fault	excess cofailure	phi	adjusted phi	kappa	wrong agreement
1	0.11828	0.0662504	0.385328	0.606349	0.348858	0.636364
2	0.0769231	0.0440768	0.297072	0.308277	0.296869	0.285714
3	0.48	0.048	0.218218	0.285714	0.210526	0.25
4	0.32	0.0932	0.37888	0.482402	0.368088	0.5625

Column panel 4/4 - row numbers identify the same source rows.

Row	wrong chance uniform	both wrong n	capa onehot
1	NA	11	NA
2	NA	7	NA
3	0.111111	12	0.141806
4	0.115079	16	0.346143

ladder.csv

Column panel 1/2 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	relationship
1	gsm8k	numeric	direct	17	0	same_family
2	gsm8k	numeric	direct	17	0	different_size
3	gsm8k	numeric	reasoned	17	0	same_family
4	gsm8k	numeric	reasoned	17	0	same_family
5	gsm8k	numeric	reasoned	17	0	different_size
6	gsm8k	numeric	reasoned	17	0	different_size
7	mmlu_pro	mc	direct	17	0	same_family
8	mmlu_pro	mc	direct	17	0	different_size
9	mmlu_pro	mc	reasoned	17	0	same_family
10	mmlu_pro	mc	reasoned	17	0	different_size

Column panel 2/2 - row numbers identify the same source rows.

Row	matched only	pairs	mean adjusted phi	low	high
1	False	1	0.606349	0.311147	0.914919
2	False	1	0.606349	0.246847	1
3	False	1	0.308277	0.0290087	0.624456
4	True	1	0.308277	0.00703981	0.632712
5	False	1	0.308277	0.0674105	0.586855
6	True	1	0.308277	0.0649447	0.628615
7	False	1	0.285714	-0.222619	1
8	False	1	0.285714	-0.286745	1
9	False	1	0.482402	0.19495	0.865081
10	False	1	0.482402	0.11354	0.819725

teams.csv

No rows were recorded in this output.

repeat_pairs.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	model	benchmark	format	condition a	condition b	n
1	LLA02	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	92
2	LLA02	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	62
3	LLA03	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	92
4	LLA03	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	16

Column panel 2/3 - row numbers identify the same source rows.

Row	error a	error b	double fault	excess cofailure	phi	adjusted phi
1	0.152174	0.184783	0.119565	0.0914461	0.655956	0.737143
2	0.5	0.419355	0.354839	0.145161	0.588348	0.692308
3	0.336957	0.217391	0.173913	0.100662	0.516312	0.698361
4	0.75	0.6875	0.625	0.109375	0.544949	0.636364

Column panel 3/3 - row numbers identify the same source rows.

Row	kappa	wrong agreement	wrong chance uniform	both wrong n	capa onehot
1	0.651515	0.545455	NA	11	NA
2	0.580645	0.545455	0.122655	22	0.457062
3	0.493627	0.25	NA	16	NA
4	0.538462	0.4	0.111111	10	0.348289

self_consistency.csv

No rows were recorded in this output.

baseline_predictions.csv

No rows were recorded in this output.

7.7 Mistral - saved analysis tables

pairs.csv

Column panel 1/4 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	a
1	gsm8k	numeric	direct	17	0	MIS02
2	gsm8k	numeric	direct	17	0	MIS01
3	gsm8k	numeric	reasoned	17	0	MIS02
4	gsm8k	numeric	reasoned	17	0	MIS01
5	mmlu_pro	mc	direct	17	0	MIS02
6	mmlu_pro	mc	reasoned	17	0	MIS02

Column panel 2/4 - row numbers identify the same source rows.

Row	b	relationship	accuracy matched	n	error a	error b
1	MIS03	same_family|different_size	False	1	1	0
2	MIS03	same_family|different_size	True	1	1	1
3	MIS03	same_family|different_size	False	5	0.4	0
4	MIS03	same_family|different_size	False	1	1	0
5	MIS03	same_family|different_size	True	10	0.5	0.5
6	MIS03	same_family|different_size	True	10	0.5	0.5

Column panel 3/4 - row numbers identify the same source rows.

Row	double fault	excess cofailure	phi	adjusted phi	kappa	wrong agreement
1	0	0	NA	NA	0	NA
2	1	0	NA	NA	NA	0
3	0	0	NA	NA	0	NA
4	0	0	NA	NA	0	NA
5	0.4	0.15	0.6	0.6	0.6	0.25
6	0.5	0.25	1	1	1	0.4

Column panel 4/4 - row numbers identify the same source rows.

Row	wrong chance uniform	both wrong n	capa onehot
1	NA	0	NA
2	NA	1	NA
3	NA	0	NA
4	NA	0	NA
5	0.119048	4	0.304492
6	0.11746	5	0.583155

ladder.csv

Column panel 1/2 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	relationship
1	gsm8k	numeric	direct	17	0	same_family
2	gsm8k	numeric	direct	17	0	same_family
3	gsm8k	numeric	direct	17	0	different_size
4	gsm8k	numeric	direct	17	0	different_size
5	gsm8k	numeric	reasoned	17	0	same_family
6	gsm8k	numeric	reasoned	17	0	different_size
7	mmlu_pro	mc	direct	17	0	same_family
8	mmlu_pro	mc	direct	17	0	same_family
9	mmlu_pro	mc	direct	17	0	different_size
10	mmlu_pro	mc	direct	17	0	different_size
11	mmlu_pro	mc	reasoned	17	0	same_family
12	mmlu_pro	mc	reasoned	17	0	same_family
13	mmlu_pro	mc	reasoned	17	0	different_size
14	mmlu_pro	mc	reasoned	17	0	different_size

Column panel 2/2 - row numbers identify the same source rows.

Row	matched only	pairs	mean adjusted phi	low	high
1	False	2	NA	NA	NA
2	True	1	NA	NA	NA

Row	matched only	pairs	mean adjusted phi	low	high
3	False	2	NA	NA	NA
4	True	1	NA	NA	NA
5	False	2	NA	NA	NA
6	False	2	NA	NA	NA
7	False	1	0.6	-0.0116667	1
8	True	1	0.6	-0.13125	1
9	False	1	0.6	-0.22	1
10	True	1	0.6	0.11	1
11	False	1	1	1	1
12	True	1	1	1	1
13	False	1	1	1	1
14	True	1	1	1	1

teams.csv

No rows were recorded in this output.

repeat_pairs.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	model	benchmark	format	condition a	condition b	n
1	MIS02	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	1
2	MIS02	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	8
3	MIS03	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	15
4	MIS03	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	89

Column panel 2/3 - row numbers identify the same source rows.

Row	error a	error b	double fault	excess cofailure	phi	adjusted phi
1	1	1	1	0	NA	NA
2	0.5	0.5	0.5	0.25	1	1
3	0.8	0.0666667	0.0666667	0.0133333	0.133631	1

Row	error a	error b	double fault	excess cofailure	phi	adjusted phi
4	0.58427	0.539326	0.494382	0.17927	0.729749	0.79955

Column panel 3/3 - row numbers identify the same source rows.

Row	kappa	wrong agreement	wrong chance uniform	both wrong n	capa onehot
1	NA	1	NA	1	NA
2	1	1	0.126984	4	1
3	0.0350877	0	NA	1	NA
4	0.726714	0.886364	0.119408	44	0.751873

self_consistency.csv

No rows were recorded in this output.

baseline_predictions.csv

No rows were recorded in this output.

7.8 OLMo - saved analysis tables

pairs.csv

Column panel 1/4 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	a
1	gsm8k	numeric	direct	17	0	OLM03
2	gsm8k	numeric	reasoned	17	0	OLM03
3	mmlu_pro	mc	direct	17	0	OLM03
4	mmlu_pro	mc	reasoned	17	0	OLM03

Column panel 2/4 - row numbers identify the same source rows.

Row	b	relationship	accuracy matched	n	error a	error b
1	OLM02	same_family|different_size	True	5	0.2	0.2
2	OLM02	same_family|different_size	True	15	0.333333	0.333333
3	OLM02	same_family|different_size	False	5	0.8	0.6
4	OLM02	same_family|different_size	True	21	0.761905	0.761905

Column panel 3/4 - row numbers identify the same source rows.

Row	double fault	excess cofailure	phi	adjusted phi	kappa	wrong agreement
1	0	-0.04	-0.25	-0.25	-0.25	NA
2	0.2	0.0888889	0.4	0.4	0.4	0
3	0.6	0.12	0.612372	1	0.545455	0
4	0.619048	0.0385488	0.2125	0.2125	0.2125	0.230769

Column panel 4/4 - row numbers identify the same source rows.

Row	wrong chance uniform	both wrong n	capa onehot
1	NA	0	NA
2	NA	3	NA
3	0.111111	3	0.0769231
4	0.111111	13	0.130912

ladder.csv

Column panel 1/2 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	relationship
1	gsm8k	numeric	direct	17	0	same_family
2	gsm8k	numeric	direct	17	0	same_family
3	gsm8k	numeric	direct	17	0	different_size
4	gsm8k	numeric	direct	17	0	different_size
5	gsm8k	numeric	reasoned	17	0	same_family
6	gsm8k	numeric	reasoned	17	0	same_family
7	gsm8k	numeric	reasoned	17	0	different_size
8	gsm8k	numeric	reasoned	17	0	different_size
9	mmlu_pro	mc	direct	17	0	same_family
10	mmlu_pro	mc	direct	17	0	different_size
11	mmlu_pro	mc	reasoned	17	0	same_family
12	mmlu_pro	mc	reasoned	17	0	same_family
13	mmlu_pro	mc	reasoned	17	0	different_size
14	mmlu_pro	mc	reasoned	17	0	different_size

Column panel 2/2 - row numbers identify the same source rows.

Row	matched only	pairs	mean adjusted phi	low	high
1	False	1	-0.25	-2	-0.166667
2	True	1	-0.25	-1	-0.142857
3	False	1	-0.25	-1.6375	-0.142857
4	True	1	-0.25	-1.6	-0.166667
5	False	1	0.4	-0.269524	1
6	True	1	0.4	-0.277273	1
7	False	1	0.4	-0.25	1
8	True	1	0.4	-0.25125	1
9	False	1	1	1	1
10	False	1	1	1	1

Row	matched only	pairs	mean adjusted phi	low	high
11	False	1	0.2125	-0.361429	1
12	True	1	0.2125	-0.400588	1
13	False	1	0.2125	-0.375481	1
14	True	1	0.2125	-0.335119	1

teams.csv

No rows were recorded in this output.

repeat_pairs.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	model	benchmark	format	condition a	condition b	n
1	OLM03	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	4
2	OLM03	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	5
3	OLM02	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	12
4	OLM02	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	21

Column panel 2/3 - row numbers identify the same source rows.

Row	error a	error b	double fault	excess cofailure	phi	adjusted phi
1	0.25	0.25	0.25	0.1875	1	1
2	0.8	0.8	0.8	0.16	1	1
3	0	0.333333	0	0	NA	NA
4	0.619048	0.761905	0.619048	0.147392	0.71261	1

Column panel 3/3 - row numbers identify the same source rows.

Row	kappa	wrong agreement	wrong chance uniform	both wrong n	capa onehot
1	1	0	NA	1	NA
2	1	0.25	0.111111	4	0.325
3	-3.33067e-16	NA	NA	0	NA
4	0.673575	0.230769	0.113553	13	0.276097

self_consistency.csv

No rows were recorded in this output.

baseline_predictions.csv

No rows were recorded in this output.

7.9 Phi - saved analysis tables

pairs.csv

Column panel 1/4 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	a
1	gsm8k	numeric	direct	17	0	PHI01
2	gsm8k	numeric	direct	17	0	PHI01
3	gsm8k	numeric	direct	17	0	PHI01
4	gsm8k	numeric	direct	17	0	PHI03
5	gsm8k	numeric	direct	17	0	PHI03
6	gsm8k	numeric	direct	17	0	PHI04
7	gsm8k	numeric	reasoned	17	0	PHI01
8	gsm8k	numeric	reasoned	17	0	PHI01
9	gsm8k	numeric	reasoned	17	0	PHI01
10	gsm8k	numeric	reasoned	17	0	PHI03
11	gsm8k	numeric	reasoned	17	0	PHI03
12	gsm8k	numeric	reasoned	17	0	PHI04
13	mmlu_pro	mc	direct	17	0	PHI01
14	mmlu_pro	mc	direct	17	0	PHI01
15	mmlu_pro	mc	direct	17	0	PHI01
16	mmlu_pro	mc	direct	17	0	PHI03
17	mmlu_pro	mc	direct	17	0	PHI03
18	mmlu_pro	mc	direct	17	0	PHI04
19	mmlu_pro	mc	reasoned	17	0	PHI01
20	mmlu_pro	mc	reasoned	17	0	PHI01
21	mmlu_pro	mc	reasoned	17	0	PHI01
22	mmlu_pro	mc	reasoned	17	0	PHI03
23	mmlu_pro	mc	reasoned	17	0	PHI03
24	mmlu_pro	mc	reasoned	17	0	PHI04

Column panel 2/4 - row numbers identify the same source rows.

Row	b	relationship	accuracy matched	n	error a	error b
1	PHI03	same_family	True	1	0	0
2	PHI04	same_family	False	4	0.75	0
3	PHI02	same_family|different_size	False	14	0.428571	0
4	PHI04	same_family	True	1	0	0
5	PHI02	same_family|different_size	False	8	0	0.125
6	PHI02	same_family|different_size	True	26	0.0384615	0.0384615
7	PHI03	same_family	False	3	0	0.333333
8	PHI04	same_family	False	3	0.333333	0
9	PHI02	same_family|different_size	False	7	0.142857	0
10	PHI04	same_family	False	14	0.142857	0.214286
11	PHI02	same_family|different_size	False	42	0.0714286	0.142857
12	PHI02	same_family|different_size	True	33	0.121212	0.121212
13	PHI03	same_family	False	31	0.580645	0.483871
14	PHI04	same_family	False	26	0.692308	0.307692
15	PHI02	same_family|different_size	False	54	0.574074	0.240741
16	PHI04	same_family	False	24	0.5	0.333333
17	PHI02	same_family|different_size	False	51	0.509804	0.215686
18	PHI02	same_family|different_size	True	45	0.266667	0.311111
19	PHI03	same_family	False	34	0.411765	0.5
20	PHI04	same_family	False	25	0.44	0.32
21	PHI02	same_family|different_size	True	52	0.442308	0.403846
22	PHI04	same_family	False	18	0.5	0.333333
23	PHI02	same_family|different_size	False	45	0.466667	0.333333
24	PHI02	same_family|different_size	True	31	0.258065	0.290323

Column panel 3/4 - row numbers identify the same source rows.

Row	double fault	excess cofailure	phi	adjusted phi	kappa	wrong agreement
1	0	0	NA	NA	NA	NA
2	0	0	NA	NA	0	NA
3	0	0	NA	NA	0	NA

Row	double fault	excess cofailure	phi	adjusted phi	kappa	wrong agreement
4	0	0	NA	NA	NA	NA
5	0	0	NA	NA	0	NA
6	0	-0.00147929	-0.04	-0.04	-0.04	NA
7	0	0	NA	NA	-3.33067e-16	NA
8	0	0	NA	NA	-3.33067e-16	NA
9	0	0	NA	NA	-7.77156e-16	NA
10	0.0714286	0.0408163	0.284268	0.363636	0.275862	0
11	0	-0.0102041	-0.113228	-0.166667	-0.105263	NA
12	0.030303	0.0156107	0.146552	0.146552	0.146552	1
13	0.419355	0.138398	0.561226	0.682051	0.550725	0.384615
14	0.269231	0.056213	0.263889	0.59375	0.195876	0.285714
15	0.185185	0.0469822	0.222235	0.458194	0.174522	0.3
16	0.291667	0.125	0.53033	0.75	0.5	0.428571
17	0.156863	0.046905	0.228127	0.443636	0.185551	0.5
18	0.177778	0.0948148	0.463136	0.516129	0.460432	0.5
19	0.323529	0.117647	0.478091	0.571429	0.470588	0.454545
20	0.28	0.1392	0.601159	0.776786	0.58194	0.428571
21	0.307692	0.129068	0.529629	0.573071	0.527988	0.25
22	0.277778	0.111111	0.471405	0.666667	0.444444	0.8
23	0.244444	0.0888889	0.377964	0.5	0.363636	0.363636
24	0.225806	0.150884	0.759673	0.823864	0.75718	0.714286

Column panel 4/4 - row numbers identify the same source rows.

Row	wrong chance uniform	both wrong n	capa onehot
1	NA	0	NA
2	NA	0	NA
3	NA	0	NA
4	NA	0	NA
5	NA	0	NA
6	NA	0	NA

Row	wrong chance uniform	both wrong n	capa onehot
7	NA	0	NA
8	NA	0	NA
9	NA	0	NA
10	NA	1	NA
11	NA	0	NA
12	NA	1	NA
13	0.117827	13	0.350152
14	0.142857	7	0.140133
15	0.111111	10	0.12853
16	0.111111	7	0.354646
17	0.111111	8	0.184802
18	0.111111	8	0.358288
19	0.113997	11	0.350422
20	0.111111	7	0.401862
21	0.111111	16	0.285392
22	0.111111	5	0.483636
23	0.111111	11	0.25416
24	0.111111	7	0.652563

ladder.csv

Column panel 1/2 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	relationship
1	gsm8k	numeric	direct	17	0	same_family
2	gsm8k	numeric	direct	17	0	same_family
3	gsm8k	numeric	direct	17	0	different_size
4	gsm8k	numeric	direct	17	0	different_size
5	gsm8k	numeric	reasoned	17	0	same_family
6	gsm8k	numeric	reasoned	17	0	same_family
7	gsm8k	numeric	reasoned	17	0	different_size

Row	benchmark	format	template	seed	temperature	relationship
8	gsm8k	numeric	reasoned	17	0	different_size
9	mmlu_pro	mc	direct	17	0	same_family
10	mmlu_pro	mc	direct	17	0	same_family
11	mmlu_pro	mc	direct	17	0	different_size
12	mmlu_pro	mc	direct	17	0	different_size
13	mmlu_pro	mc	reasoned	17	0	same_family
14	mmlu_pro	mc	reasoned	17	0	same_family
15	mmlu_pro	mc	reasoned	17	0	different_size
16	mmlu_pro	mc	reasoned	17	0	different_size

Column panel 2/2 - row numbers identify the same source rows.

Row	matched only	pairs	mean adjusted phi	low	high
1	False	6	-0.04	-0.172727	-0.0322581
2	True	3	-0.04	-0.155921	-0.0322581
3	False	3	-0.04	-0.149573	-0.0305804
4	True	1	-0.04	-0.156071	-0.03125
5	False	6	0.114507	-0.237451	0.587754
6	True	1	0.146552	-0.251852	1
7	False	3	-0.0100575	-0.240246	0.453602
8	True	1	0.146552	-0.24151	1
9	False	6	0.57396	0.340542	0.798349
10	True	1	0.516129	0.18007	0.892568
11	False	3	0.472653	0.155041	0.756594
12	True	1	0.516129	0.176042	0.872366
13	False	6	0.651969	0.441109	0.86525
14	True	2	0.698467	0.506918	0.898164
15	False	3	0.632311	0.447145	0.824432
16	True	2	0.698467	0.492421	0.9063

teams.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	team
1	gsm8k	numeric	direct	17	0	["PHI01", "PHI03", "PHI02"]
2	gsm8k	numeric	direct	17	0	["PHI01", "PHI04", "PHI02"]
3	gsm8k	numeric	direct	17	0	["PHI03", "PHI04", "PHI02"]
4	gsm8k	numeric	reasoned	17	0	["PHI01", "PHI03", "PHI04"]
5	gsm8k	numeric	reasoned	17	0	["PHI01", "PHI03", "PHI02"]
6	gsm8k	numeric	reasoned	17	0	["PHI01", "PHI04", "PHI02"]
7	gsm8k	numeric	reasoned	17	0	["PHI03", "PHI04", "PHI02"]
8	mmlu_pro	mc	direct	17	0	["PHI01", "PHI03", "PHI04"]
9	mmlu_pro	mc	direct	17	0	["PHI01", "PHI03", "PHI02"]
10	mmlu_pro	mc	direct	17	0	["PHI01", "PHI04", "PHI02"]
11	mmlu_pro	mc	direct	17	0	["PHI03", "PHI04", "PHI02"]
12	mmlu_pro	mc	reasoned	17	0	["PHI01", "PHI03", "PHI04"]
13	mmlu_pro	mc	reasoned	17	0	["PHI01", "PHI03", "PHI02"]
14	mmlu_pro	mc	reasoned	17	0	["PHI01", "PHI04", "PHI02"]
15	mmlu_pro	mc	reasoned	17	0	["PHI03", "PHI04", "PHI02"]

Column panel 2/3 - row numbers identify the same source rows.

Row	n	beta	beta low	beta high	selection ceiling	correctness majority
1	1	0	0	0	1	1
2	4	0	0	0	1	1
3	1	0	0	0	1	1
4	1	0	0	0	1	1
5	3	0	0	0	1	1
6	3	0	0	0	1	1
7	14	0	0	0	1	0.928571
8	13	0.307692	0.0769231	0.615385	0.692308	0.384615
9	31	0.193548	0.0959677	0.322581	0.806452	0.548387
10	25	0.16	0.04	0.281	0.84	0.6
11	24	0.208333	0.0416667	0.375	0.791667	0.666667

Row	n	beta	beta low	beta high	selection ceiling	correctness majority
12	16	0.25	0.0625	0.4375	0.75	0.5
13	32	0.3125	0.155469	0.5	0.6875	0.59375
14	24	0.291667	0.125	0.458333	0.708333	0.666667
15	16	0.25	0.0625	0.439063	0.75	0.625

Column panel 3/3 - row numbers identify the same source rows.

Row	plurality accuracy	independence beta
1	1	0
2	1	0
3	1	0
4	1	0
5	1	0
6	1	0
7	0.952381	0.00655977
8	0.410256	0.159308
9	0.602151	0.063442
10	0.64	0.060928
11	0.680556	0.0486111
12	0.541667	0.0922852
13	0.625	0.0888672
14	0.666667	0.0501302
15	0.645833	0.0585938

repeat_pairs.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	model	benchmark	format	condition a	condition b	n
1	PHI01	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	3
2	PHI01	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	39
3	PHI03	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	8

Row	model	benchmark	format	condition a	condition b	n
4	PHI03	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	37
5	PHI04	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	12
6	PHI04	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	25
7	PHI02	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	96
8	PHI02	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	88

Column panel 2/3 - row numbers identify the same source rows.

Row	error a	error b	double fault	excess cofailure	phi	adjusted phi
1	0.333333	0.333333	0.333333	0.222222	1	1
2	0.589744	0.384615	0.358974	0.13215	0.552231	0.8375
3	0	0	0	0	NA	NA
4	0.513514	0.432432	0.432432	0.210373	0.849591	1
5	0	0	0	0	NA	NA
6	0.28	0.28	0.24	0.1616	0.801587	0.801587
7	0.0416667	0.0729167	0.0416667	0.0386285	0.743502	1
8	0.295455	0.340909	0.25	0.149277	0.690241	0.766578

Column panel 3/3 - row numbers identify the same source rows.

Row	kappa	wrong agreement	wrong chance uniform	both wrong n	capa onehot
1	1	1	NA	1	NA
2	0.507576	0.714286	0.115646	14	0.500232
3	NA	NA	NA	0	NA
4	0.838428	0.8125	0.130456	16	0.766784
5	NA	NA	NA	0	NA
6	0.801587	0.833333	0.111111	6	0.746087
7	0.712	1	NA	4	NA
8	0.686461	0.727273	0.111111	22	0.609273

self_consistency.csv

No rows were recorded in this output.

baseline_predictions.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	condition	team	method	train items	test items	observed beta
1	["gsm8k", "numeric", "direct", 17, 0.0]	["PHI01", "PHI04", "PHI02"]	independence	50	3	0
2	["gsm8k", "numeric", "direct", 17, 0.0]	["PHI01", "PHI04", "PHI02"]	single_factor_tetrachoric	50	3	0
3	["gsm8k", "numeric", "direct", 17, 0.0]	["PHI03", "PHI04", "PHI02"]	independence	50	1	0
4	["gsm8k", "numeric", "direct", 17, 0.0]	["PHI03", "PHI04", "PHI02"]	single_factor_tetrachoric	50	1	0
5	["gsm8k", "numeric", "reasoned", 17, 0.0]	["PHI01", "PHI03", "PHI02"]	independence	50	1	0
6	["gsm8k", "numeric", "reasoned", 17, 0.0]	["PHI01", "PHI03", "PHI02"]	single_factor_tetrachoric	50	1	0
7	["gsm8k", "numeric", "reasoned", 17, 0.0]	["PHI03", "PHI04", "PHI02"]	independence	50	6	0
8	["gsm8k", "numeric", "reasoned", 17, 0.0]	["PHI03", "PHI04", "PHI02"]	single_factor_tetrachoric	50	6	0
9	["mmlu_pro", "mc", "direct", 17, 0.0]	["PHI01", "PHI03", "PHI04"]	independence	50	5	0.2
10	["mmlu_pro", "mc", "direct", 17, 0.0]	["PHI01", "PHI03", "PHI04"]	single_factor_tetrachoric	50	5	0.2
11	["mmlu_pro", "mc", "direct", 17, 0.0]	["PHI01", "PHI03", "PHI02"]	independence	50	11	0.272727
12	["mmlu_pro", "mc", "direct", 17, 0.0]	["PHI01", "PHI03", "PHI02"]	single_factor_tetrachoric	50	11	0.272727
13	["mmlu_pro", "mc", "direct", 17, 0.0]	["PHI01", "PHI04", "PHI02"]	independence	50	11	0.0909091
14	["mmlu_pro", "mc", "direct", 17, 0.0]	["PHI01", "PHI04", "PHI02"]	single_factor_tetrachoric	50	11	0.0909091
15	["mmlu_pro", "mc", "direct", 17, 0.0]	["PHI03", "PHI04", "PHI02"]	independence	50	12	0.166667
16	["mmlu_pro", "mc", "direct", 17, 0.0]	["PHI03", "PHI04", "PHI02"]	single_factor_tetrachoric	50	12	0.166667
17	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["PHI01", "PHI03", "PHI04"]	independence	50	4	0.25
18	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["PHI01", "PHI03", "PHI04"]	single_factor_tetrachoric	50	4	0.25
19	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["PHI01", "PHI03", "PHI02"]	independence	50	12	0.333333
20	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["PHI01", "PHI03", "PHI02"]	single_factor_tetrachoric	50	12	0.333333
21	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["PHI01", "PHI04", "PHI02"]	independence	50	9	0.333333
22	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["PHI01", "PHI04", "PHI02"]	single_factor_tetrachoric	50	9	0.333333

Row	condition	team	method	train items	test items	observed beta
23	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["PHI03", "PHI04", "PHI02"]	independence	50	5	0.2
24	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["PHI03", "PHI04", "PHI02"]	single_factor_tetrachoric	50	5	0.2

Column panel 2/3 - row numbers identify the same source rows.

Row	beta low	beta high	predicted beta	beta mae	observed majority	predicted majority
1	0	0	2.44898e-06	2.44898e-06	1	0.975451
2	0	0	2.44898e-06	2.44898e-06	1	0.9795
3	0	0	4.08163e-10	4.08163e-10	1	0.999992
4	0	0	4.08163e-10	4.08163e-10	1	1
5	0	0	6.94444e-07	6.94444e-07	1	0.993036
6	0	0	6.94444e-07	6.94444e-07	1	0.995
7	0	0	4.62963e-07	4.62963e-07	0.833333	0.995356
8	0	0	4.62963e-07	4.62963e-07	0.833333	0.9955
9	0	0.6	0.0875	0.1125	0.6	0.554167
10	0	0.6	0.206804	0.00680441	0.6	0.5435
11	0.0909091	0.547727	0.066	0.206727	0.636364	0.59
12	0.0909091	0.547727	0.169701	0.103027	0.636364	0.556
13	0	0.272727	0.0385	0.0524091	0.545455	0.705833
14	0	0.272727	0.123886	0.0329771	0.545455	0.6745
15	0	0.416667	0.0320833	0.134583	0.666667	0.744167
16	0	0.416667	0.124548	0.0421191	0.666667	0.689
17	0	0.75	0.0611111	0.188889	0.75	0.645106
18	0	0.75	0.181892	0.0681079	0.75	0.6165
19	0.0833333	0.585417	0.0663221	0.267011	0.666667	0.632214
20	0.0833333	0.585417	0.207664	0.125669	0.666667	0.5755
21	0.108333	0.666667	0.0383721	0.294961	0.666667	0.733306
22	0.108333	0.666667	0.146831	0.186503	0.666667	0.689
23	0	0.6	0.050646	0.149354	0.8	0.679242
24	0	0.6	0.161826	0.0381745	0.8	0.6515

Column panel 3/3 - row numbers identify the same source rows.

Row	majority mae	copula rho	effective voters	policy
1	0.024549	0	NA	item holdout; capabilities calibrated on training items
2	0.0205	0	NA	item holdout; capabilities calibrated on training items
3	8.17245e-06	0	NA	item holdout; capabilities calibrated on training items
4	0	0	NA	item holdout; capabilities calibrated on training items
5	0.00696389	0	NA	item holdout; capabilities calibrated on training items
6	0.005	0	NA	item holdout; capabilities calibrated on training items
7	0.162023	0	NA	item holdout; capabilities calibrated on training items
8	0.162167	0	NA	item holdout; capabilities calibrated on training items
9	0.0458333	0.61226	NA	item holdout; capabilities calibrated on training items
10	0.0565	0.61226	NA	item holdout; capabilities calibrated on training items
11	0.0463636	0.61226	NA	item holdout; capabilities calibrated on training items
12	0.0803636	0.61226	NA	item holdout; capabilities calibrated on training items
13	0.160379	0.61226	NA	item holdout; capabilities calibrated on training items
14	0.129045	0.61226	NA	item holdout; capabilities calibrated on training items
15	0.0775	0.61226	NA	item holdout; capabilities calibrated on training items
16	0.0223333	0.61226	NA	item holdout; capabilities calibrated on training items
17	0.104894	0.617284	NA	item holdout; capabilities calibrated on training items
18	0.1335	0.617284	NA	item holdout; capabilities calibrated on training items
19	0.0344531	0.617284	NA	item holdout; capabilities calibrated on training items
20	0.0911667	0.617284	NA	item holdout; capabilities calibrated on training items
21	0.066639	0.617284	NA	item holdout; capabilities calibrated on training items
22	0.0223333	0.617284	NA	item holdout; capabilities calibrated on training items
23	0.120758	0.617284	NA	item holdout; capabilities calibrated on training items
24	0.1485	0.617284	NA	item holdout; capabilities calibrated on training items

7.10 Qwen - saved analysis tables

pairs.csv

Column panel 1/4 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	a
1	gsm8k	numeric	direct	17	0	QWE03
2	gsm8k	numeric	direct	17	0	QWE03
3	gsm8k	numeric	direct	17	0	QWE03
4	gsm8k	numeric	direct	17	0	QWE03
5	gsm8k	numeric	direct	17	0	QWE03
6	gsm8k	numeric	direct	17	0	QWE03
7	gsm8k	numeric	direct	17	0	QWE03
8	gsm8k	numeric	direct	17	0	QWE01
9	gsm8k	numeric	direct	17	0	QWE01
10	gsm8k	numeric	direct	17	0	QWE01
11	gsm8k	numeric	direct	17	0	QWE01
12	gsm8k	numeric	direct	17	0	QWE01
13	gsm8k	numeric	direct	17	0	QWE02
14	gsm8k	numeric	direct	17	0	QWE02
15	gsm8k	numeric	direct	17	0	QWE02
16	gsm8k	numeric	direct	17	0	QWE02
17	gsm8k	numeric	direct	17	0	QWE02
18	gsm8k	numeric	direct	17	0	QWE04
19	gsm8k	numeric	direct	17	0	QWE04
20	gsm8k	numeric	direct	17	0	QWE04
21	gsm8k	numeric	direct	17	0	QWE04
22	gsm8k	numeric	direct	17	0	QWE07
23	gsm8k	numeric	direct	17	0	QWE07
24	gsm8k	numeric	direct	17	0	QWE07
25	gsm8k	numeric	direct	17	0	QWE06

Row	benchmark	format	template	seed	temperature	a
26	gsm8k	numeric	direct	17	0	QWE06
27	gsm8k	numeric	direct	17	0	QWE09
28	gsm8k	numeric	reasoned	17	0	QWE03
29	gsm8k	numeric	reasoned	17	0	QWE03
30	gsm8k	numeric	reasoned	17	0	QWE03
31	gsm8k	numeric	reasoned	17	0	QWE03
32	gsm8k	numeric	reasoned	17	0	QWE03
33	gsm8k	numeric	reasoned	17	0	QWE03
34	gsm8k	numeric	reasoned	17	0	QWE01
35	gsm8k	numeric	reasoned	17	0	QWE01
36	gsm8k	numeric	reasoned	17	0	QWE01
37	gsm8k	numeric	reasoned	17	0	QWE01
38	gsm8k	numeric	reasoned	17	0	QWE01
39	gsm8k	numeric	reasoned	17	0	QWE02
40	gsm8k	numeric	reasoned	17	0	QWE02
41	gsm8k	numeric	reasoned	17	0	QWE02
42	gsm8k	numeric	reasoned	17	0	QWE02
43	gsm8k	numeric	reasoned	17	0	QWE07
44	gsm8k	numeric	reasoned	17	0	QWE07
45	gsm8k	numeric	reasoned	17	0	QWE07
46	gsm8k	numeric	reasoned	17	0	QWE06
47	gsm8k	numeric	reasoned	17	0	QWE06
48	gsm8k	numeric	reasoned	17	0	QWE09
49	mmlu_pro	mc	direct	17	0	QWE03
50	mmlu_pro	mc	direct	17	0	QWE03
51	mmlu_pro	mc	direct	17	0	QWE03
52	mmlu_pro	mc	direct	17	0	QWE03
53	mmlu_pro	mc	direct	17	0	QWE03
54	mmlu_pro	mc	direct	17	0	QWE02
55	mmlu_pro	mc	direct	17	0	QWE02

Row	benchmark	format	template	seed	temperature	a
56	mmlu_pro	mc	direct	17	0	QWE02
57	mmlu_pro	mc	direct	17	0	QWE02
58	mmlu_pro	mc	direct	17	0	QWE07
59	mmlu_pro	mc	direct	17	0	QWE07
60	mmlu_pro	mc	direct	17	0	QWE07
61	mmlu_pro	mc	direct	17	0	QWE06
62	mmlu_pro	mc	direct	17	0	QWE06
63	mmlu_pro	mc	direct	17	0	QWE09
64	mmlu_pro	mc	reasoned	17	0	QWE03
65	mmlu_pro	mc	reasoned	17	0	QWE03
66	mmlu_pro	mc	reasoned	17	0	QWE03
67	mmlu_pro	mc	reasoned	17	0	QWE03
68	mmlu_pro	mc	reasoned	17	0	QWE03
69	mmlu_pro	mc	reasoned	17	0	QWE02
70	mmlu_pro	mc	reasoned	17	0	QWE02
71	mmlu_pro	mc	reasoned	17	0	QWE02
72	mmlu_pro	mc	reasoned	17	0	QWE02
73	mmlu_pro	mc	reasoned	17	0	QWE07
74	mmlu_pro	mc	reasoned	17	0	QWE07
75	mmlu_pro	mc	reasoned	17	0	QWE07
76	mmlu_pro	mc	reasoned	17	0	QWE06
77	mmlu_pro	mc	reasoned	17	0	QWE06
78	mmlu_pro	mc	reasoned	17	0	QWE09

Column panel 2/4 - row numbers identify the same source rows.

Row	b	relationship	accuracy matched	n	error a	error b
1	QWE01	same_family|different_size	True	5	0	0
2	QWE02	same_family|different_size	False	35	0.0571429	0.142857
3	QWE04	same_family|different_size	False	1	0	1
4	QWE07	same_family	True	90	0.0333333	0.0555556

Row	b	relationship	accuracy matched	n	error a	error b
5	QWE06	same_family|different_size	True	89	0.0337079	0.011236
6	QWE09	student/base-it|same_family	True	80	0.0375	0.025
7	QWE08	same_family|different_size	True	89	0.0337079	0.0561798
8	QWE02	base/tuned|same_family	True	2	0	0
9	QWE07	same_family|different_size	False	5	0	0.2
10	QWE06	same_family|different_size	True	5	0	0
11	QWE09	same_family|different_size	True	5	0	0
12	QWE08	same_family	True	5	0	0
13	QWE04	same_family	True	1	1	1
14	QWE07	same_family|different_size	False	34	0.147059	0.0882353
15	QWE06	same_family|different_size	False	32	0.15625	0
16	QWE09	same_family|different_size	False	30	0.166667	0.0333333
17	QWE08	same_family	False	32	0.15625	0.0625
18	QWE07	same_family|different_size	False	1	1	0
19	QWE06	same_family|different_size	False	1	1	0
20	QWE09	same_family|different_size	False	1	1	0
21	QWE08	base/tuned|same_family	False	1	1	0
22	QWE06	same_family|different_size	True	88	0.0568182	0.0113636
23	QWE09	same_family	True	80	0.0625	0.025
24	QWE08	same_family|different_size	True	87	0.0574713	0.0574713
25	QWE09	same_family|different_size	True	76	0.0131579	0.0263158
26	QWE08	same_family|different_size	True	87	0.0114943	0.0574713
27	QWE08	same_family|different_size|same_lineage	True	78	0.025641	0.0641026
28	QWE01	same_family|different_size	False	8	0	0.25
29	QWE02	same_family|different_size	True	79	0.0506329	0.0632911
30	QWE07	same_family	True	91	0.043956	0.021978
31	QWE06	same_family|different_size	True	82	0.0609756	0.0243902
32	QWE09	student/base-it|same_family	True	83	0.060241	0.0361446
33	QWE08	same_family|different_size	True	31	0.129032	0.129032

Row	b	relationship	accuracy matched	n	error a	error b
34	QWE02	base/tuned|same_family	False	4	0.25	0
35	QWE07	same_family|different_size	False	6	0.166667	0
36	QWE06	same_family|different_size	False	7	0.285714	0
37	QWE09	same_family|different_size	False	7	0.285714	0
38	QWE08	same_family	False	4	0.5	0
39	QWE07	same_family|different_size	True	78	0.0641026	0.025641
40	QWE06	same_family|different_size	True	67	0.0447761	0.0149254
41	QWE09	same_family|different_size	True	73	0.0684932	0.0273973
42	QWE08	same_family	True	25	0.12	0.12
43	QWE06	same_family|different_size	True	77	0.012987	0.012987
44	QWE09	same_family	True	79	0.0253165	0.0253165
45	QWE08	same_family|different_size	False	28	0.0714286	0.142857
46	QWE09	same_family|different_size	True	73	0.0273973	0.0410959
47	QWE08	same_family|different_size	False	24	0.0416667	0.166667
48	QWE08	same_family|different_size|same_lineage	False	27	0.037037	0.148148
49	QWE02	same_family|different_size	False	83	0.240964	0.493976
50	QWE07	same_family	True	85	0.270588	0.258824
51	QWE06	same_family|different_size	False	59	0.20339	0.305085
52	QWE09	student/base-it|same_family	True	89	0.247191	0.280899
53	QWE08	same_family|different_size	False	56	0.232143	0.464286
54	QWE07	same_family|different_size	False	68	0.470588	0.25
55	QWE06	same_family|different_size	False	44	0.386364	0.272727
56	QWE09	same_family|different_size	False	72	0.486111	0.263889
57	QWE08	same_family	True	42	0.452381	0.428571
58	QWE06	same_family|different_size	False	57	0.245614	0.315789
59	QWE09	same_family	True	81	0.246914	0.271605
60	QWE08	same_family|different_size	False	54	0.259259	0.481481
61	QWE09	same_family|different_size	False	58	0.293103	0.224138
62	QWE08	same_family|different_size	False	43	0.232558	0.488372

Row	b	relationship	accuracy matched	n	error a	error b
63	QWE08	same_family|different_size|same_lineage	False	56	0.196429	0.464286
64	QWE02	same_family|different_size	False	95	0.294737	0.378947
65	QWE07	same_family	False	76	0.302632	0.236842
66	QWE06	same_family|different_size	True	51	0.254902	0.254902
67	QWE09	student/base-it|same_family	True	82	0.280488	0.243902
68	QWE08	same_family|different_size	False	49	0.306122	0.571429
69	QWE07	same_family|different_size	False	75	0.373333	0.24
70	QWE06	same_family|different_size	False	50	0.36	0.26
71	QWE09	same_family|different_size	False	81	0.345679	0.259259
72	QWE08	same_family	False	49	0.367347	0.591837
73	QWE06	same_family|different_size	True	51	0.254902	0.254902
74	QWE09	same_family	False	72	0.222222	0.277778
75	QWE08	same_family|different_size	False	49	0.204082	0.571429
76	QWE09	same_family|different_size	True	49	0.244898	0.244898
77	QWE08	same_family|different_size	False	34	0.235294	0.588235
78	QWE08	same_family|different_size|same_lineage	False	49	0.22449	0.571429

Column panel 3/4 - row numbers identify the same source rows.

Row	double fault	excess cofailure	phi	adjusted phi	kappa	wrong agreement
1	0	0	NA	NA	NA	NA
2	0.0285714	0.0204082	0.251259	0.416667	0.222222	0
3	0	0	NA	NA	0	NA
4	0.0111111	0.00925926	0.225189	0.294118	0.217391	1
5	0.011236	0.0108572	0.570752	1	0.491429	1
6	0.0125	0.0115625	0.389819	0.480519	0.381443	1
7	0	-0.0018937	-0.0455677	-0.0595238	-0.0439883	NA
8	0	0	NA	NA	NA	NA
9	0	0	NA	NA	0	NA
10	0	0	NA	NA	NA	NA

Row	double fault	excess cofailure	phi	adjusted phi	kappa	wrong agreement
11	0	0	NA	NA	NA	NA
12	0	0	NA	NA	NA	NA
13	1	0	NA	NA	NA	0
14	0	-0.0129758	-0.129171	-0.172414	-0.123967	NA
15	0	0	NA	NA	0	NA
16	0	-0.00555556	-0.0830455	-0.2	-0.0588235	NA
17	0	-0.00976562	-0.111111	-0.185185	-0.0980392	NA
18	0	0	NA	NA	0	NA
19	0	0	NA	NA	0	NA
20	0	0	NA	NA	0	NA
21	0	0	NA	NA	0	NA
22	0.0113636	0.010718	0.436812	1	0.320463	1
23	0.025	0.0234375	0.620174	1	0.555556	1
24	0.0114943	0.00819131	0.15122	0.15122	0.15122	1
25	0.0131579	0.0128116	0.702377	1	0.660714	1
26	0	-0.000660589	-0.0266274	-0.0609756	-0.0195312	NA
27	0	-0.00164366	-0.0424553	-0.0684932	-0.0380228	NA
28	0	0	NA	NA	0	NA
29	0.0126582	0.00945361	0.177089	0.199324	0.175857	1
30	0.010989	0.0100229	0.333487	0.477011	0.313208	1
31	0.0243902	0.022903	0.620484	1	0.555957	1
32	0.0240964	0.021919	0.493559	0.645299	0.476341	1
33	0.0967742	0.0801249	0.712963	0.712963	0.712963	0.666667
34	0	0	NA	NA	0	NA
35	0	0	NA	NA	0	NA
36	0	0	NA	NA	0	NA
37	0	0	NA	NA	0	NA
38	0	0	NA	NA	0	NA
39	0.0128205	0.0111769	0.288696	0.465753	0.258555	1
40	0	-0.0006683	-0.0266501	-0.046875	-0.0229008	NA

Row	double fault	excess cofailure	phi	adjusted phi	kappa	wrong agreement
41	0	-0.00187652	-0.045511	-0.0735294	-0.0407332	NA
42	0.04	0.0256	0.242424	0.242424	0.242424	1
43	0	-0.000168663	-0.0131579	-0.0131579	-0.0131579	NA
44	0	-0.000640923	-0.025974	-0.025974	-0.025974	NA
45	0.0357143	0.0255102	0.283069	0.416667	0.263158	1
46	0.0273973	0.0262713	0.810726	1	0.793201	1
47	0.0416667	0.0347222	0.466252	1	0.357143	1
48	0.037037	0.0315501	0.47027	1	0.362205	1
49	0.192771	0.0737407	0.344875	0.604762	0.296816	0.375
50	0.188235	0.118201	0.607458	0.6261	0.607181	0.5
51	0.118644	0.0565929	0.30535	0.400407	0.294469	0.714286
52	0.202247	0.132812	0.685026	0.747159	0.682452	0.611111
53	0.178571	0.0707908	0.336202	0.569231	0.29443	0.5
54	0.191176	0.0735294	0.340207	0.555556	0.30303	0.461538
55	0.204545	0.0991736	0.45733	0.592593	0.442396	0.222222
56	0.180556	0.0522762	0.237312	0.385491	0.211884	0.461538
57	0.261905	0.0680272	0.276184	0.289855	0.275862	0.363636
58	0.175439	0.0978763	0.489171	0.582418	0.481818	0.7
59	0.17284	0.105777	0.551496	0.588136	0.550357	0.857143
60	0.203704	0.0788752	0.36022	0.586735	0.321229	0.272727
61	0.137931	0.0722354	0.38055	0.45591	0.374422	0.5
62	0.186047	0.0724716	0.343184	0.609091	0.293538	0.625
63	0.142857	0.0516582	0.260714	0.490909	0.216	0.5
64	0.2	0.0883102	0.399268	0.482446	0.392224	0.473684
65	0.184211	0.112535	0.576184	0.681342	0.568182	0.571429
66	0.156863	0.0918877	0.483806	0.483806	0.483806	0.5
67	0.182927	0.114515	0.593595	0.652542	0.590944	0.6
68	0.244898	0.0699708	0.306786	0.533333	0.265193	0.583333
69	0.16	0.0704	0.340795	0.468085	0.324324	0.5
70	0.18	0.0864	0.410365	0.519231	0.399261	0.444444

Row	double fault	excess cofailure	phi	adjusted phi	kappa	wrong agreement
71	0.185185	0.0955647	0.458527	0.563342	0.44898	0.466667
72	0.306122	0.088713	0.374411	0.591667	0.338364	0.533333
73	0.215686	0.150711	0.793522	0.793522	0.793522	0.454545
74	0.152778	0.0910494	0.488958	0.567308	0.483607	0.545455
75	0.183673	0.0670554	0.336206	0.766667	0.247312	0.555556
76	0.142857	0.0828821	0.448198	0.448198	0.448198	0.571429
77	0.176471	0.0380623	0.182323	0.392857	0.139241	0.833333
78	0.204082	0.0758017	0.367109	0.787879	0.281081	0.2

Column panel 4/4 - row numbers identify the same source rows.

Row	wrong chance uniform	both wrong n	capa onehot
1	NA	0	NA
2	NA	1	NA
3	NA	0	NA
4	NA	1	NA
5	NA	1	NA
6	NA	1	NA
7	NA	0	NA
8	NA	0	NA
9	NA	0	NA
10	NA	0	NA
11	NA	0	NA
12	NA	0	NA
13	NA	1	NA
14	NA	0	NA
15	NA	0	NA
16	NA	0	NA
17	NA	0	NA
18	NA	0	NA
19	NA	0	NA

Row	wrong chance uniform	both wrong n	capa onehot
20	NA	0	NA
21	NA	0	NA
22	NA	1	NA
23	NA	2	NA
24	NA	1	NA
25	NA	1	NA
26	NA	0	NA
27	NA	0	NA
28	NA	0	NA
29	NA	1	NA
30	NA	1	NA
31	NA	2	NA
32	NA	2	NA
33	NA	3	NA
34	NA	0	NA
35	NA	0	NA
36	NA	0	NA
37	NA	0	NA
38	NA	0	NA
39	NA	1	NA
40	NA	0	NA
41	NA	0	NA
42	NA	1	NA
43	NA	0	NA
44	NA	0	NA
45	NA	1	NA
46	NA	2	NA
47	NA	1	NA
48	NA	1	NA
49	0.111111	16	0.218536

Row	wrong chance uniform	both wrong n	capa onehot
50	0.111979	16	0.451855
51	0.111111	7	0.304555
52	0.111883	18	0.55067
53	0.111111	10	0.254465
54	0.113553	13	0.249926
55	0.118166	9	0.242859
56	0.111111	13	0.197472
57	0.116883	11	0.208054
58	0.114286	10	0.444976
59	0.112103	14	0.554229
60	0.113997	11	0.197336
61	0.115079	8	0.299989
62	0.142857	8	0.294023
63	0.111111	8	0.19972
64	0.111111	19	0.308909
65	0.112103	14	0.455135
66	0.111111	8	0.371595
67	0.112037	15	0.482235
68	0.111111	12	0.280012
69	0.113757	12	0.271552
70	0.114638	9	0.300785
71	0.113228	15	0.338908
72	0.131746	15	0.314019
73	0.113997	11	0.551139
74	0.112374	11	0.387056
75	0.111111	9	0.239446
76	0.111111	7	0.371986
77	0.111111	6	0.251766
78	0.114286	10	0.153911

ladder.csv

Column panel 1/2 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	relationship
1	gsm8k	numeric	direct	17	0	base/tuned
2	gsm8k	numeric	direct	17	0	base/tuned
3	gsm8k	numeric	direct	17	0	same_family
4	gsm8k	numeric	direct	17	0	same_family
5	gsm8k	numeric	direct	17	0	different_size
6	gsm8k	numeric	direct	17	0	different_size
7	gsm8k	numeric	direct	17	0	same_lineage
8	gsm8k	numeric	direct	17	0	same_lineage
9	gsm8k	numeric	direct	17	0	student/base-it
10	gsm8k	numeric	direct	17	0	student/base-it
11	gsm8k	numeric	reasoned	17	0	base/tuned
12	gsm8k	numeric	reasoned	17	0	same_family
13	gsm8k	numeric	reasoned	17	0	same_family
14	gsm8k	numeric	reasoned	17	0	different_size
15	gsm8k	numeric	reasoned	17	0	different_size
16	gsm8k	numeric	reasoned	17	0	same_lineage
17	gsm8k	numeric	reasoned	17	0	student/base-it
18	gsm8k	numeric	reasoned	17	0	student/base-it
19	mmlu_pro	mc	direct	17	0	same_family
20	mmlu_pro	mc	direct	17	0	same_family
21	mmlu_pro	mc	direct	17	0	different_size
22	mmlu_pro	mc	direct	17	0	same_lineage
23	mmlu_pro	mc	direct	17	0	student/base-it
24	mmlu_pro	mc	direct	17	0	student/base-it
25	mmlu_pro	mc	reasoned	17	0	same_family
26	mmlu_pro	mc	reasoned	17	0	same_family
27	mmlu_pro	mc	reasoned	17	0	different_size

Row	benchmark	format	template	seed	temperature	relationship
28	mmlu_pro	mc	reasoned	17	0	different_size
29	mmlu_pro	mc	reasoned	17	0	same_lineage
30	mmlu_pro	mc	reasoned	17	0	student/base-it
31	mmlu_pro	mc	reasoned	17	0	student/base-it

Column panel 2/2 - row numbers identify the same source rows.

Row	matched only	pairs	mean adjusted phi	low	high
1	False	2	NA	NA	NA
2	True	1	NA	NA	NA
3	False	27	0.328281	-0.0388237	0.505301
4	True	16	0.473686	-0.0397364	0.613628
5	False	19	0.300648	-0.104773	0.479498
6	True	10	0.423175	-0.0645729	0.487568
7	False	1	-0.0684932	-0.133003	-0.0264386
8	True	1	-0.0684932	-0.136782	-0.0247762
9	False	1	0.480519	-0.0463085	1
10	True	1	0.480519	-0.050538	1
11	False	1	NA	NA	NA
12	False	21	0.46666	-0.0428083	0.754328
13	True	12	0.381937	-0.049698	0.679494
14	False	15	0.51465	-0.0541094	0.82901
15	True	8	0.40556	-0.0136164	0.743002
16	False	1	1	1	1
17	False	1	0.645299	-0.0568046	1
18	True	1	0.645299	-0.0564926	1
19	False	15	0.538957	0.374693	0.699252
20	True	4	0.562812	0.423861	0.747627
21	False	11	0.530282	0.318806	0.749925
22	False	1	0.490909	0.0568083	1
23	False	1	0.747159	0.557089	1

Row	matched only	pairs	mean adjusted phi	low	high
24	True	1	0.747159	0.575632	1
25	False	15	0.582148	0.438322	0.753541
26	True	4	0.594517	0.388152	0.829141
27	False	11	0.567215	0.396023	0.783938
28	True	3	0.575175	0.450166	0.840624
29	False	1	0.787879	0.368314	1
30	False	1	0.652542	0.424846	0.884978
31	True	1	0.652542	0.449659	0.880445

teams.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	team
1	gsm8k	numeric	direct	17	0	["QWE03", "QWE01", "QWE02"]
2	gsm8k	numeric	direct	17	0	["QWE03", "QWE01", "QWE07"]
3	gsm8k	numeric	direct	17	0	["QWE03", "QWE01", "QWE06"]
4	gsm8k	numeric	direct	17	0	["QWE03", "QWE01", "QWE09"]
5	gsm8k	numeric	direct	17	0	["QWE03", "QWE01", "QWE08"]
6	gsm8k	numeric	direct	17	0	["QWE03", "QWE02", "QWE04"]
7	gsm8k	numeric	direct	17	0	["QWE03", "QWE02", "QWE07"]
8	gsm8k	numeric	direct	17	0	["QWE03", "QWE02", "QWE06"]
9	gsm8k	numeric	direct	17	0	["QWE03", "QWE02", "QWE09"]
10	gsm8k	numeric	direct	17	0	["QWE03", "QWE02", "QWE08"]
11	gsm8k	numeric	direct	17	0	["QWE03", "QWE04", "QWE07"]
12	gsm8k	numeric	direct	17	0	["QWE03", "QWE04", "QWE06"]
13	gsm8k	numeric	direct	17	0	["QWE03", "QWE04", "QWE09"]
14	gsm8k	numeric	direct	17	0	["QWE03", "QWE04", "QWE08"]
15	gsm8k	numeric	direct	17	0	["QWE03", "QWE07", "QWE06"]
16	gsm8k	numeric	direct	17	0	["QWE03", "QWE07", "QWE09"]
17	gsm8k	numeric	direct	17	0	["QWE03", "QWE07", "QWE08"]

Row	benchmark	format	template	seed	temperature	team
18	gsm8k	numeric	direct	17	0	["QWE03", "QWE06", "QWE09"]
19	gsm8k	numeric	direct	17	0	["QWE03", "QWE06", "QWE08"]
20	gsm8k	numeric	direct	17	0	["QWE03", "QWE09", "QWE08"]
21	gsm8k	numeric	direct	17	0	["QWE01", "QWE02", "QWE07"]
22	gsm8k	numeric	direct	17	0	["QWE01", "QWE02", "QWE06"]
23	gsm8k	numeric	direct	17	0	["QWE01", "QWE02", "QWE09"]
24	gsm8k	numeric	direct	17	0	["QWE01", "QWE02", "QWE08"]
25	gsm8k	numeric	direct	17	0	["QWE01", "QWE07", "QWE06"]
26	gsm8k	numeric	direct	17	0	["QWE01", "QWE07", "QWE09"]
27	gsm8k	numeric	direct	17	0	["QWE01", "QWE07", "QWE08"]
28	gsm8k	numeric	direct	17	0	["QWE01", "QWE06", "QWE09"]
29	gsm8k	numeric	direct	17	0	["QWE01", "QWE06", "QWE08"]
30	gsm8k	numeric	direct	17	0	["QWE01", "QWE09", "QWE08"]
31	gsm8k	numeric	direct	17	0	["QWE02", "QWE04", "QWE07"]
32	gsm8k	numeric	direct	17	0	["QWE02", "QWE04", "QWE06"]
33	gsm8k	numeric	direct	17	0	["QWE02", "QWE04", "QWE09"]
34	gsm8k	numeric	direct	17	0	["QWE02", "QWE04", "QWE08"]
35	gsm8k	numeric	direct	17	0	["QWE02", "QWE07", "QWE06"]
36	gsm8k	numeric	direct	17	0	["QWE02", "QWE07", "QWE09"]
37	gsm8k	numeric	direct	17	0	["QWE02", "QWE07", "QWE08"]
38	gsm8k	numeric	direct	17	0	["QWE02", "QWE06", "QWE09"]
39	gsm8k	numeric	direct	17	0	["QWE02", "QWE06", "QWE08"]
40	gsm8k	numeric	direct	17	0	["QWE02", "QWE09", "QWE08"]
41	gsm8k	numeric	direct	17	0	["QWE04", "QWE07", "QWE06"]
42	gsm8k	numeric	direct	17	0	["QWE04", "QWE07", "QWE09"]
43	gsm8k	numeric	direct	17	0	["QWE04", "QWE07", "QWE08"]
44	gsm8k	numeric	direct	17	0	["QWE04", "QWE06", "QWE09"]
45	gsm8k	numeric	direct	17	0	["QWE04", "QWE06", "QWE08"]
46	gsm8k	numeric	direct	17	0	["QWE04", "QWE09", "QWE08"]
47	gsm8k	numeric	direct	17	0	["QWE07", "QWE06", "QWE09"]

Row	benchmark	format	template	seed	temperature	team
48	gsm8k	numeric	direct	17	0	["QWE07", "QWE06", "QWE08"]
49	gsm8k	numeric	direct	17	0	["QWE07", "QWE09", "QWE08"]
50	gsm8k	numeric	direct	17	0	["QWE06", "QWE09", "QWE08"]
51	gsm8k	numeric	direct	17	0	["QWE03", "QWE01", "QWE02", "QWE07", "QWE09"]
52	gsm8k	numeric	direct	17	0	["QWE03", "QWE01", "QWE02", "QWE07", "QWE08"]
53	gsm8k	numeric	direct	17	0	["QWE03", "QWE01", "QWE02", "QWE06", "QWE09"]
54	gsm8k	numeric	direct	17	0	["QWE03", "QWE01", "QWE02", "QWE09", "QWE08"]
55	gsm8k	numeric	direct	17	0	["QWE03", "QWE01", "QWE07", "QWE06", "QWE09"]
56	gsm8k	numeric	direct	17	0	["QWE03", "QWE01", "QWE07", "QWE09", "QWE08"]
57	gsm8k	numeric	direct	17	0	["QWE03", "QWE01", "QWE06", "QWE09", "QWE08"]
58	gsm8k	numeric	direct	17	0	["QWE03", "QWE02", "QWE04", "QWE07", "QWE06"]
59	gsm8k	numeric	direct	17	0	["QWE03", "QWE02", "QWE04", "QWE07", "QWE09"]
60	gsm8k	numeric	direct	17	0	["QWE03", "QWE02", "QWE04", "QWE07", "QWE08"]
61	gsm8k	numeric	direct	17	0	["QWE03", "QWE02", "QWE04", "QWE06", "QWE08"]
62	gsm8k	numeric	direct	17	0	["QWE03", "QWE02", "QWE04", "QWE09", "QWE08"]
63	gsm8k	numeric	direct	17	0	["QWE03", "QWE02", "QWE07", "QWE06", "QWE09"]
64	gsm8k	numeric	direct	17	0	["QWE03", "QWE02", "QWE07", "QWE06", "QWE08"]
65	gsm8k	numeric	direct	17	0	["QWE03", "QWE02", "QWE07", "QWE09", "QWE08"]
66	gsm8k	numeric	direct	17	0	["QWE03", "QWE04", "QWE07", "QWE06", "QWE09"]
67	gsm8k	numeric	direct	17	0	["QWE03", "QWE04", "QWE07", "QWE09", "QWE08"]
68	gsm8k	numeric	direct	17	0	["QWE03", "QWE04", "QWE06", "QWE09", "QWE08"]

Row	benchmark	format	template	seed	temperature	team
69	gsm8k	numeric	direct	17	0	["QWE03", "QWE07", "QWE06", "QWE09", "QWE08"]
70	gsm8k	numeric	direct	17	0	["QWE01", "QWE02", "QWE07", "QWE06", "QWE09"]
71	gsm8k	numeric	direct	17	0	["QWE01", "QWE02", "QWE07", "QWE06", "QWE08"]
72	gsm8k	numeric	direct	17	0	["QWE01", "QWE02", "QWE07", "QWE09", "QWE08"]
73	gsm8k	numeric	direct	17	0	["QWE01", "QWE02", "QWE06", "QWE09", "QWE08"]
74	gsm8k	numeric	direct	17	0	["QWE01", "QWE07", "QWE06", "QWE09", "QWE08"]
75	gsm8k	numeric	direct	17	0	["QWE02", "QWE04", "QWE07", "QWE06", "QWE09"]
76	gsm8k	numeric	direct	17	0	["QWE02", "QWE04", "QWE07", "QWE06", "QWE08"]
77	gsm8k	numeric	direct	17	0	["QWE02", "QWE04", "QWE07", "QWE09", "QWE08"]
78	gsm8k	numeric	direct	17	0	["QWE04", "QWE07", "QWE06", "QWE09", "QWE08"]
79	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE01", "QWE02"]
80	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE01", "QWE07"]
81	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE01", "QWE06"]
82	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE01", "QWE09"]
83	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE01", "QWE08"]
84	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE02", "QWE07"]
85	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE02", "QWE06"]
86	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE02", "QWE09"]
87	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE02", "QWE08"]
88	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE07", "QWE06"]
89	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE07", "QWE09"]
90	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE07", "QWE08"]
91	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE06", "QWE09"]
92	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE06", "QWE08"]
93	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE09", "QWE08"]

Row	benchmark	format	template	seed	temperature	team
94	gsm8k	numeric	reasoned	17	0	["QWE01", "QWE02", "QWE07"]
95	gsm8k	numeric	reasoned	17	0	["QWE01", "QWE02", "QWE06"]
96	gsm8k	numeric	reasoned	17	0	["QWE01", "QWE02", "QWE09"]
97	gsm8k	numeric	reasoned	17	0	["QWE01", "QWE02", "QWE08"]
98	gsm8k	numeric	reasoned	17	0	["QWE01", "QWE07", "QWE06"]
99	gsm8k	numeric	reasoned	17	0	["QWE01", "QWE07", "QWE09"]
100	gsm8k	numeric	reasoned	17	0	["QWE01", "QWE07", "QWE08"]
101	gsm8k	numeric	reasoned	17	0	["QWE01", "QWE06", "QWE09"]
102	gsm8k	numeric	reasoned	17	0	["QWE01", "QWE06", "QWE08"]
103	gsm8k	numeric	reasoned	17	0	["QWE01", "QWE09", "QWE08"]
104	gsm8k	numeric	reasoned	17	0	["QWE02", "QWE07", "QWE06"]
105	gsm8k	numeric	reasoned	17	0	["QWE02", "QWE07", "QWE09"]
106	gsm8k	numeric	reasoned	17	0	["QWE02", "QWE07", "QWE08"]
107	gsm8k	numeric	reasoned	17	0	["QWE02", "QWE06", "QWE09"]
108	gsm8k	numeric	reasoned	17	0	["QWE02", "QWE06", "QWE08"]
109	gsm8k	numeric	reasoned	17	0	["QWE02", "QWE09", "QWE08"]
110	gsm8k	numeric	reasoned	17	0	["QWE07", "QWE06", "QWE09"]
111	gsm8k	numeric	reasoned	17	0	["QWE07", "QWE06", "QWE08"]
112	gsm8k	numeric	reasoned	17	0	["QWE07", "QWE09", "QWE08"]
113	gsm8k	numeric	reasoned	17	0	["QWE06", "QWE09", "QWE08"]
114	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE01", "QWE02", "QWE07", "QWE06"]
115	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE01", "QWE02", "QWE07", "QWE09"]
116	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE01", "QWE02", "QWE07", "QWE08"]
117	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE01", "QWE02", "QWE06", "QWE09"]
118	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE01", "QWE02", "QWE06", "QWE08"]
119	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE01", "QWE02", "QWE09", "QWE08"]

Row	benchmark	format	template	seed	temperature	team
120	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE01", "QWE07", "QWE06", "QWE08"]
121	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE01", "QWE07", "QWE09", "QWE08"]
122	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE01", "QWE06", "QWE09", "QWE08"]
123	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE02", "QWE07", "QWE06", "QWE09"]
124	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE02", "QWE07", "QWE06", "QWE08"]
125	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE02", "QWE07", "QWE09", "QWE08"]
126	gsm8k	numeric	reasoned	17	0	["QWE03", "QWE07", "QWE06", "QWE09", "QWE08"]
127	gsm8k	numeric	reasoned	17	0	["QWE01", "QWE02", "QWE07", "QWE06", "QWE09"]
128	gsm8k	numeric	reasoned	17	0	["QWE01", "QWE02", "QWE07", "QWE06", "QWE08"]
129	gsm8k	numeric	reasoned	17	0	["QWE01", "QWE02", "QWE06", "QWE09", "QWE08"]
130	gsm8k	numeric	reasoned	17	0	["QWE01", "QWE07", "QWE06", "QWE09", "QWE08"]
131	gsm8k	numeric	reasoned	17	0	["QWE02", "QWE07", "QWE06", "QWE09", "QWE08"]
132	mmlu_pro	mc	direct	17	0	["QWE03", "QWE02", "QWE07"]
133	mmlu_pro	mc	direct	17	0	["QWE03", "QWE02", "QWE06"]
134	mmlu_pro	mc	direct	17	0	["QWE03", "QWE02", "QWE09"]
135	mmlu_pro	mc	direct	17	0	["QWE03", "QWE02", "QWE08"]
136	mmlu_pro	mc	direct	17	0	["QWE03", "QWE07", "QWE06"]
137	mmlu_pro	mc	direct	17	0	["QWE03", "QWE07", "QWE09"]
138	mmlu_pro	mc	direct	17	0	["QWE03", "QWE07", "QWE08"]
139	mmlu_pro	mc	direct	17	0	["QWE03", "QWE06", "QWE09"]
140	mmlu_pro	mc	direct	17	0	["QWE03", "QWE06", "QWE08"]
141	mmlu_pro	mc	direct	17	0	["QWE03", "QWE09", "QWE08"]
142	mmlu_pro	mc	direct	17	0	["QWE02", "QWE07", "QWE06"]
143	mmlu_pro	mc	direct	17	0	["QWE02", "QWE07", "QWE09"]

Row	benchmark	format	template	seed	temperature	team
144	mmlu_pro	mc	direct	17	0	["QWE02", "QWE07", "QWE08"]
145	mmlu_pro	mc	direct	17	0	["QWE02", "QWE06", "QWE09"]
146	mmlu_pro	mc	direct	17	0	["QWE02", "QWE06", "QWE08"]
147	mmlu_pro	mc	direct	17	0	["QWE02", "QWE09", "QWE08"]
148	mmlu_pro	mc	direct	17	0	["QWE07", "QWE06", "QWE09"]
149	mmlu_pro	mc	direct	17	0	["QWE07", "QWE06", "QWE08"]
150	mmlu_pro	mc	direct	17	0	["QWE07", "QWE09", "QWE08"]
151	mmlu_pro	mc	direct	17	0	["QWE06", "QWE09", "QWE08"]
152	mmlu_pro	mc	direct	17	0	["QWE03", "QWE02", "QWE07", "QWE09", "QWE08"]
153	mmlu_pro	mc	direct	17	0	["QWE03", "QWE02", "QWE06", "QWE09", "QWE08"]
154	mmlu_pro	mc	direct	17	0	["QWE03", "QWE07", "QWE06", "QWE09", "QWE08"]
155	mmlu_pro	mc	direct	17	0	["QWE02", "QWE07", "QWE06", "QWE09", "QWE08"]
156	mmlu_pro	mc	reasoned	17	0	["QWE03", "QWE02", "QWE07"]
157	mmlu_pro	mc	reasoned	17	0	["QWE03", "QWE02", "QWE06"]
158	mmlu_pro	mc	reasoned	17	0	["QWE03", "QWE02", "QWE09"]
159	mmlu_pro	mc	reasoned	17	0	["QWE03", "QWE02", "QWE08"]
160	mmlu_pro	mc	reasoned	17	0	["QWE03", "QWE07", "QWE06"]
161	mmlu_pro	mc	reasoned	17	0	["QWE03", "QWE07", "QWE09"]
162	mmlu_pro	mc	reasoned	17	0	["QWE03", "QWE07", "QWE08"]
163	mmlu_pro	mc	reasoned	17	0	["QWE03", "QWE06", "QWE09"]
164	mmlu_pro	mc	reasoned	17	0	["QWE03", "QWE06", "QWE08"]
165	mmlu_pro	mc	reasoned	17	0	["QWE03", "QWE09", "QWE08"]
166	mmlu_pro	mc	reasoned	17	0	["QWE02", "QWE07", "QWE06"]
167	mmlu_pro	mc	reasoned	17	0	["QWE02", "QWE07", "QWE09"]
168	mmlu_pro	mc	reasoned	17	0	["QWE02", "QWE07", "QWE08"]
169	mmlu_pro	mc	reasoned	17	0	["QWE02", "QWE06", "QWE09"]
170	mmlu_pro	mc	reasoned	17	0	["QWE02", "QWE06", "QWE08"]
171	mmlu_pro	mc	reasoned	17	0	["QWE02", "QWE09", "QWE08"]

Row	benchmark	format	template	seed	temperature	team
172	mmlu_pro	mc	reasoned	17	0	["QWE07", "QWE06", "QWE09"]
173	mmlu_pro	mc	reasoned	17	0	["QWE07", "QWE06", "QWE08"]
174	mmlu_pro	mc	reasoned	17	0	["QWE07", "QWE09", "QWE08"]
175	mmlu_pro	mc	reasoned	17	0	["QWE06", "QWE09", "QWE08"]
176	mmlu_pro	mc	reasoned	17	0	["QWE03", "QWE02", "QWE07", "QWE06", "QWE08"]
177	mmlu_pro	mc	reasoned	17	0	["QWE03", "QWE02", "QWE07", "QWE09", "QWE08"]
178	mmlu_pro	mc	reasoned	17	0	["QWE03", "QWE02", "QWE06", "QWE09", "QWE08"]
179	mmlu_pro	mc	reasoned	17	0	["QWE03", "QWE07", "QWE06", "QWE09", "QWE08"]

Column panel 2/3 - row numbers identify the same source rows.

Row	n	beta	beta low	beta high	selection ceiling	correctness majority
1	2	0	0	0	1	1
2	5	0	0	0	1	1
3	5	0	0	0	1	1
4	5	0	0	0	1	1
5	5	0	0	0	1	1
6	1	0	0	0	1	0
7	34	0	0	0	1	0.970588
8	32	0	0	0	1	0.96875
9	30	0	0	0	1	0.966667
10	32	0	0	0	1	0.96875
11	1	0	0	0	1	1
12	1	0	0	0	1	1
13	1	0	0	0	1	1
14	1	0	0	0	1	1
15	86	0.0116279	0	0.0348837	0.988372	0.988372
16	79	0.0126582	0	0.0379747	0.987342	0.974684
17	85	0	0	0	1	0.976471

Row	n	beta	beta low	beta high	selection ceiling	correctness majority
18	75	0.0133333	0	0.04	0.986667	0.986667
19	84	0	0	0	1	0.988095
20	77	0	0	0	1	0.987013
21	2	0	0	0	1	1
22	2	0	0	0	1	1
23	2	0	0	0	1	1
24	2	0	0	0	1	1
25	5	0	0	0	1	1
26	5	0	0	0	1	1
27	5	0	0	0	1	1
28	5	0	0	0	1	1
29	5	0	0	0	1	1
30	5	0	0	0	1	1
31	1	0	0	0	1	0
32	1	0	0	0	1	0
33	1	0	0	0	1	0
34	1	0	0	0	1	0
35	32	0	0	0	1	1
36	30	0	0	0	1	0.966667
37	32	0	0	0	1	1
38	28	0	0	0	1	1
39	30	0	0	0	1	1
40	29	0	0	0	1	1
41	1	0	0	0	1	1
42	1	0	0	0	1	1
43	1	0	0	0	1	1
44	1	0	0	0	1	1
45	1	0	0	0	1	1
46	1	0	0	0	1	1
47	76	0.0131579	0	0.0526316	0.986842	0.973684

Row	n	beta	beta low	beta high	selection ceiling	correctness majority
48	83	0	0	0	1	0.975904
49	77	0	0	0	1	0.961039
50	73	0	0	0	1	0.986301
51	2	0	0	0	1	1
52	2	0	0	0	1	1
53	2	0	0	0	1	1
54	2	0	0	0	1	1
55	5	0	0	0	1	1
56	5	0	0	0	1	1
57	5	0	0	0	1	1
58	1	0	0	0	1	1
59	1	0	0	0	1	1
60	1	0	0	0	1	1
61	1	0	0	0	1	1
62	1	0	0	0	1	1
63	28	0	0	0	1	1
64	30	0	0	0	1	1
65	29	0	0	0	1	1
66	1	0	0	0	1	1
67	1	0	0	0	1	1
68	1	0	0	0	1	1
69	72	0	0	0	1	0.986111
70	2	0	0	0	1	1
71	2	0	0	0	1	1
72	2	0	0	0	1	1
73	2	0	0	0	1	1
74	5	0	0	0	1	1
75	1	0	0	0	1	1
76	1	0	0	0	1	1
77	1	0	0	0	1	1

Row	n	beta	beta low	beta high	selection ceiling	correctness majority
78	1	0	0	0	1	1
79	4	0	0	0	1	1
80	6	0	0	0	1	1
81	7	0	0	0	1	1
82	7	0	0	0	1	1
83	4	0	0	0	1	1
84	78	0.0128205	0	0.0384615	0.987179	0.987179
85	67	0	0	0	1	0.970149
86	73	0	0	0	1	0.972603
87	25	0.04	0	0.12	0.96	0.92
88	77	0	0	0	1	0.974026
89	79	0	0	0	1	0.974684
90	28	0.0357143	0	0.107143	0.964286	0.892857
91	73	0.0273973	0	0.0684932	0.972603	0.972603
92	24	0.0416667	0	0.125	0.958333	0.875
93	27	0.037037	0	0.111111	0.962963	0.888889
94	4	0	0	0	1	1
95	4	0	0	0	1	1
96	4	0	0	0	1	1
97	2	0	0	0	1	1
98	5	0	0	0	1	1
99	5	0	0	0	1	1
100	3	0	0	0	1	1
101	6	0	0	0	1	1
102	3	0	0	0	1	1
103	4	0	0	0	1	1
104	66	0	0	0	1	0.984848
105	72	0	0	0	1	0.986111
106	25	0.04	0	0.12	0.96	0.96
107	64	0	0	0	1	0.984375

Row	n	beta	beta low	beta high	selection ceiling	correctness majority
108	19	0	0	0	1	0.947368
109	22	0	0	0	1	0.954545
110	69	0	0	0	1	0.985507
111	21	0	0	0	1	0.904762
112	25	0	0	0	1	0.92
113	22	0.0454545	0	0.136364	0.954545	0.954545
114	4	0	0	0	1	1
115	4	0	0	0	1	1
116	2	0	0	0	1	1
117	4	0	0	0	1	1
118	2	0	0	0	1	1
119	2	0	0	0	1	1
120	2	0	0	0	1	1
121	3	0	0	0	1	1
122	3	0	0	0	1	1
123	63	0	0	0	1	0.984127
124	19	0	0	0	1	0.947368
125	22	0	0	0	1	0.954545
126	20	0	0	0	1	0.9
127	4	0	0	0	1	1
128	2	0	0	0	1	1
129	2	0	0	0	1	1
130	2	0	0	0	1	1
131	18	0	0	0	1	0.944444
132	68	0.132353	0.0588235	0.205882	0.867647	0.720588
133	44	0.0681818	0	0.159091	0.931818	0.727273
134	72	0.138889	0.0694444	0.236111	0.861111	0.75
135	42	0.142857	0.047619	0.261905	0.857143	0.738095
136	57	0.122807	0.0350877	0.210526	0.877193	0.807018
137	81	0.148148	0.0864198	0.222531	0.851852	0.753086

Row	n	beta	beta low	beta high	selection ceiling	correctness majority
138	54	0.148148	0.0921296	0.240741	0.851852	0.722222
139	58	0.0862069	0.0172414	0.155172	0.913793	0.793103
140	43	0.0930233	0.0226744	0.186047	0.906977	0.744186
141	56	0.142857	0.0535714	0.25	0.857143	0.785714
142	42	0.119048	0.047619	0.238095	0.880952	0.690476
143	64	0.109375	0.046875	0.203125	0.890625	0.71875
144	40	0.2	0.099375	0.300625	0.8	0.725
145	43	0.0697674	0	0.139535	0.930233	0.744186
146	30	0.133333	0.0333333	0.233333	0.866667	0.7
147	42	0.119048	0.0238095	0.238095	0.880952	0.738095
148	56	0.107143	0.0357143	0.196429	0.892857	0.767857
149	42	0.119048	0.047619	0.238095	0.880952	0.714286
150	54	0.12963	0.0550926	0.204167	0.87037	0.722222
151	43	0.0697674	0	0.162791	0.930233	0.767442
152	40	0.1	0.025	0.2	0.9	0.75
153	30	0.0333333	0	0.1	0.966667	0.833333
154	42	0.0714286	0	0.142857	0.928571	0.833333
155	29	0.0344828	0	0.10431	0.965517	0.758621
156	74	0.135135	0.0540541	0.216216	0.864865	0.716216
157	50	0.14	0.0595	0.24	0.86	0.74
158	80	0.1375	0.0625	0.200625	0.8625	0.725
159	48	0.208333	0.104167	0.3125	0.791667	0.625
160	51	0.156863	0.0588235	0.255392	0.843137	0.764706
161	71	0.15493	0.0704225	0.239789	0.84507	0.774648
162	48	0.145833	0.0619792	0.25	0.854167	0.708333
163	49	0.142857	0.0408163	0.245408	0.857143	0.816327
164	34	0.117647	0.0294118	0.206618	0.882353	0.705882
165	48	0.145833	0.0619792	0.25	0.854167	0.708333
166	50	0.18	0.08	0.3	0.82	0.78
167	70	0.114286	0.0567857	0.2	0.885714	0.728571

Row	n	beta	beta low	beta high	selection ceiling	correctness majority
168	48	0.104167	0.0208333	0.229167	0.895833	0.604167
169	48	0.125	0.0416667	0.208854	0.875	0.75
170	33	0.0909091	0	0.212121	0.909091	0.666667
171	48	0.1875	0.103646	0.3125	0.8125	0.666667
172	49	0.142857	0.0612245	0.245408	0.857143	0.795918
173	34	0.147059	0.0294118	0.294118	0.852941	0.735294
174	49	0.102041	0.0204082	0.204082	0.897959	0.693878
175	34	0.0882353	0	0.176471	0.911765	0.735294
176	33	0.0909091	0	0.181818	0.909091	0.727273
177	47	0.0851064	0.0207447	0.170213	0.914894	0.702128
178	33	0.0606061	0	0.151515	0.939394	0.757576
179	34	0.0882353	0	0.176471	0.911765	0.764706

Column panel 3/3 - row numbers identify the same source rows.

Row	plurality accuracy	independence beta
1	1	0
2	1	0
3	1	0
4	1	0
5	1	0
6	0.333333	0
7	0.980392	0.000763281
8	0.979167	0
9	0.977778	0.00037037
10	0.979167	0.000610352
11	1	0
12	1	0
13	1	0
14	1	0
15	0.988372	2.35828e-05

Row	plurality accuracy	independence beta
16	0.974684	6.08471e-05
17	0.976471	0.000122125
18	0.986667	1.42222e-05
19	0.988095	2.53077e-05
20	0.987013	6.57127e-05
21	1	0
22	1	0
23	1	0
24	1	0
25	1	0
26	1	0
27	1	0
28	1	0
29	1	0
30	1	0
31	0.333333	0
32	0.333333	0
33	0.333333	0
34	0.333333	0
35	1	0
36	0.966667	0.000555556
37	1	0.000915527
38	1	0
39	1	0
40	1	0.000410021
41	1	0
42	1	0
43	1	0
44	1	0
45	1	0

Row	plurality accuracy	independence beta
46	1	0
47	0.973684	2.27803e-05
48	0.975904	4.37226e-05
49	0.961039	0.000109521
50	0.986301	2.57058e-05
51	1	0
52	1	0
53	1	0
54	1	0
55	1	0
56	1	0
57	1	0
58	1	0
59	1	0
60	1	0
61	1	0
62	1	0
63	1	0
64	1	0
65	1	2.92524e-06
66	1	0
67	1	0
68	1	0
69	0.986111	7.75227e-08
70	1	0
71	1	0
72	1	0
73	1	0
74	1	0
75	1	0

Row	plurality accuracy	independence beta
76	1	0
77	1	0
78	1	0
79	1	0
80	1	0
81	1	0
82	1	0
83	1	0
84	0.987179	6.32175e-05
85	0.970149	3.98985e-05
86	0.972603	0.000102823
87	0.933333	0.001728
88	0.974026	8.76169e-06
89	0.974684	3.24518e-05
90	0.904762	0.00145773
91	0.972603	7.71175e-05
92	0.888889	0.00115741
93	0.901235	0.000812884
94	1	0
95	1	0
96	1	0
97	1	0
98	1	0
99	1	0
100	1	0
101	1	0
102	1	0
103	1	0
104	0.984848	0
105	0.986111	2.67918e-05

Row	plurality accuracy	independence beta
106	0.96	0.001152
107	0.984375	2.28882e-05
108	0.947368	0
109	0.954545	0
110	0.985507	6.08811e-06
111	0.904762	0.000431919
112	0.92	0.000512
113	0.954545	0.000375657
114	1	0
115	1	0
116	1	0
117	1	0
118	1	0
119	1	0
120	1	0
121	1	0
122	1	0
123	0.984127	0
124	0.947368	0
125	0.954545	0
126	0.9	5e-06
127	1	0
128	1	0
129	1	0
130	1	0
131	0.944444	0
132	0.740196	0.0276817
133	0.772727	0.0167637
134	0.768519	0.0267249
135	0.777778	0.0323129

Row	plurality accuracy	independence beta
136	0.818713	0.0163289
137	0.769547	0.0165588
138	0.746914	0.0300513
139	0.827586	0.0124595
140	0.767442	0.0211302
141	0.797619	0.0211712
142	0.738095	0.0275348
143	0.744792	0.0274658
144	0.741667	0.050625
145	0.782946	0.0199228
146	0.733333	0.0342222
147	0.785714	0.0323129
148	0.779762	0.0163596
149	0.722222	0.0283447
150	0.740741	0.025428
151	0.782946	0.0158477
152	0.8	0.00155039
153	0.866667	0.000342222
154	0.845238	0.000771284
155	0.827586	0.000946022
156	0.738739	0.0263854
157	0.773333	0.024336
158	0.75	0.0232031
159	0.666667	0.0645616
160	0.777778	0.0165623
161	0.793427	0.0186862
162	0.722222	0.0341797
163	0.829932	0.0159117
164	0.72549	0.0366375
165	0.743056	0.0341797

Row	plurality accuracy	independence beta
166	0.786667	0.024336
167	0.747619	0.0223907
168	0.659722	0.0430411
169	0.763889	0.0208333
170	0.727273	0.0400701
171	0.694444	0.0473452
172	0.802721	0.0146878
173	0.745098	0.0325667
174	0.727891	0.0261796
175	0.754902	0.024425
176	0.80303	0.00264926
177	0.776596	0.00263708
178	0.80303	0.00198695
179	0.808824	0.00152128

repeat_pairs.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	model	benchmark	format	condition a	condition b	n
1	QWE03	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	95
2	QWE03	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	99
3	QWE01	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	1
4	QWE02	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	29
5	QWE02	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	80
6	QWE07	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	90
7	QWE07	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	75
8	QWE06	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	78
9	QWE06	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	47
10	QWE09	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	79
11	QWE09	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	81

Row	model	benchmark	format	condition a	condition b	n
12	QWE08	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	30
13	QWE08	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	39

Column panel 2/3 - row numbers identify the same source rows.

Row	error a	error b	double fault	excess cofailure	phi	adjusted phi
1	0.0315789	0.0526316	0.0210526	0.0193906	0.496566	0.648148
2	0.262626	0.30303	0.222222	0.142639	0.7053	0.779264
3	0	0	0	0	NA	NA
4	0.172414	0	0	0	NA	NA
5	0.5125	0.4	0.325	0.12	0.490051	0.615385
6	0.0444444	0.0222222	0.0222222	0.0212346	0.699025	1
7	0.266667	0.24	0.226667	0.162667	0.861293	0.924242
8	0.0128205	0.025641	0.0128205	0.0124918	0.7025	1
9	0.297872	0.212766	0.191489	0.128112	0.684486	0.857576
10	0.0253165	0.0379747	0.0253165	0.0243551	0.811177	1
11	0.259259	0.246914	0.160494	0.0964792	0.51055	0.5275
12	0.0333333	0.133333	0.0333333	0.0288889	0.473432	1
13	0.461538	0.538462	0.435897	0.187377	0.753968	0.87963

Column panel 3/3 - row numbers identify the same source rows.

Row	kappa	wrong agreement	wrong chance uniform	both wrong n	capa onehot
1	0.479452	0.5	NA	2	NA
2	0.701807	0.818182	0.111742	22	0.660737
3	NA	NA	NA	0	NA
4	0	NA	NA	0	NA
5	0.477612	0.461538	0.11569	26	0.358707
6	0.656489	1	NA	2	NA
7	0.859155	0.823529	0.113796	17	0.785266
8	0.66087	1	NA	1	NA
9	0.667453	0.888889	0.13933	9	0.66108

Row	kappa	wrong agreement	wrong chance uniform	both wrong n	capa onehot
10	0.793734	1	NA	2	NA
11	0.510278	0.846154	0.114621	13	0.516668
12	0.366197	1	NA	1	NA
13	0.745098	0.823529	0.14099	17	0.714577

self_consistency.csv

No rows were recorded in this output.

baseline_predictions.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	condition	team	method	train items	test items	observed beta
1	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE01", "QWE02"]	independence	50	2	0
2	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE01", "QWE02"]	single_factor_tetrachoric	50	2	0
3	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE01", "QWE07"]	independence	50	4	0
4	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE01", "QWE07"]	single_factor_tetrachoric	50	4	0
5	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE01", "QWE06"]	independence	50	4	0
6	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE01", "QWE06"]	single_factor_tetrachoric	50	4	0
7	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE01", "QWE09"]	independence	50	4	0
8	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE01", "QWE09"]	single_factor_tetrachoric	50	4	0
9	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE01", "QWE08"]	independence	50	4	0
10	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE01", "QWE08"]	single_factor_tetrachoric	50	4	0
11	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE07"]	independence	50	18	0
12	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE07"]	single_factor_tetrachoric	50	18	0
13	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE06"]	independence	50	17	0
14	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE06"]	single_factor_tetrachoric	50	17	0
15	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE09"]	independence	50	17	0
16	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE09"]	single_factor_tetrachoric	50	17	0
17	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE08"]	independence	50	16	0
18	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE08"]	single_factor_tetrachoric	50	16	0

Row	condition	team	method	train items	test items	observed beta
19	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE07", "QWE06"]	independence	50	42	0
20	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE07", "QWE06"]	single_factor_tetrachoric	50	42	0
21	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE07", "QWE09"]	independence	50	42	0
22	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE07", "QWE09"]	single_factor_tetrachoric	50	42	0
23	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE07", "QWE08"]	independence	50	42	0
24	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE07", "QWE08"]	single_factor_tetrachoric	50	42	0
25	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE06", "QWE09"]	independence	50	39	0
26	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE06", "QWE09"]	single_factor_tetrachoric	50	39	0
27	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE06", "QWE08"]	independence	50	41	0
28	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE06", "QWE08"]	single_factor_tetrachoric	50	41	0
29	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE09", "QWE08"]	independence	50	40	0
30	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE03", "QWE09", "QWE08"]	single_factor_tetrachoric	50	40	0
31	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE02", "QWE07"]	independence	50	2	0
32	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE02", "QWE07"]	single_factor_tetrachoric	50	2	0
33	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE02", "QWE06"]	independence	50	2	0
34	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE02", "QWE06"]	single_factor_tetrachoric	50	2	0
35	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE02", "QWE09"]	independence	50	2	0
36	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE02", "QWE09"]	single_factor_tetrachoric	50	2	0
37	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE02", "QWE08"]	independence	50	2	0
38	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE02", "QWE08"]	single_factor_tetrachoric	50	2	0
39	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE07", "QWE06"]	independence	50	4	0
40	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE07", "QWE06"]	single_factor_tetrachoric	50	4	0
41	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE07", "QWE09"]	independence	50	4	0
42	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE07", "QWE09"]	single_factor_tetrachoric	50	4	0
43	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE07", "QWE08"]	independence	50	4	0
44	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE07", "QWE08"]	single_factor_tetrachoric	50	4	0
45	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE06", "QWE09"]	independence	50	4	0
46	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE06", "QWE09"]	single_factor_tetrachoric	50	4	0
47	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE06", "QWE08"]	independence	50	4	0
48	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE06", "QWE08"]	single_factor_tetrachoric	50	4	0

Row	condition	team	method	train items	test items	observed beta
49	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE09", "QWE08"]	independence	50	4	0
50	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE01", "QWE09", "QWE08"]	single_factor_tetrachoric	50	4	0
51	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE02", "QWE07", "QWE06"]	independence	50	17	0
52	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE02", "QWE07", "QWE06"]	single_factor_tetrachoric	50	17	0
53	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE02", "QWE07", "QWE09"]	independence	50	17	0
54	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE02", "QWE07", "QWE09"]	single_factor_tetrachoric	50	17	0
55	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE02", "QWE07", "QWE08"]	independence	50	16	0
56	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE02", "QWE07", "QWE08"]	single_factor_tetrachoric	50	16	0
57	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE02", "QWE06", "QWE09"]	independence	50	16	0
58	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE02", "QWE06", "QWE09"]	single_factor_tetrachoric	50	16	0
59	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE02", "QWE06", "QWE08"]	independence	50	15	0
60	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE02", "QWE06", "QWE08"]	single_factor_tetrachoric	50	15	0
61	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE02", "QWE09", "QWE08"]	independence	50	16	0
62	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE02", "QWE09", "QWE08"]	single_factor_tetrachoric	50	16	0
63	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE07", "QWE06", "QWE09"]	independence	50	39	0
64	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE07", "QWE06", "QWE09"]	single_factor_tetrachoric	50	39	0
65	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE07", "QWE06", "QWE08"]	independence	50	39	0
66	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE07", "QWE06", "QWE08"]	single_factor_tetrachoric	50	39	0
67	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE07", "QWE09", "QWE08"]	independence	50	40	0
68	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE07", "QWE09", "QWE08"]	single_factor_tetrachoric	50	40	0
69	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE06", "QWE09", "QWE08"]	independence	50	37	0
70	["gsm8k", "numeric", "direct", 17, 0.0]	["QWE06", "QWE09", "QWE08"]	single_factor_tetrachoric	50	37	0
71	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE01", "QWE02"]	independence	50	2	0
72	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE01", "QWE02"]	single_factor_tetrachoric	50	2	0
73	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE01", "QWE07"]	independence	50	3	0
74	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE01", "QWE07"]	single_factor_tetrachoric	50	3	0
75	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE01", "QWE06"]	independence	50	4	0

Row	condition	team	method	train items	test items	observed beta
76	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE01", "QWE06"]	single_factor_tetrachoric	50	4	0
77	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE01", "QWE09"]	independence	50	5	0
78	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE01", "QWE09"]	single_factor_tetrachoric	50	5	0
79	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE01", "QWE08"]	independence	50	4	0
80	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE01", "QWE08"]	single_factor_tetrachoric	50	4	0
81	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE07"]	independence	50	38	0.0263158
82	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE07"]	single_factor_tetrachoric	50	38	0.0263158
83	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE06"]	independence	50	31	0
84	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE06"]	single_factor_tetrachoric	50	31	0
85	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE09"]	independence	50	35	0
86	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE09"]	single_factor_tetrachoric	50	35	0
87	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE08"]	independence	50	13	0.0769231
88	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE08"]	single_factor_tetrachoric	50	13	0.0769231
89	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE07", "QWE06"]	independence	50	36	0
90	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE07", "QWE06"]	single_factor_tetrachoric	50	36	0
91	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE07", "QWE09"]	independence	50	40	0
92	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE07", "QWE09"]	single_factor_tetrachoric	50	40	0
93	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE07", "QWE08"]	independence	50	16	0.0625
94	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE07", "QWE08"]	single_factor_tetrachoric	50	16	0.0625
95	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE06", "QWE09"]	independence	50	38	0.0263158

Row	condition	team	method	train items	test items	observed beta
96	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE06", "QWE09"]	single_factor_tetrachoric	50	38	0.0263158
97	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE06", "QWE08"]	independence	50	13	0.0769231
98	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE06", "QWE08"]	single_factor_tetrachoric	50	13	0.0769231
99	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE09", "QWE08"]	independence	50	16	0.0625
100	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE03", "QWE09", "QWE08"]	single_factor_tetrachoric	50	16	0.0625
101	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE02", "QWE07"]	independence	50	2	0
102	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE02", "QWE07"]	single_factor_tetrachoric	50	2	0
103	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE02", "QWE06"]	independence	50	2	0
104	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE02", "QWE06"]	single_factor_tetrachoric	50	2	0
105	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE02", "QWE09"]	independence	50	2	0
106	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE02", "QWE09"]	single_factor_tetrachoric	50	2	0
107	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE02", "QWE08"]	independence	50	2	0
108	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE02", "QWE08"]	single_factor_tetrachoric	50	2	0
109	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE07", "QWE06"]	independence	50	2	0
110	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE07", "QWE06"]	single_factor_tetrachoric	50	2	0
111	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE07", "QWE09"]	independence	50	3	0
112	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE07", "QWE09"]	single_factor_tetrachoric	50	3	0
113	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE07", "QWE08"]	independence	50	3	0
114	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE07", "QWE08"]	single_factor_tetrachoric	50	3	0
115	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE06", "QWE09"]	independence	50	4	0

Row	condition	team	method	train items	test items	observed beta
116	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE06", "QWE09"]	single_factor_tetrachoric	50	4	0
117	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE06", "QWE08"]	independence	50	3	0
118	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE06", "QWE08"]	single_factor_tetrachoric	50	3	0
119	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE09", "QWE08"]	independence	50	4	0
120	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE01", "QWE09", "QWE08"]	single_factor_tetrachoric	50	4	0
121	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE02", "QWE07", "QWE06"]	independence	50	31	0
122	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE02", "QWE07", "QWE06"]	single_factor_tetrachoric	50	31	0
123	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE02", "QWE07", "QWE09"]	independence	50	35	0
124	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE02", "QWE07", "QWE09"]	single_factor_tetrachoric	50	35	0
125	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE02", "QWE07", "QWE08"]	independence	50	13	0.0769231
126	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE02", "QWE07", "QWE08"]	single_factor_tetrachoric	50	13	0.0769231
127	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE02", "QWE06", "QWE09"]	independence	50	31	0
128	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE02", "QWE06", "QWE09"]	single_factor_tetrachoric	50	31	0
129	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE02", "QWE06", "QWE08"]	independence	50	9	0
130	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE02", "QWE06", "QWE08"]	single_factor_tetrachoric	50	9	0
131	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE02", "QWE09", "QWE08"]	independence	50	11	0
132	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE02", "QWE09", "QWE08"]	single_factor_tetrachoric	50	11	0
133	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE07", "QWE06", "QWE09"]	independence	50	35	0
134	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE07", "QWE06", "QWE09"]	single_factor_tetrachoric	50	35	0
135	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE07", "QWE06", "QWE08"]	independence	50	11	0

Row	condition	team	method	train items	test items	observed beta
136	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE07", "QWE06", "QWE08"]	single_factor_tetrachoric	50	11	0
137	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE07", "QWE09", "QWE08"]	independence	50	14	0
138	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE07", "QWE09", "QWE08"]	single_factor_tetrachoric	50	14	0
139	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE06", "QWE09", "QWE08"]	independence	50	13	0.0769231
140	["gsm8k", "numeric", "reasoned", 17, 0.0]	["QWE06", "QWE09", "QWE08"]	single_factor_tetrachoric	50	13	0.0769231
141	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE07"]	independence	50	35	0.2
142	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE07"]	single_factor_tetrachoric	50	35	0.2
143	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE06"]	independence	50	26	0.115385
144	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE06"]	single_factor_tetrachoric	50	26	0.115385
145	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE09"]	independence	50	35	0.2
146	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE09"]	single_factor_tetrachoric	50	35	0.2
147	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE08"]	independence	50	20	0.25
148	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE02", "QWE08"]	single_factor_tetrachoric	50	20	0.25
149	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE07", "QWE06"]	independence	50	35	0.142857
150	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE07", "QWE06"]	single_factor_tetrachoric	50	35	0.142857
151	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE07", "QWE09"]	independence	50	43	0.162791
152	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE07", "QWE09"]	single_factor_tetrachoric	50	43	0.162791
153	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE07", "QWE08"]	independence	50	29	0.172414
154	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE07", "QWE08"]	single_factor_tetrachoric	50	29	0.172414
155	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE06", "QWE09"]	independence	50	35	0.0857143
156	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE06", "QWE09"]	single_factor_tetrachoric	50	35	0.0857143
157	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE06", "QWE08"]	independence	50	24	0.125
158	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE06", "QWE08"]	single_factor_tetrachoric	50	24	0.125
159	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE09", "QWE08"]	independence	50	29	0.172414
160	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE03", "QWE09", "QWE08"]	single_factor_tetrachoric	50	29	0.172414
161	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE02", "QWE07", "QWE06"]	independence	50	25	0.16
162	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE02", "QWE07", "QWE06"]	single_factor_tetrachoric	50	25	0.16

Row	condition	team	method	train items	test items	observed beta
163	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE02", "QWE07", "QWE09"]	independence	50	33	0.151515
164	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE02", "QWE07", "QWE09"]	single_factor_tetrachoric	50	33	0.151515
165	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE02", "QWE07", "QWE08"]	independence	50	20	0.3
166	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE02", "QWE07", "QWE08"]	single_factor_tetrachoric	50	20	0.3
167	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE02", "QWE06", "QWE09"]	independence	50	25	0.08
168	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE02", "QWE06", "QWE09"]	single_factor_tetrachoric	50	25	0.08
169	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE02", "QWE06", "QWE08"]	independence	50	15	0.266667
170	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE02", "QWE06", "QWE08"]	single_factor_tetrachoric	50	15	0.266667
171	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE02", "QWE09", "QWE08"]	independence	50	20	0.2
172	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE02", "QWE09", "QWE08"]	single_factor_tetrachoric	50	20	0.2
173	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE07", "QWE06", "QWE09"]	independence	50	34	0.0882353
174	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE07", "QWE06", "QWE09"]	single_factor_tetrachoric	50	34	0.0882353
175	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE07", "QWE06", "QWE08"]	independence	50	24	0.166667
176	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE07", "QWE06", "QWE08"]	single_factor_tetrachoric	50	24	0.166667
177	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE07", "QWE09", "QWE08"]	independence	50	29	0.137931
178	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE07", "QWE09", "QWE08"]	single_factor_tetrachoric	50	29	0.137931
179	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE06", "QWE09", "QWE08"]	independence	50	24	0.0833333
180	["mmlu_pro", "mc", "direct", 17, 0.0]	["QWE06", "QWE09", "QWE08"]	single_factor_tetrachoric	50	24	0.0833333
181	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE07"]	independence	50	39	0.0512821
182	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE07"]	single_factor_tetrachoric	50	39	0.0512821
183	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE06"]	independence	50	27	0.0740741
184	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE06"]	single_factor_tetrachoric	50	27	0.0740741
185	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE09"]	independence	50	44	0.113636
186	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE09"]	single_factor_tetrachoric	50	44	0.113636
187	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE08"]	independence	50	28	0.107143
188	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE02", "QWE08"]	single_factor_tetrachoric	50	28	0.107143

Row	condition	team	method	train items	test items	observed beta
189	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE07", "QWE06"]	independence	50	27	0.0740741
190	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE07", "QWE06"]	single_factor_tetrachoric	50	27	0.0740741
191	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE07", "QWE09"]	independence	50	39	0.102564
192	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE07", "QWE09"]	single_factor_tetrachoric	50	39	0.102564
193	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE07", "QWE08"]	independence	50	28	0
194	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE07", "QWE08"]	single_factor_tetrachoric	50	28	0
195	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE06", "QWE09"]	independence	50	27	0.0740741
196	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE06", "QWE09"]	single_factor_tetrachoric	50	27	0.0740741
197	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE06", "QWE08"]	independence	50	20	0
198	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE06", "QWE08"]	single_factor_tetrachoric	50	20	0
199	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE09", "QWE08"]	independence	50	28	0.0714286
200	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE03", "QWE09", "QWE08"]	single_factor_tetrachoric	50	28	0.0714286
201	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE02", "QWE07", "QWE06"]	independence	50	27	0.111111
202	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE02", "QWE07", "QWE06"]	single_factor_tetrachoric	50	27	0.111111
203	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE02", "QWE07", "QWE09"]	independence	50	38	0.0526316
204	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE02", "QWE07", "QWE09"]	single_factor_tetrachoric	50	38	0.0526316
205	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE02", "QWE07", "QWE08"]	independence	50	28	0
206	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE02", "QWE07", "QWE08"]	single_factor_tetrachoric	50	28	0
207	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE02", "QWE06", "QWE09"]	independence	50	27	0.0740741
208	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE02", "QWE06", "QWE09"]	single_factor_tetrachoric	50	27	0.0740741

Row	condition	team	method	train items	test items	observed beta
209	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE02", "QWE06", "QWE08"]	independence	50	20	0
210	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE02", "QWE06", "QWE08"]	single_factor_tetrachoric	50	20	0
211	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE02", "QWE09", "QWE08"]	independence	50	28	0.107143
212	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE02", "QWE09", "QWE08"]	single_factor_tetrachoric	50	28	0.107143
213	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE07", "QWE06", "QWE09"]	independence	50	27	0.0740741
214	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE07", "QWE06", "QWE09"]	single_factor_tetrachoric	50	27	0.0740741
215	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE07", "QWE06", "QWE08"]	independence	50	20	0.05
216	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE07", "QWE06", "QWE08"]	single_factor_tetrachoric	50	20	0.05
217	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE07", "QWE09", "QWE08"]	independence	50	28	0
218	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE07", "QWE09", "QWE08"]	single_factor_tetrachoric	50	28	0
219	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE06", "QWE09", "QWE08"]	independence	50	20	0
220	["mmlu_pro", "mc", "reasoned", 17, 0.0]	["QWE06", "QWE09", "QWE08"]	single_factor_tetrachoric	50	20	0

Column panel 2/3 - row numbers identify the same source rows.

Row	beta low	beta high	predicted beta	beta mae	observed majority	predicted majority
1	0	0	1.06383e-06	1.06383e-06	1	0.989335
2	0	0	1.06383e-06	1.06383e-06	1	0.993
3	0	0	1.81077e-07	1.81077e-07	1	0.998181
4	0	0	1.81077e-07	1.81077e-07	1	0.996
5	0	0	9.05387e-08	9.05387e-08	1	0.999088
6	0	0	9.05387e-08	9.05387e-08	1	0.999
7	0	0	2.18221e-07	2.18221e-07	1	0.997809
8	0	0	2.18221e-07	2.18221e-07	1	0.999
9	0	0	2.71616e-07	2.71616e-07	1	0.997274

Row	beta low	beta high	predicted beta	beta mae	observed majority	predicted majority
10	0	0	2.71616e-07	2.71616e-07	1	0.999
11	0	0	0.000452694	0.000452694	1	0.977818
12	0	0	0.000452694	0.000452694	1	0.976
13	0	0	0.000226347	0.000226347	1	0.98359
14	0	0	0.000226347	0.000226347	1	0.9845
15	0	0	0.000545554	0.000545554	1	0.97545
16	0	0	0.000545554	0.000545554	1	0.973
17	0	0	0.00067904	0.00067904	1	0.972046
18	0	0	0.00067904	0.00067904	1	0.9695
19	0	0	3.85271e-05	3.85271e-05	1	0.996456
20	0	0	3.85271e-05	3.85271e-05	1	0.996
21	0	0	9.28602e-05	9.28602e-05	1	0.994011
22	0	0	9.28602e-05	9.28602e-05	1	0.995
23	0	0	0.000115581	0.000115581	0.97619	0.992988
24	0	0	0.000115581	0.000115581	0.97619	0.995
25	0	0	4.64301e-05	4.64301e-05	1	0.995914
26	0	0	4.64301e-05	4.64301e-05	1	0.9965
27	0	0	5.77907e-05	5.77907e-05	1	0.995136
28	0	0	5.77907e-05	5.77907e-05	1	0.995
29	0	0	0.00013929	0.00013929	1	0.992107
30	0	0	0.00013929	0.00013929	1	0.993
31	0	0	1.06383e-06	1.06383e-06	1	0.989335
32	0	0	1.06383e-06	1.06383e-06	1	0.989
33	0	0	5.31915e-07	5.31915e-07	1	0.994655
34	0	0	5.31915e-07	5.31915e-07	1	0.994
35	0	0	1.28205e-06	1.28205e-06	1	0.987152
36	0	0	1.28205e-06	1.28205e-06	1	0.989
37	0	0	1.59574e-06	1.59574e-06	1	0.984014
38	0	0	1.59574e-06	1.59574e-06	1	0.9845
39	0	0	9.05387e-08	9.05387e-08	1	0.999088

Row	beta low	beta high	predicted beta	beta mae	observed majority	predicted majority
40	0	0	9.05387e-08	9.05387e-08	1	0.9985
41	0	0	2.18221e-07	2.18221e-07	1	0.997809
42	0	0	2.18221e-07	2.18221e-07	1	0.9985
43	0	0	2.71616e-07	2.71616e-07	1	0.997274
44	0	0	2.71616e-07	2.71616e-07	1	0.9975
45	0	0	1.09111e-07	1.09111e-07	1	0.998902
46	0	0	1.09111e-07	1.09111e-07	1	0.998
47	0	0	1.35808e-07	1.35808e-07	1	0.998634
48	0	0	1.35808e-07	1.35808e-07	1	0.9995
49	0	0	3.27332e-07	3.27332e-07	1	0.996716
50	0	0	3.27332e-07	3.27332e-07	1	0.9975
51	0	0	0.000226347	0.000226347	1	0.98359
52	0	0	0.000226347	0.000226347	1	0.983
53	0	0	0.000545554	0.000545554	1	0.97545
54	0	0	0.000545554	0.000545554	1	0.973
55	0	0	0.00067904	0.00067904	1	0.972046
56	0	0	0.00067904	0.00067904	1	0.973
57	0	0	0.000272777	0.000272777	1	0.981315
58	0	0	0.000272777	0.000272777	1	0.979
59	0	0	0.00033952	0.00033952	1	0.978044
60	0	0	0.00033952	0.00033952	1	0.975
61	0	0	0.000818331	0.000818331	1	0.969585
62	0	0	0.000818331	0.000818331	1	0.9725
63	0	0	4.64301e-05	4.64301e-05	1	0.995914
64	0	0	4.64301e-05	4.64301e-05	1	0.9965
65	0	0	5.77907e-05	5.77907e-05	0.974359	0.995136
66	0	0	5.77907e-05	5.77907e-05	0.974359	0.996
67	0	0	0.00013929	0.00013929	0.975	0.992107
68	0	0	0.00013929	0.00013929	0.975	0.9925
69	0	0	6.96452e-05	6.96452e-05	1	0.994417

Row	beta low	beta high	predicted beta	beta mae	observed majority	predicted majority
70	0	0	6.96452e-05	6.96452e-05	1	0.9945
71	0	0	9.9552e-08	9.9552e-08	1	0.998998
72	0	0	9.9552e-08	9.9552e-08	1	0.999
73	0	0	4.43656e-08	4.43656e-08	1	0.999552
74	0	0	4.43656e-08	4.43656e-08	1	1
75	0	0	4.74608e-08	4.74608e-08	1	0.999521
76	0	0	4.74608e-08	4.74608e-08	1	0.999
77	0	0	1.02041e-07	1.02041e-07	1	0.998973
78	0	0	1.02041e-07	1.02041e-07	1	0.9975
79	0	0	2.04082e-10	2.04082e-10	1	0.999996
80	0	0	2.04082e-10	2.04082e-10	1	1
81	0	0.0789474	2.16417e-05	0.0262941	0.973684	0.997544
82	0	0.0789474	2.16417e-05	0.0262941	0.973684	0.996
83	0	0	2.31516e-05	2.31516e-05	0.967742	0.997442
84	0	0	2.31516e-05	2.31516e-05	0.967742	0.999
85	0	0	4.9776e-05	4.9776e-05	0.971429	0.995645
86	0	0	4.9776e-05	4.9776e-05	0.971429	0.9965
87	0	0.230769	9.9552e-08	0.076923	0.846154	0.998998
88	0	0.230769	9.9552e-08	0.076923	0.846154	1
89	0	0	1.03176e-05	1.03176e-05	0.944444	0.998597
90	0	0	1.03176e-05	1.03176e-05	0.944444	0.9985
91	0	0	2.21828e-05	2.21828e-05	0.95	0.997493
92	0	0	2.21828e-05	2.21828e-05	0.95	0.999
93	0	0.25	4.43656e-08	0.0625	0.8125	0.999552
94	0	0.25	4.43656e-08	0.0625	0.8125	1
95	0	0.0789474	2.37304e-05	0.0262921	0.973684	0.99739
96	0	0.0789474	2.37304e-05	0.0262921	0.973684	0.997
97	0	0.230769	4.74608e-08	0.076923	0.769231	0.999521
98	0	0.230769	4.74608e-08	0.076923	0.769231	0.9995
99	0	0.1875	1.02041e-07	0.0624999	0.8125	0.998973

Row	beta low	beta high	predicted beta	beta mae	observed majority	predicted majority
100	0	0.1875	1.02041e-07	0.0624999	0.8125	0.999
101	0	0	1.06045e-07	1.06045e-07	1	0.998933
102	0	0	1.06045e-07	1.06045e-07	1	0.9995
103	0	0	1.13443e-07	1.13443e-07	1	0.998859
104	0	0	1.13443e-07	1.13443e-07	1	0.9985
105	0	0	2.43902e-07	2.43902e-07	1	0.997552
106	0	0	2.43902e-07	2.43902e-07	1	0.9975
107	0	0	4.87805e-10	4.87805e-10	1	0.99999
108	0	0	4.87805e-10	4.87805e-10	1	1
109	0	0	5.05561e-08	5.05561e-08	1	0.99949
110	0	0	5.05561e-08	5.05561e-08	1	1
111	0	0	1.08696e-07	1.08696e-07	1	0.998906
112	0	0	1.08696e-07	1.08696e-07	1	0.9995
113	0	0	2.17391e-10	2.17391e-10	1	0.999996
114	0	0	2.17391e-10	2.17391e-10	1	1
115	0	0	1.16279e-07	1.16279e-07	1	0.99883
116	0	0	1.16279e-07	1.16279e-07	1	1
117	0	0	2.32558e-10	2.32558e-10	1	0.999995
118	0	0	2.32558e-10	2.32558e-10	1	1
119	0	0	5e-10	5e-10	1	0.99999
120	0	0	5e-10	5e-10	1	1
121	0	0	2.46615e-05	2.46615e-05	0.967742	0.997349
122	0	0	2.46615e-05	2.46615e-05	0.967742	0.9985
123	0	0	5.30223e-05	5.30223e-05	0.971429	0.99552
124	0	0	5.30223e-05	5.30223e-05	0.971429	0.9935
125	0	0.230769	1.06045e-07	0.076923	0.923077	0.998933
126	0	0.230769	1.06045e-07	0.076923	0.923077	0.9995
127	0	0	5.67215e-05	5.67215e-05	1	0.995377
128	0	0	5.67215e-05	5.67215e-05	1	0.9965
129	0	0	1.13443e-07	1.13443e-07	0.888889	0.998859

Row	beta low	beta high	predicted beta	beta mae	observed majority	predicted majority
130	0	0	1.13443e-07	1.13443e-07	0.888889	0.999
131	0	0	2.43902e-07	2.43902e-07	0.909091	0.997552
132	0	0	2.43902e-07	2.43902e-07	0.909091	0.9975
133	0	0	2.52781e-05	2.52781e-05	0.971429	0.997295
134	0	0	2.52781e-05	2.52781e-05	0.971429	0.995
135	0	0	5.05561e-08	5.05561e-08	0.818182	0.99949
136	0	0	5.05561e-08	5.05561e-08	0.818182	0.9995
137	0	0	1.08696e-07	1.08696e-07	0.857143	0.998906
138	0	0	1.08696e-07	1.08696e-07	0.857143	0.9985
139	0	0.230769	1.16279e-07	0.076923	0.923077	0.99883
140	0	0.230769	1.16279e-07	0.076923	0.923077	0.998
141	0.0857143	0.342857	0.0194419	0.180558	0.657143	0.809302
142	0.0857143	0.342857	0.0864266	0.113573	0.657143	0.755
143	0	0.269231	0.0253589	0.0900257	0.615385	0.780849
144	0	0.269231	0.0982805	0.0171041	0.615385	0.739
145	0.085	0.314286	0.0265116	0.173488	0.714286	0.775307
146	0.085	0.314286	0.10353	0.0964696	0.714286	0.725
147	0.09875	0.45	0.0360034	0.213997	0.55	0.729664
148	0.09875	0.45	0.126341	0.123659	0.55	0.6825
149	0.0285714	0.257143	0.0114783	0.131379	0.771429	0.869391
150	0.0285714	0.257143	0.0671281	0.075729	0.771429	0.801
151	0.0697674	0.27907	0.012	0.150791	0.72093	0.865455
152	0.0697674	0.27907	0.0734824	0.0893083	0.72093	0.812
153	0.0689655	0.310345	0.0162963	0.156117	0.655172	0.833037
154	0.0689655	0.310345	0.082103	0.0903108	0.655172	0.7735
155	0	0.172143	0.0156522	0.0700621	0.771429	0.842767
156	0	0.172143	0.0777458	0.00796852	0.771429	0.775
157	0.040625	0.25	0.021256	0.103744	0.625	0.807021
158	0.040625	0.25	0.0904342	0.0345658	0.625	0.7655
159	0.0344828	0.344828	0.0222222	0.150192	0.724138	0.801953

Row	beta low	beta high	predicted beta	beta mae	observed majority	predicted majority
160	0.0344828	0.344828	0.0860133	0.0864005	0.724138	0.7645
161	0.04	0.32	0.0230536	0.136946	0.6	0.790293
162	0.04	0.32	0.0972211	0.0627789	0.6	0.7275
163	0.030303	0.272727	0.0241015	0.127414	0.636364	0.784778
164	0.030303	0.272727	0.099142	0.0523732	0.636364	0.7335
165	0.1	0.5	0.0327304	0.26727	0.55	0.739363
166	0.1	0.5	0.112515	0.187485	0.55	0.6955
167	0	0.2	0.0314367	0.0485633	0.64	0.755952
168	0	0.2	0.121589	0.0415886	0.64	0.7095
169	0.0666667	0.533333	0.0426918	0.223975	0.533333	0.709845
170	0.0666667	0.533333	0.137259	0.129408	0.533333	0.6645
171	0.05	0.40125	0.0446324	0.155368	0.55	0.704095
172	0.05	0.40125	0.144958	0.0550423	0.55	0.6525
173	0	0.176471	0.0142292	0.074006	0.735294	0.850593
174	0	0.176471	0.0790281	0.00920716	0.735294	0.7795
175	0.0416667	0.333333	0.0193237	0.147343	0.625	0.815781
176	0.0416667	0.333333	0.0870916	0.079575	0.625	0.755
177	0.0344828	0.275862	0.020202	0.117729	0.655172	0.810774
178	0.0344828	0.275862	0.0924774	0.0454536	0.655172	0.751
179	0	0.208333	0.0263505	0.0569829	0.666667	0.783926
180	0	0.208333	0.100724	0.0173907	0.666667	0.743
181	0	0.128205	0.0480327	0.00324932	0.769231	0.695475
182	0	0.128205	0.178655	0.127373	0.769231	0.6565
183	0	0.185185	0.0538549	0.0202192	0.777778	0.678146
184	0	0.185185	0.209373	0.135299	0.777778	0.626
185	0.0454545	0.205114	0.0510204	0.062616	0.75	0.686583
186	0.0454545	0.205114	0.181241	0.0676046	0.75	0.656
187	0.0348214	0.215179	0.110158	0.00301484	0.75	0.510571
188	0.0348214	0.215179	0.264561	0.157419	0.75	0.525
189	0	0.185185	0.0404486	0.0336255	0.851852	0.724398

Row	beta low	beta high	predicted beta	beta mae	observed majority	predicted majority
190	0	0.185185	0.177706	0.103632	0.851852	0.661
191	0.025641	0.205128	0.0383197	0.0642444	0.846154	0.732517
192	0.025641	0.205128	0.163985	0.061421	0.846154	0.679
193	0	0	0.0827358	0.0827358	0.821429	0.56313
194	0	0	0.218303	0.218303	0.821429	0.5665
195	0	0.185185	0.0429646	0.0311095	0.888889	0.715718
196	0	0.185185	0.179572	0.105498	0.888889	0.6635
197	0	0	0.0927644	0.0927644	0.85	0.543908
198	0	0	0.238136	0.238136	0.85	0.5555
199	0	0.178571	0.087882	0.0164535	0.857143	0.553266
200	0	0.178571	0.226069	0.154641	0.857143	0.5665
201	0	0.259259	0.0392267	0.0718844	0.851852	0.72973
202	0	0.259259	0.167812	0.0567004	0.851852	0.6835
203	0	0.131579	0.0371622	0.0154694	0.763158	0.737761
204	0	0.131579	0.155074	0.102442	0.763158	0.694
205	0	0	0.0802365	0.0802365	0.714286	0.570204
206	0	0	0.227202	0.227202	0.714286	0.5485
207	0	0.185185	0.0416667	0.0324074	0.777778	0.721126
208	0	0.185185	0.171146	0.0970716	0.777778	0.6785
209	0	0	0.0899621	0.0899621	0.75	0.550821
210	0	0	0.234245	0.234245	0.75	0.5635
211	0	0.214286	0.0852273	0.0219156	0.785714	0.560257
212	0	0.214286	0.235761	0.128618	0.785714	0.5615
213	0	0.185185	0.0312945	0.0427796	0.851852	0.764343
214	0	0.185185	0.15764	0.0835661	0.851852	0.6885
215	0	0.15	0.0675676	0.0175676	0.9	0.606061
216	0	0.15	0.19807	0.14807	0.9	0.609
217	0	0	0.0640114	0.0640114	0.821429	0.616126
218	0	0	0.183677	0.183677	0.821429	0.6115
219	0	0	0.0717703	0.0717703	0.85	0.595694

Row	beta low	beta high	predicted beta	beta mae	observed majority	predicted majority
220	0	0	0.203475	0.203475	0.85	0.5885

Column panel 3/3 - row numbers identify the same source rows.

Row	majority mae	copula rho	effective voters	policy
1	0.0106654	0	NA	item holdout; capabilities calibrated on training items
2	0.007	0	NA	item holdout; capabilities calibrated on training items
3	0.00181892	0	NA	item holdout; capabilities calibrated on training items
4	0.004	0	NA	item holdout; capabilities calibrated on training items
5	0.000911589	0	NA	item holdout; capabilities calibrated on training items
6	0.001	0	NA	item holdout; capabilities calibrated on training items
7	0.00219116	0	NA	item holdout; capabilities calibrated on training items
8	0.001	0	NA	item holdout; capabilities calibrated on training items
9	0.00272626	0	NA	item holdout; capabilities calibrated on training items
10	0.001	0	NA	item holdout; capabilities calibrated on training items
11	0.022182	0	NA	item holdout; capabilities calibrated on training items
12	0.024	0	NA	item holdout; capabilities calibrated on training items
13	0.0164101	0	NA	item holdout; capabilities calibrated on training items
14	0.0155	0	NA	item holdout; capabilities calibrated on training items
15	0.0245499	0	NA	item holdout; capabilities calibrated on training items
16	0.027	0	NA	item holdout; capabilities calibrated on training items
17	0.0279538	0	NA	item holdout; capabilities calibrated on training items
18	0.0305	0	NA	item holdout; capabilities calibrated on training items
19	0.00354449	0	NA	item holdout; capabilities calibrated on training items
20	0.004	0	NA	item holdout; capabilities calibrated on training items
21	0.00598948	0	NA	item holdout; capabilities calibrated on training items
22	0.005	0	NA	item holdout; capabilities calibrated on training items
23	0.0167976	0	NA	item holdout; capabilities calibrated on training items
24	0.0188095	0	NA	item holdout; capabilities calibrated on training items
25	0.00408585	0	NA	item holdout; capabilities calibrated on training items
26	0.0035	0	NA	item holdout; capabilities calibrated on training items

Row	majority mae	copula rho	effective voters	policy
27	0.00486405	0	NA	item holdout; capabilities calibrated on training items
28	0.005	0	NA	item holdout; capabilities calibrated on training items
29	0.00789312	0	NA	item holdout; capabilities calibrated on training items
30	0.007	0	NA	item holdout; capabilities calibrated on training items
31	0.0106654	0	NA	item holdout; capabilities calibrated on training items
32	0.011	0	NA	item holdout; capabilities calibrated on training items
33	0.00534521	0	NA	item holdout; capabilities calibrated on training items
34	0.006	0	NA	item holdout; capabilities calibrated on training items
35	0.0128481	0	NA	item holdout; capabilities calibrated on training items
36	0.011	0	NA	item holdout; capabilities calibrated on training items
37	0.0159856	0	NA	item holdout; capabilities calibrated on training items
38	0.0155	0	NA	item holdout; capabilities calibrated on training items
39	0.000911589	0	NA	item holdout; capabilities calibrated on training items
40	0.0015	0	NA	item holdout; capabilities calibrated on training items
41	0.00219116	0	NA	item holdout; capabilities calibrated on training items
42	0.0015	0	NA	item holdout; capabilities calibrated on training items
43	0.00272626	0	NA	item holdout; capabilities calibrated on training items
44	0.0025	0	NA	item holdout; capabilities calibrated on training items
45	0.00109815	0	NA	item holdout; capabilities calibrated on training items
46	0.002	0	NA	item holdout; capabilities calibrated on training items
47	0.00136632	0	NA	item holdout; capabilities calibrated on training items
48	0.0005	0	NA	item holdout; capabilities calibrated on training items
49	0.00328418	0	NA	item holdout; capabilities calibrated on training items
50	0.0025	0	NA	item holdout; capabilities calibrated on training items
51	0.0164101	0	NA	item holdout; capabilities calibrated on training items
52	0.017	0	NA	item holdout; capabilities calibrated on training items
53	0.0245499	0	NA	item holdout; capabilities calibrated on training items
54	0.027	0	NA	item holdout; capabilities calibrated on training items
55	0.0279538	0	NA	item holdout; capabilities calibrated on training items
56	0.027	0	NA	item holdout; capabilities calibrated on training items

Row	majority mae	copula rho	effective voters	policy
57	0.0186852	0	NA	item holdout; capabilities calibrated on training items
58	0.021	0	NA	item holdout; capabilities calibrated on training items
59	0.0219556	0	NA	item holdout; capabilities calibrated on training items
60	0.025	0	NA	item holdout; capabilities calibrated on training items
61	0.0304146	0	NA	item holdout; capabilities calibrated on training items
62	0.0275	0	NA	item holdout; capabilities calibrated on training items
63	0.00408585	0	NA	item holdout; capabilities calibrated on training items
64	0.0035	0	NA	item holdout; capabilities calibrated on training items
65	0.020777	0	NA	item holdout; capabilities calibrated on training items
66	0.021641	0	NA	item holdout; capabilities calibrated on training items
67	0.0171069	0	NA	item holdout; capabilities calibrated on training items
68	0.0175	0	NA	item holdout; capabilities calibrated on training items
69	0.00558322	0	NA	item holdout; capabilities calibrated on training items
70	0.0055	0	NA	item holdout; capabilities calibrated on training items
71	0.00100224	0	NA	item holdout; capabilities calibrated on training items
72	0.001	0	NA	item holdout; capabilities calibrated on training items
73	0.000447782	0	NA	item holdout; capabilities calibrated on training items
74	0	0	NA	item holdout; capabilities calibrated on training items
75	0.00047888	0	NA	item holdout; capabilities calibrated on training items
76	0.001	0	NA	item holdout; capabilities calibrated on training items
77	0.00102724	0	NA	item holdout; capabilities calibrated on training items
78	0.0025	0	NA	item holdout; capabilities calibrated on training items
79	4.09122e-06	0	NA	item holdout; capabilities calibrated on training items
80	0	0	NA	item holdout; capabilities calibrated on training items
81	0.0238595	0	NA	item holdout; capabilities calibrated on training items
82	0.0223158	0	NA	item holdout; capabilities calibrated on training items
83	0.0296998	0	NA	item holdout; capabilities calibrated on training items
84	0.0312581	0	NA	item holdout; capabilities calibrated on training items
85	0.024216	0	NA	item holdout; capabilities calibrated on training items
86	0.0250714	0	NA	item holdout; capabilities calibrated on training items

Row	majority mae	copula rho	effective voters	policy
87	0.152844	0	NA	item holdout; capabilities calibrated on training items
88	0.153846	0	NA	item holdout; capabilities calibrated on training items
89	0.0541524	0	NA	item holdout; capabilities calibrated on training items
90	0.0540556	0	NA	item holdout; capabilities calibrated on training items
91	0.0474933	0	NA	item holdout; capabilities calibrated on training items
92	0.049	0	NA	item holdout; capabilities calibrated on training items
93	0.187052	0	NA	item holdout; capabilities calibrated on training items
94	0.1875	0	NA	item holdout; capabilities calibrated on training items
95	0.0237054	0	NA	item holdout; capabilities calibrated on training items
96	0.0233158	0	NA	item holdout; capabilities calibrated on training items
97	0.23029	0	NA	item holdout; capabilities calibrated on training items
98	0.230269	0	NA	item holdout; capabilities calibrated on training items
99	0.186473	0	NA	item holdout; capabilities calibrated on training items
100	0.1865	0	NA	item holdout; capabilities calibrated on training items
101	0.00106729	0	NA	item holdout; capabilities calibrated on training items
102	0.0005	0	NA	item holdout; capabilities calibrated on training items
103	0.00114141	0	NA	item holdout; capabilities calibrated on training items
104	0.0015	0	NA	item holdout; capabilities calibrated on training items
105	0.00244841	0	NA	item holdout; capabilities calibrated on training items
106	0.0025	0	NA	item holdout; capabilities calibrated on training items
107	9.76512e-06	0	NA	item holdout; capabilities calibrated on training items
108	0	0	NA	item holdout; capabilities calibrated on training items
109	0.00050996	0	NA	item holdout; capabilities calibrated on training items
110	0	0	NA	item holdout; capabilities calibrated on training items
111	0.00109391	0	NA	item holdout; capabilities calibrated on training items
112	0.0005	0	NA	item holdout; capabilities calibrated on training items
113	4.35739e-06	0	NA	item holdout; capabilities calibrated on training items
114	0	0	NA	item holdout; capabilities calibrated on training items
115	0.00116988	0	NA	item holdout; capabilities calibrated on training items
116	0	0	NA	item holdout; capabilities calibrated on training items

Row	majority mae	copula rho	effective voters	policy
117	4.6607e-06	0	NA	item holdout; capabilities calibrated on training items
118	0	0	NA	item holdout; capabilities calibrated on training items
119	1.0009e-05	0	NA	item holdout; capabilities calibrated on training items
120	0	0	NA	item holdout; capabilities calibrated on training items
121	0.029607	0	NA	item holdout; capabilities calibrated on training items
122	0.0307581	0	NA	item holdout; capabilities calibrated on training items
123	0.024091	0	NA	item holdout; capabilities calibrated on training items
124	0.0220714	0	NA	item holdout; capabilities calibrated on training items
125	0.0758558	0	NA	item holdout; capabilities calibrated on training items
126	0.0764231	0	NA	item holdout; capabilities calibrated on training items
127	0.0046228	0	NA	item holdout; capabilities calibrated on training items
128	0.0035	0	NA	item holdout; capabilities calibrated on training items
129	0.10997	0	NA	item holdout; capabilities calibrated on training items
130	0.110111	0	NA	item holdout; capabilities calibrated on training items
131	0.0884607	0	NA	item holdout; capabilities calibrated on training items
132	0.0884091	0	NA	item holdout; capabilities calibrated on training items
133	0.0258667	0	NA	item holdout; capabilities calibrated on training items
134	0.0235714	0	NA	item holdout; capabilities calibrated on training items
135	0.181308	0	NA	item holdout; capabilities calibrated on training items
136	0.181318	0	NA	item holdout; capabilities calibrated on training items
137	0.141763	0	NA	item holdout; capabilities calibrated on training items
138	0.141357	0	NA	item holdout; capabilities calibrated on training items
139	0.0757532	0	NA	item holdout; capabilities calibrated on training items
140	0.0749231	0	NA	item holdout; capabilities calibrated on training items
141	0.152159	0.552432	NA	item holdout; capabilities calibrated on training items
142	0.0978571	0.552432	NA	item holdout; capabilities calibrated on training items
143	0.165465	0.552432	NA	item holdout; capabilities calibrated on training items
144	0.123615	0.552432	NA	item holdout; capabilities calibrated on training items
145	0.0610208	0.552432	NA	item holdout; capabilities calibrated on training items
146	0.0107143	0.552432	NA	item holdout; capabilities calibrated on training items

Row	majority mae	copula rho	effective voters	policy
147	0.179664	0.552432	NA	item holdout; capabilities calibrated on training items
148	0.1325	0.552432	NA	item holdout; capabilities calibrated on training items
149	0.0979627	0.552432	NA	item holdout; capabilities calibrated on training items
150	0.0295714	0.552432	NA	item holdout; capabilities calibrated on training items
151	0.144524	0.552432	NA	item holdout; capabilities calibrated on training items
152	0.0910698	0.552432	NA	item holdout; capabilities calibrated on training items
153	0.177865	0.552432	NA	item holdout; capabilities calibrated on training items
154	0.118328	0.552432	NA	item holdout; capabilities calibrated on training items
155	0.0713382	0.552432	NA	item holdout; capabilities calibrated on training items
156	0.00357143	0.552432	NA	item holdout; capabilities calibrated on training items
157	0.182021	0.552432	NA	item holdout; capabilities calibrated on training items
158	0.1405	0.552432	NA	item holdout; capabilities calibrated on training items
159	0.0778149	0.552432	NA	item holdout; capabilities calibrated on training items
160	0.0403621	0.552432	NA	item holdout; capabilities calibrated on training items
161	0.190293	0.552432	NA	item holdout; capabilities calibrated on training items
162	0.1275	0.552432	NA	item holdout; capabilities calibrated on training items
163	0.148414	0.552432	NA	item holdout; capabilities calibrated on training items
164	0.0971364	0.552432	NA	item holdout; capabilities calibrated on training items
165	0.189363	0.552432	NA	item holdout; capabilities calibrated on training items
166	0.1455	0.552432	NA	item holdout; capabilities calibrated on training items
167	0.115952	0.552432	NA	item holdout; capabilities calibrated on training items
168	0.0695	0.552432	NA	item holdout; capabilities calibrated on training items
169	0.176512	0.552432	NA	item holdout; capabilities calibrated on training items
170	0.131167	0.552432	NA	item holdout; capabilities calibrated on training items
171	0.154095	0.552432	NA	item holdout; capabilities calibrated on training items
172	0.1025	0.552432	NA	item holdout; capabilities calibrated on training items
173	0.115299	0.552432	NA	item holdout; capabilities calibrated on training items
174	0.0442059	0.552432	NA	item holdout; capabilities calibrated on training items
175	0.190781	0.552432	NA	item holdout; capabilities calibrated on training items
176	0.13	0.552432	NA	item holdout; capabilities calibrated on training items

Row	majority mae	copula rho	effective voters	policy
177	0.155602	0.552432	NA	item holdout; capabilities calibrated on training items
178	0.0958276	0.552432	NA	item holdout; capabilities calibrated on training items
179	0.11726	0.552432	NA	item holdout; capabilities calibrated on training items
180	0.0763333	0.552432	NA	item holdout; capabilities calibrated on training items
181	0.073756	0.702224	NA	item holdout; capabilities calibrated on training items
182	0.112731	0.702224	NA	item holdout; capabilities calibrated on training items
183	0.0996315	0.702224	NA	item holdout; capabilities calibrated on training items
184	0.151778	0.702224	NA	item holdout; capabilities calibrated on training items
185	0.0634175	0.702224	NA	item holdout; capabilities calibrated on training items
186	0.094	0.702224	NA	item holdout; capabilities calibrated on training items
187	0.239429	0.702224	NA	item holdout; capabilities calibrated on training items
188	0.225	0.702224	NA	item holdout; capabilities calibrated on training items
189	0.127454	0.702224	NA	item holdout; capabilities calibrated on training items
190	0.190852	0.702224	NA	item holdout; capabilities calibrated on training items
191	0.113637	0.702224	NA	item holdout; capabilities calibrated on training items
192	0.167154	0.702224	NA	item holdout; capabilities calibrated on training items
193	0.258299	0.702224	NA	item holdout; capabilities calibrated on training items
194	0.254929	0.702224	NA	item holdout; capabilities calibrated on training items
195	0.173171	0.702224	NA	item holdout; capabilities calibrated on training items
196	0.225389	0.702224	NA	item holdout; capabilities calibrated on training items
197	0.306092	0.702224	NA	item holdout; capabilities calibrated on training items
198	0.2945	0.702224	NA	item holdout; capabilities calibrated on training items
199	0.303877	0.702224	NA	item holdout; capabilities calibrated on training items
200	0.290643	0.702224	NA	item holdout; capabilities calibrated on training items
201	0.122122	0.702224	NA	item holdout; capabilities calibrated on training items
202	0.168352	0.702224	NA	item holdout; capabilities calibrated on training items
203	0.0253971	0.702224	NA	item holdout; capabilities calibrated on training items
204	0.0691579	0.702224	NA	item holdout; capabilities calibrated on training items
205	0.144082	0.702224	NA	item holdout; capabilities calibrated on training items
206	0.165786	0.702224	NA	item holdout; capabilities calibrated on training items

Row	majority mae	copula rho	effective voters	policy
207	0.056652	0.702224	NA	item holdout; capabilities calibrated on training items
208	0.0992778	0.702224	NA	item holdout; capabilities calibrated on training items
209	0.199179	0.702224	NA	item holdout; capabilities calibrated on training items
210	0.1865	0.702224	NA	item holdout; capabilities calibrated on training items
211	0.225457	0.702224	NA	item holdout; capabilities calibrated on training items
212	0.224214	0.702224	NA	item holdout; capabilities calibrated on training items
213	0.0875086	0.702224	NA	item holdout; capabilities calibrated on training items
214	0.163352	0.702224	NA	item holdout; capabilities calibrated on training items
215	0.293939	0.702224	NA	item holdout; capabilities calibrated on training items
216	0.291	0.702224	NA	item holdout; capabilities calibrated on training items
217	0.205303	0.702224	NA	item holdout; capabilities calibrated on training items
218	0.209929	0.702224	NA	item holdout; capabilities calibrated on training items
219	0.254306	0.702224	NA	item holdout; capabilities calibrated on training items
220	0.2615	0.702224	NA	item holdout; capabilities calibrated on training items

7.11 Yi - saved analysis tables

pairs.csv

Column panel 1/4 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	a
1	gsm8k	numeric	direct	17	0	YI01
2	mmlu_pro	mc	direct	17	0	YI01
3	mmlu_pro	mc	reasoned	17	0	YI01

Column panel 2/4 - row numbers identify the same source rows.

Row	b	relationship	accuracy matched	n	error a	error b
1	YI02	same_family	False	7	1	0.142857
2	YI02	same_family	False	3	1	0.666667
3	YI02	same_family	False	9	0.666667	0.444444

Column panel 3/4 - row numbers identify the same source rows.

Row	double fault	excess cofailure	phi	adjusted phi	kappa	wrong agreement
1	0.142857	0	NA	NA	0	0
2	0.666667	0	NA	NA	0	0
3	0.444444	0.148148	0.632456	1	0.571429	0.5

Column panel 4/4 - row numbers identify the same source rows.

Row	wrong chance uniform	both wrong n	capa onehot
1	NA	1	NA
2	0.111111	2	-0.08
3	0.125	4	0.430246

ladder.csv

Column panel 1/2 - row numbers identify the same source rows.

Row	benchmark	format	template	seed	temperature	relationship
1	gsm8k	numeric	direct	17	0	same_family
2	mmlu_pro	mc	direct	17	0	same_family
3	mmlu_pro	mc	reasoned	17	0	same_family

Column panel 2/2 - row numbers identify the same source rows.

Row	matched only	pairs	mean adjusted phi	low	high
1	False	1	NA	NA	NA
2	False	1	NA	NA	NA
3	False	1	1	1	1

teams.csv

No rows were recorded in this output.

repeat_pairs.csv

Column panel 1/3 - row numbers identify the same source rows.

Row	model	benchmark	format	condition a	condition b	n
1	YI02	gsm8k	numeric	direct|17|0.0	reasoned|17|0.0	58
2	YI02	mmlu_pro	mc	direct|17|0.0	reasoned|17|0.0	49

Column panel 2/3 - row numbers identify the same source rows.

Row	error a	error b	double fault	excess cofailure	phi	adjusted phi
1	0.0689655	0.103448	0.0689655	0.0618312	0.801234	1
2	0.530612	0.530612	0.510204	0.228655	0.91806	0.91806

Column panel 3/3 - row numbers identify the same source rows.

Row	kappa	wrong agreement	wrong chance uniform	both wrong n	capa onehot
1	0.781955	0.5	NA	4	NA
2	0.91806	0.6	0.114921	25	0.671054

self_consistency.csv

No rows were recorded in this output.

baseline_predictions.csv

No rows were recorded in this output.

8. Definitions, provenance, and publication checklist

Metric dictionary

Row	metric	meaning
1	n	Number of attempted responses (accuracy) or jointly usable questions (pair/team).
2	usable	Answer parsed and not marked truncated; allow_truncated_analysis=False in every run.
3	pipeline_accuracy	Mean correct with parse failures treated as 0 and every truncated response forced to 0.
4	parsed_accuracy	Mean correct among parsed answers, including parsed but truncated outputs. This differs from usable_accuracy.
5	usable_accuracy	Mean correctness among parsed, non-truncated outputs only. This is conditional and can be selection-biased.
6	parse_failure	Fraction with parse_status other than ok. May overlap with truncation.
7	Wilson interval	Saved 95% binomial interval for pipeline or usable accuracy; not adjusted for multiple comparisons or benchmark sampling design.
8	double_fault / beta	Fraction of jointly usable questions where both models / all team members are wrong.
9	selection_ceiling	1 - beta: observed oracle ceiling if a policy must choose one member answer. Not an achievable router score.
10	phi	Pearson correlation of binary errors on the jointly usable subset; undefined with constant error indicators.
11	adjusted_phi	phi divided by the maximum positive correlation compatible with the two marginal error rates.
12	excess_cofailure	Observed double-fault probability minus the product of marginal error probabilities on the same overlap.
13	kappa	Chance-corrected agreement of binary correctness/error states, not answer strings.
14	wrong_agreement	Answer-string agreement restricted to questions where both models are wrong.
15	capa_onehot	Saved one-hot chance-adjusted answer agreement; not soft-probability CAPA. Numeric rows have NA for MC-uniform chance quantities.
16	correctness_majority	Fraction where more than half the members are correct. This uses gold correctness indicators and is not necessarily actual answer plurality.
17	plurality_accuracy	Most frequent answer string, with fractional credit when the gold answer is tied for top count.
18	independence_beta	Product of observed team-member error rates on the team complete-case set.
19	matched_only / accuracy_matched	Uses configured error-rate tolerance in saved analysis; pooled exploratory graph uses a fixed .05 tolerance.
20	NA	Undefined, missing, or not measured. An empty file does not imply a zero-valued result.

Recorded hardware and software environments

Column panel 1/2 - row numbers identify the same source rows.

Row	family	environment file	gpu	vram GiB	capability	cuda
1	Falcon	be80b5bb358ae8ba_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
2	Falcon	c6a0ffcaead602dd_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
3	Falcon	fd0b90b318279afc_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
4	GLM	53681ebbb8d901fd_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
5	GLM	e6ff9ef2f4540778_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
6	Gemma	1db40ff7162175d1_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
7	Gemma	2f8b0d91cb9fd059_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
8	Gemma	590084f68168bc4b_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
9	Granite	29fd7d17b9e281e6_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
10	Granite	7b920552836f81f0_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
11	InternLM	6f1164416a3afdd2_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
12	InternLM	d198a2f6c8214163_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
13	Llama	042d25014942545e_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
14	Llama	21a327f4e17abb33_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
15	Llama	556f689cdce57e30_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
16	Llama	eca6acf5083af268_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
17	Mistral	14eea238693e6d64_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
18	Mistral	28582105f61718f1_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
19	Mistral	6cf128bfb9d28111_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
20	OLMo	ad712a6a59979c84_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2

Row	family	environment file	gpu	vram GiB	capability	cuda
21	OLMo	e24db3528e44505b_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
22	OLMo	fe1bfd7b96d44935_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
23	Phi	1b758908c6879282_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
24	Phi	1e3cc480ab9c4c69_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
25	Phi	cef067ff07620a03_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
26	Phi	dbb27d093b433fc8_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
27	Qwen	0455031383e6ebd5_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
28	Qwen	0b6a685354bed51d_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
29	Qwen	38cb82bdf42ebca7_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
30	Qwen	5a08821ba6730666_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
31	Qwen	72f3909e83e91b06_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
32	Qwen	790e210fa7d463c8_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
33	Qwen	eca68f9e6ce5ee94_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
34	Qwen	f63376572c9909b1_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
35	Qwen	fa298a60157294dd_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
36	Yi	a64710ec40e5c4c3_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2
37	Yi	d2ded87d134371bf_environment.json	NVIDIA RTX PRO 6000 Blackwell Server Edition	94.9724	[12,0]	13.2

Column panel 2/2 - row numbers identify the same source rows.

Row	torch	python	vllm	transformers	pandas
1	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3

Row	torch	python	vllm	transformers	pandas
2	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
3	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
4	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
5	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
6	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
7	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
8	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
9	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
10	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
11	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
12	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
13	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
14	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
15	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
16	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
17	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.31.0	5.14.1	3.0.3
18	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.31.0	5.14.1	3.0.3
19	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.31.0	5.14.1	3.0.3
20	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
21	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3

Row	torch	python	vllm	transformers	pandas
22	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
23	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
24	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
25	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
26	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
27	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.31.0	5.14.1	3.0.3
28	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.31.0	5.14.1	3.0.3
29	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.31.0	5.14.1	3.0.3
30	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.31.0	5.14.1	3.0.3
31	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.31.0	5.14.1	3.0.3
32	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.31.0	5.14.1	3.0.3
33	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.31.0	5.14.1	3.0.3
34	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.31.0	5.14.1	3.0.3
35	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.31.0	5.14.1	3.0.3
36	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3
37	2.13.0+cu132	3.13.11 (main, Feb 3 2026, 03:08:54) [GCC 14.2.0]	0.32.0	5.14.1	3.0.3

The model inventories preserve immutable revisions and repository license labels. license_reviewed is recorded as supplied; this report does not treat a saved label as a completed license review. Dates in response records are UTC completion timestamps; report preparation date is 20 October 2026. Environment records include GPU and software provenance; measured generation times are not full end-to-end runtime or reproducible billing estimates.

Priority 1: retain original outputs, then implement a versioned answer extractor and manually validate a stratified sample per model, benchmark, template, and failure type. Recover unambiguous answers without selecting a number from intermediate reasoning or silently treating ambiguous outputs as correct. Report strict-format compliance separately from answer correctness.

Priority 2: rescore every family with the same validated policy and compare extraction changes. Fix base-model prompt continuation and truncation through a separately versioned pilot, rather than merely increasing token budgets. Do not silently combine altered prompts or budgets with these original runs.

Priority 3: after rescoring, recompute aligned pairs and actual answer-plurality teams with common-question coverage rules, uncertainty, question-level paired tests, and explicit capability matching. Define model-selection rules before examining final held-out outcomes; family labels alone do not establish diversity.

Priority 4: audit lineage from model documentation, add teacher observations or documented cross-family descendants where needed, check identifiability, fit hierarchical difficulty/family/lineage models, and perform held-out item/family/lineage validation. Verification and sampled self-consistency require new experiments because those data are absent here.

Source map: all original run paths are listed in Section 6. Accuracy tables come from accuracy.csv and were independently reconciled with scored.json. Pair/team/baseline tables in Section 7 are original saved outputs. Figures 1-3 derive from the reconciled accuracy records; Figure 4 and Figure 5 are new descriptive computations from scored.json. Integrity results come from manifests, response JSONL files, metadata JSON, and immutable model revisions. The repository README supplied research aims, not empirical evidence for conclusions. No external literature claims are adjudicated in this archive audit.

Companion artifacts: family_report_data/all_accuracy.csv, family_summary.csv, integrity_audit.json, pooled_pairwise.csv, paired_prompt_changes.csv, within_family_team_gains.csv; family_report_figures/ contains the figures, with split versions for PDF readability. Keep these folders next to the Markdown if you want embedded figures and machine-readable drill-down. The PDF embeds every figure and is self-contained. Neither report reproduces all raw response text; those 14,800 records remain in the original archive.